

4. Relational Databases

Fundamentals of Data Management

Dr Charaka Palansuriya

- What is a database
- Keys and constraints
- Design considerations
- Querying a relational database
- Accessing relational databases via applications
- Transactions
- Performance optimisation
- Backup and restore

- Means of storing and accessing data efficiently
- Usually contains a Database Management System (DBMS)
- DBMS provides:
 - Mechanisms to create data structure (e.g., Tables) and content
 - A means of querying and modifying content
 - Ways to optimise performance
 - A tool to backup or archive
 - Ways to allow applications to access data

- There are various types of databases
 - Relational Databases
 - NoSQL Databases
 - Other databases
 - Hybrids,
 - PostgreSQL (object-relational)
- This course will look at two of the most popular and widely used types:
 - Relational
 - NoSQL
- This lecture focus on Relational Databases
 - By far the most common
 - Examples: Oracle, Microsoft SQL Server, MySQL

What is a Relational Database?

- A Collection of tables (i.e., Relations) with associated relationships

Employee

| Name | Department |
|--------|------------|
| Andrew | Computers |
| Bob | Finance |
| Carol | Computers |
| David | Management |

Department

| Department | HeadOfDept | Location |
|------------|------------|-----------|
| Computers | Carol | Edinburgh |
| Finance | Mike | Aberdeen |
| Management | Ethel | Cardiff |

TABLES (*aka Relations*)

RELATIONSHIPS (*aka Associations, References*)

FIELD (*aka Column*)

RECORD (*aka Row, Entry, Tuple*)



- A Primary Key is a field which is guaranteed to hold a unique value in every entry in a table
 - A unique key may be generated automatically by the database management system
- An index is a means of accessing the entries in a table efficiently
 - i.e., indexing is used to improve the performance of querying a database


- Domain Constraints
 - A value for a field must be picked from a particular set (or range) of pre-defined values
- Uniqueness Constraints
 - If a field F is designated as a key, it is not possible to add a record which has a value of F equal to that of any existing record
- Null Constraint
 - Specifies whether or not a field may be Null
- Referential integrity
 - When a table refers to a primary key of another table, i.e., foreign key, a new record can only be added if the record refers to a valid primary key of the other table. For example, if a foreign key is deleted then all records referring to this foreign key must be deleted.

- Sematic Constraints
 - e.g., The salary of an employee must be less than that of his boss
- Dynamic Constraints
 - e.g., The salary of an employee can only increase

- Minimise data redundancy


| Name | Company | Address | Postcode |
|-----------|-------------------|------------------|----------|
| A. Smith | Smith & Son Ltd | 51 High Street | PC7 4LT |
| J. Smith | Smith & Son Ltd | 51 High Street | PC7 4LT |
| T. Jones | Smith & Son Ltd | 51 High St | PC7 4LT |
| N. Dupont | Flash Lighting Co | 14 Howe Crescent | RN4 8PU |

Redundancy




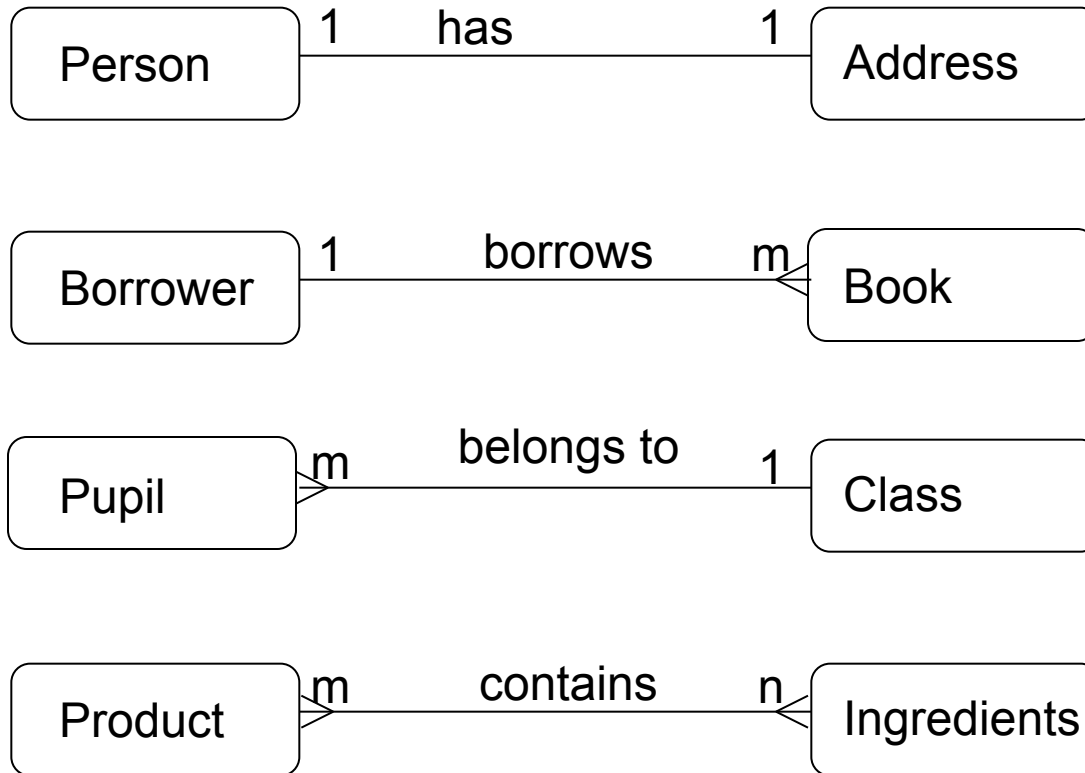
- First Normal Form
 - Eliminate repeating sets of related data in an individual table
 - Create a separate table for each set of related data
 - Identify each set of related data with a primary key
- Second Normal Form
 - Non-primary key columns must depend on the entire primary key, not just part of the primary key
 - Applies when a primary key is based on more than one column
- Third Normal Form
 - Non-primary key columns must depend only on primary key
- Higher Normal Forms
 - Fourth (Boyce Codd) Normal Form and Fifth Normal Form do exist in theory but rarely used in practice

- Binary Relationship




- Tertiary Relationship





- Produces a data model representing a real-world situation
 - Identifies important entities (tables)
 - Relationship between the entities.
- Allows simplification of the data model
 - E.g., Remove many-to-many relationships

Removing Many-to-Many Relationships



- Relational databases are queried using SQL (Structured Query Language)
- Used by all major DBMSs
- Standards exist
 - SQL:2003
 - Dialects exists too
- The result of an SQL query is another table

- `SELECT Name, Department FROM Employee WHERE Department='Computers'`

| Name | Department |
|--------|------------|
| Andrew | Computers |
| Carol | Computers |

- Print out all fields using the * notation:
`SELECT * FROM Employee WHERE Department='Computers'`

- Joins enable extraction of information from more than one table
- `SELECT Name, Employee.Department, Department.Location AS Name, Dept, Location FROM Employee, Department WHERE Employee.Department=Department.Department`

| Name | Dept | Location |
|--------|------------|-----------|
| Andrew | Computers | Edinburgh |
| Bob | Finance | Aberdeen |
| Carol | Computers | Edinburgh |
| David | Management | Cardiff |

- Boolean operators “AND” and “OR” can be used:
 - **SELECT name,department FROM Employee WHERE name='Andrew' OR name='Bob'**
- Brackets can be used to group conditions in the usual way.
- For numerical fields, all of the usual operators exist:
<, >, =, >=, <=, <> (some implementations also accept **!=**)
- Negation can also be performed with NOT.
 - **SELECT student FROM class WHERE NOT(mark1>50)**
is equivalent to
 - **SELECT student FROM class WHERE mark1<=50**

- SQL includes simple pattern matching/wildcards
- Wildcards **%** (zero or more characters) and **_** (exactly one character)
 - Don't work for equality expressions; need to use LIKE keyword.
 - `SELECT name FROM addressbook WHERE name LIKE 'Car%'`
 - Would pick out Caroline, Carl, Carol, cAroLinE
 - `SELECT name FROM addressbook WHERE name LIKE 'A_a_'`
 - Would pick out Adam, Alan (but not Armstrong)

- SQL includes aggregation functions
- SUM, AVG, COUNT, MIN, MAX
- `SELECT count(*) FROM Customers`
 - Will count the number of customers
- Some implementations provide many other functions

- Results (*not* the original tables) can be sorted by one or more fields
 - `SELECT name, salary FROM employees ORDER BY salary, name`
 - Lists all employees ordered by (increasing) salary, and then those with the same salary are listed alphabetically
- Reverse order sorting can be performed with the DESC keyword (for DESCending)
 - `SELECT name, salary FROM employees ORDER BY salary DESC, name`

- Arithmetic can be performed on numerical fields with the usual operators: +, -, *, /.
- The BETWEEN keyword can be used
 - SELECT name,salary FROM employee WHERE salary BETWEEN 10000 AND 20000
 - Above is equivalent to:
 - SELECT name,salary FROM employee WHERE salary >=10000 AND salary <=20000
 - Note that limits *are* included

Accessing Relational Databases via applications

- Databases could be accessed using the Command Line Tool (CLT) provided by a RDBMS
- Typically access to the relational databases are restricted and access provided via applications
- Standards exists for accessing relational databases via applications
 - JDBC: for Java applications
 - ODBC: for example used by Microsoft databases and applications
- Each RDBMS provides drivers which an application should use to access the database
 - For example: MySQL JDBC driver and ODBC driver

- Use Index
 - Index in a book provide a quick way to locate information.
 - Similarly, indexes in a table give a quick way to find relevant rows with specific column values in a query.
 - i.e., instead of scanning entire table row at a time, the query executioner can jump straight to the relevant rows using the index.
 - Typically single column is used for indexing; e.g.
 - Primary key
 - Important field that is most likely to be use in WHERE clause
 - E.g., name in an Employee table

- A sequence of SQL statements executed as a single unit or they are all undone
 - Commonly referred to as a transaction is **committed** or **rolled back**
 - E.g., moving money between bank accounts
- Used for
 - Data consistency
 - When two or more users updating the same data at the same time
 - Data concurrency
 - Allowing two or more users to update data simultaneously

- Transactions are ACID:
 - Atomicity
 - All commands of a transaction is performed or none of them
 - Consistency
 - A transaction takes a database from one consistent state to another consistent state.
 - Isolation
 - When transactions are executed in parallel, the effects of one transaction must not visible to another transaction until the transaction is committed.
 - Durability
 - After a transaction is committed, the changes made by that transaction is permanent. This is important, for example, if a system failure occurs.

- Relational databases support backup and restoring databases
 - Useful for
 - recovering from system failures
 - Archiving data
- For example
 - A MySQL databases can be backed up by using “mysqldump” command
 - Backs up a database to a single file
 - Restore the database using the “source” command

- Good database design requires thought!
- SQL is used to query relational databases
- Database access is generally done via application
- Large databases would require performance optimisation

- Database design basic, <http://office.microsoft.com/en-us/access-help/database-design-basics-HA001224247.aspx>
- Oracle database documentation,
http://docs.oracle.com/cd/E11882_01/
- A relational model of data for large shared data banks, <http://www.seas.upenn.edu/~zives/03f/cis550/codd.pdf>