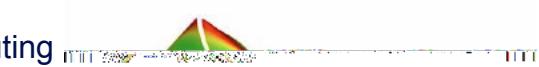


Introduction



- Most common algorithms or computational patterns in scientific computing
- Common across different applications
- Which, when, why (**why not**) in serial and in parallel
- Algorithmic complexity
- Error in numerical computing 
- Verification
 - is my program/code correct?
- Significant mathematical content
 - simple as possible to understand numerical algorithms

S. Lawson, M. Woodgate, R. Steijl, and G. Barakos
<http://www.hector.ac.uk/casestudies/ucav.php>

PNA L01 Introduction 2

Course structure

|epcc|

- 10 weeks - two lectures per week + one exercise class
 - Monday 1000 – 1050 lecture – rm 6301
 - Tuesday 1000 – 1050 exercise – rm 1028
 - Tomorrow general coding session – not PNA specific
 - Wednesday 1000 - 1050 lecture – rm 6301
- Lectures
 - available in Learn before class
 - Print if desired and annotate and elaborate
- Exercise classes
 - available in Learn before class
 - You are expected to complete the problems or practical exercises in your own time
 - Bring your work to the class, Tutors will help you when you get stuck
- Assessment: Exam at the end of the semester

Lecturers

|epcc|

- Four lecturers
 - Chris Johnson - rm 3407 (Course organiser, Matrices, PDEs)
 - Rupert Nash - rm 3409 (Representing numbers on computers)
 - Iain Bethune - rm 2404 (Fast Fourier Transforms)
 - Toni Collis - rm 2410 (Many body systems)

Timetable

- Floating Point numbers:
 - 2 lectures, 1 exercise
- Fast Fourier Transforms
 - 3 lectures, 1 exercise
- Dense linear algebra
 - 2 lectures, 1 exercise
- PDEs
 - 3 lectures, 1 exercise
- N-body methods
 - 3 lectures, 1 exercise
- Sparse linear algebra
 - 3 lectures, 2 exercises
- Full timetable in Learn

PNA L01 Introduction

5

Full PNA timetable

Mon 21-Sep	10:00	Lect:	Introduction to PNA [C]	Mon 26-Oct	10:00	Lect:	Time-dependent PDEs (pollution) [C]
Tues 22-Sep	10:00	Prac:	General MSc practical	Tues 27-Oct	10:00	Prac:	Ex Pollution PDE [C]
Wed 23-Sep	10:00	Lect:	Representing nos. on computers [RN]	Wed 28-Oct	10:00	Lect:	Intro to N-Body methods 1 [C]
Mon 28-Sep	10:00	Lect:	Representing nos. contd. [RN]	Mon 02-Nov	10:00	Lect:	Intro to N-Body methods 2 [TC]
Tues 29-Sep	10:00	Prac:	Ex FP errors, etc	Tues 03-Nov	10:00	Prac:	Ex Orbits, 2-body gravitational system
Wed 30-Sep	10:00	Lect:	Introduction to FFTs [IB]	Wed 04-Nov	10:00	Lect:	Computational chemistry [TC]
Mon 05-Oct	10:00	Lect:	Numerical libraries for FFTs (FFTW) [IB]	Mon 09-Nov	10:00	Lect:	Matrix splitting, Sparse matrices [C]
Tues 06-Oct	10:00	Prac:	Ex FFT image processing	Tues 10-Nov	10:00	Prac:	Ex Sparse mat storage, CG and BiCGSTAB
Wed 07-Oct	10:00	Lect:	Parallel FFTs [IB]	Wed 11-Nov	10:00	Lect:	Intro to Krylov subspace methods [C]
Mon 12-Oct	10:00	Lect:	Dense LA intro [C]	Mon 16-Nov	10:00	Lect:	Advanced KS methods [C]
Tues 13-Oct	10:00	Prac:	Ex LU factorisation by hand/libs	Tues 17-Nov	10:00	Prac:	Ex Time dependent PDE
Wed 14-Oct	10:00	Lect:	LA libraries [C]	Tues 24-Nov	10:00	Prac:	Ex Parallel Matrix-vector exercise [C]
Mon 19-Oct	10:00	Lect:	Discretised PDEs. Pollution model [C]				
Tues 20-Oct	10:00	Prac:	Ex LU fact contd.				
Wed 21-Oct	10:00	Lect:	BVPs. Relaxation methods [C]				

PNA L01 Overview

6

Exercise class problems

|epcc|

- Problems can be done in C or Fortran
 - Only some can be done in Java
 - competent Java programmer should be able to complete exercises in C
 - demonstrate and develop ideas essential to the course
- Significant mathematical content to the course
 - This not examinable, but you are expected to perform simple algebraic manipulation
- Complete all the exercises
 - you will complete the exam

```
Subroutine times_vector_CSR(M, x, y)
! This routine calculates y = M*x for CSR-matrix M and vectors x and y
type(csr), intent(in) :: M
real(kindle16), dimension(M%ncol), intent(in) :: x
real(kindle16), dimension(M%row), intent(out) :: y
integer :: row, col, ctr
integer
! zero the vector before multiplying
y = 0.0
ctr = 1
do row = 1,M%row
    do col = M%rowStart(row), M%rowStart(row+1)-1
        y(row) = y(row) + M%value(ctr)*x(M%colIndex(ctr))
        ctr = ctr + 1
    end do
end do
end subroutine times_vector_CSR
```

PNA L01 Introduction

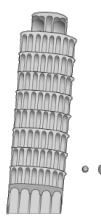
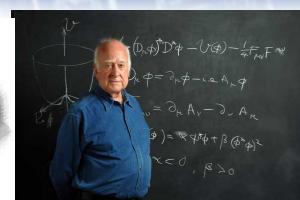
7

The scientific method

|epcc|

Theory: Mathematical equations provide a description or model

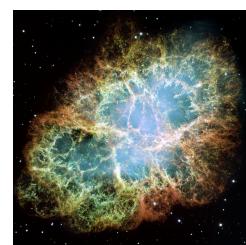
$$\begin{aligned}\nabla \times \mathbf{E} &= -\frac{\partial \mathbf{B}}{\partial t} \\ \nabla \times \mathbf{H} &= \mathbf{J} + \frac{\partial \mathbf{D}}{\partial t} \\ \nabla \cdot \mathbf{D} &= \rho \\ \nabla \cdot \mathbf{B} &= 0\end{aligned}$$



- Experiment
 - Inference from data
 - Test hypothesis

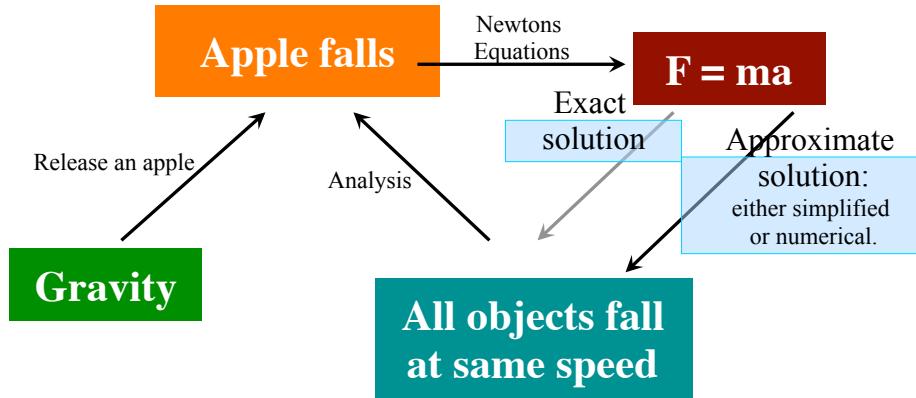
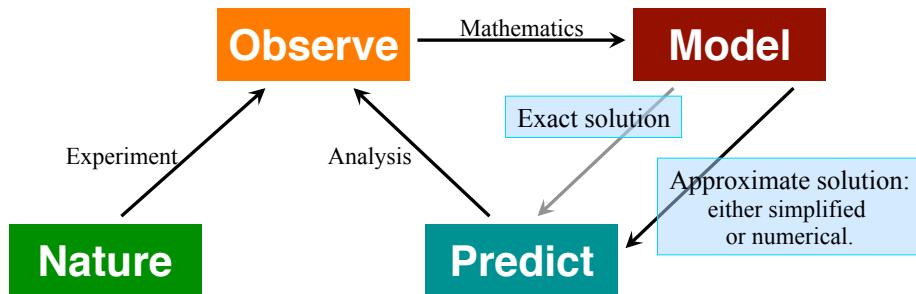
- Computer simulation

- Too big, small, difficult or dangerous for standard methods
- Explore the space of solutions



PNA L01 Introduction

8



The Seven Dwarfs of HPC

|epcc|

- Dense Linear Algebra }
- Sparse Linear Algebra } Arise from discretising PDEs,
solved with different methods
- Spectral Methods
 - transform domain problem, e.g. Fourier Transform (FT)
- N-body Methods
 - Particle-particle e.g. orbits, molecular dynamics
- Structured Grids
 - Data in a regular multidimensional, data updated in sequence
 - Lattice methods
- Unstructured Grids
 - Data in irregular multidimensional grid, triangulation, graph partitionings
- MapReduce
 - Repeated application of a function, final results aggregated
 - Monte Carlo, statistical methods

PNA L01 Introduction

11

What is tractable?

|epcc|

- How big a problem can we tackle?
 - On standard computers
 - Run out of time
 - Run out of space
 - On state-of-the-art machines
 - Run out of time
 - Run out of space
 - Run out of precision



Scientific computations are limited by available computing resources

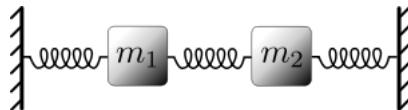
PNA L01 Introduction

12

Complexity

|epcc|

- What size of problem can be tackled depends on the **complexity** of the problem
 - How much **resource** consumed to solve
 - Depends on **problem**
 - Depends on **algorithm**
- System size, N
 - number of **degrees of freedom**
 - number of points simulated
 - dimension of problem (DoF & simulated points)
- 1-d system
- 2 grid points - 1DoF
- **N = 2**



PNA L01 Introduction

13

Big O notation

|epcc|

- How many operations does an algorithm take to solve a system of size N?
 $\Omega = f(N)$ say "**omega is some function of N**"
- We only care about the **dominant** or **leading** behaviour of f
 $f(N) = 3N^3 - 8N^2 + 27N - 3762$
 $f(N) \sim \mathcal{O}(N^3)$ say "**of order N cubed**"
- Examples: For a system of size N
 - Conjugate gradient $\mathcal{O}(N)$ flops (to solve N unknowns of a linear system)
 - LU factorisation takes $\mathcal{O}(N^3)$
 - Bubble sort takes $\mathcal{O}(N^2)$ to sort N numbers
 - Fast Fourier Transform $\mathcal{O}(N \log N)$ to transform a vector of length N

PNA L01 Introduction

14

Performance

|epcc|

- Algorithmic/computation time also depends on number of processing elements (nodes/cores/threads etc), P
- Let $T(N, P)$ be the time taken
 - system of size N
 - P elements
- Speed up, S

$$S(N, P) = \frac{T(N, 1)}{T(N, P)} \quad \cancel{< P} \quad \ll P$$

- Parallel efficiency, E

$$E(N, P) = \frac{S(N, P)}{P} = \frac{T(N, 1)}{P \times T(N, P)} \quad \cancel{< 1} \quad \ll 1$$

PNA L01 Introduction

15

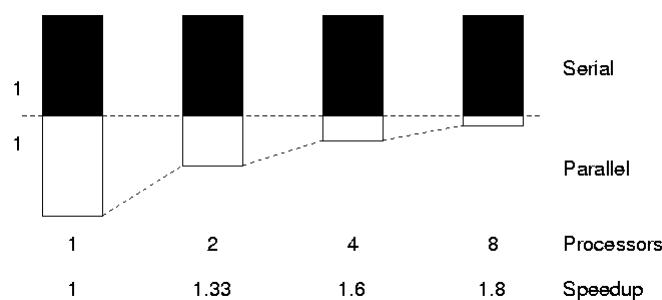
The Serial Component

|epcc|

- Amdahl's law

"the performance improvement to be gained by parallelisation is limited by the proportion of the code which is serial"

Gene Amdahl, 1967



PNA L01 Introduction

16

Amdahl's law

|epcc|

- Assume a fraction α is completely serial
 - time is sum of serial and potentially parallel

- Parallel time

$$T(N, P) = \alpha T(N, 1) + \frac{(1-\alpha)T(N, 1)}{P}$$

- parallel part 100% efficient

- Parallel speedup

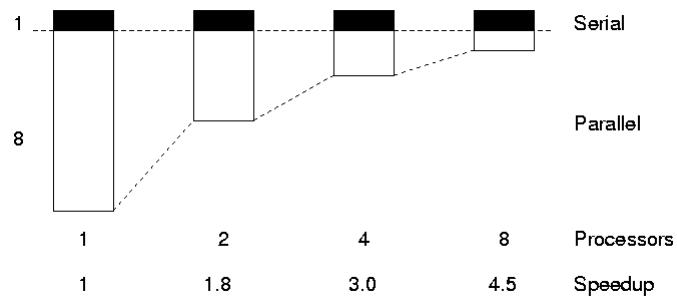
$$S(N, P) = \frac{T(N, 1)}{T(N, P)} = \frac{P}{\alpha P + (1 - \alpha)}$$

- for $\alpha = 0$, $S = P$ as expected (ie $E = 100\%$)
- otherwise, speedup limited by $1/\alpha$ for any P
- impossible to effectively utilise large parallel machines?

Gustafson's Law

|epcc|

- Need larger problems for larger numbers of CPUs



Strong versus Weak scaling

|epcc|

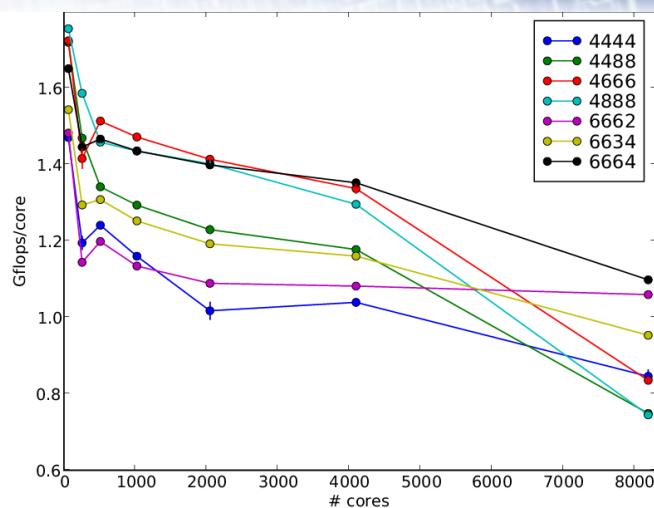
- Two different ways to present scaling data
 - Weak scaling: local volume is fixed
 - system size changes
 - different problem
 - constant efficiency → flat lines on the plot (*c.f.* Gustafson's Law)
 - Strong scaling: global volume is fixed
 - system size is fixed
 - local volume is changing
- $$P \uparrow: \frac{V}{P} \downarrow: \frac{W}{P} \downarrow: C \uparrow$$
- Achieving weak scaling is **hard**
 - Achieving strong scaling is **harder**

PNA L01 Introduction

19

Weak scaling on HECToR XT4

|epcc|



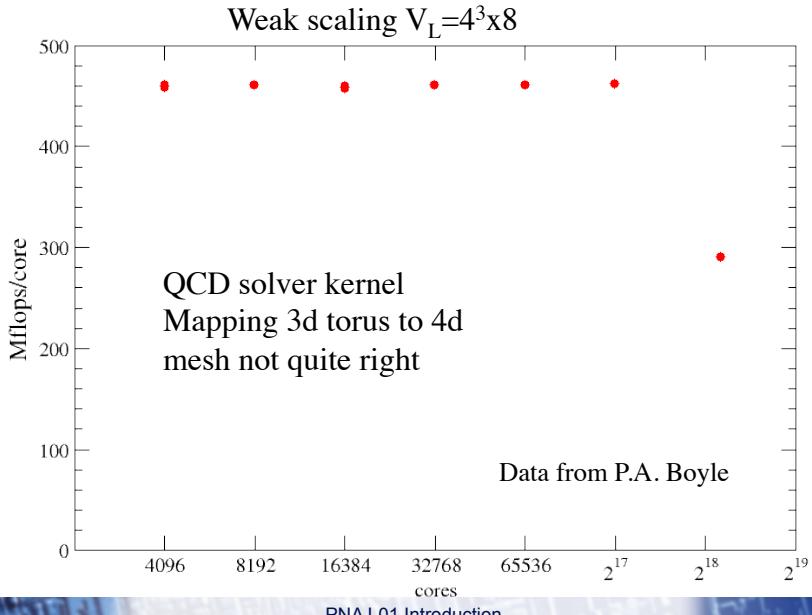
4 dimensional QCD code called Chroma
performance of solver for different local volumes

PNA L01 Introduction

20

Extreme scaling: BG/P @ FZJ Germany

|epcc|



Limits of computation: Complexity theory

|epcc|

- Polynomial time algorithms
 - computation is **polynomial function** of problem size, N
 - Feasible** calculations May still be intractable on an actual computer
- Exponential time algorithms
 - These are **unfeasible** calculations
 - Classic example is integer factorisation
 - Find the two prime factors of an integer $n = pq$
 - Actually **sub-exponential** $O(2^{(\log N)^{1/3}})$
 - Cryptography
 - Public-key system such as RSA
 - 200 digit number to equivalent of 75 years to factorise
 - Quantum computer** reduces this to polynomial time
 - Break public-key encryption!

- No analytic solution
 - No exact answer
- Finite problem, finite model, finite resource
 - Rounding errors
 - Uncertainty in input data - **GiGo**
- How accurate do we need our answer to be
 - How do we know how accurate we are?
- Use examples to illustrate this point
 - Planetary orbits
 - Pollution problem

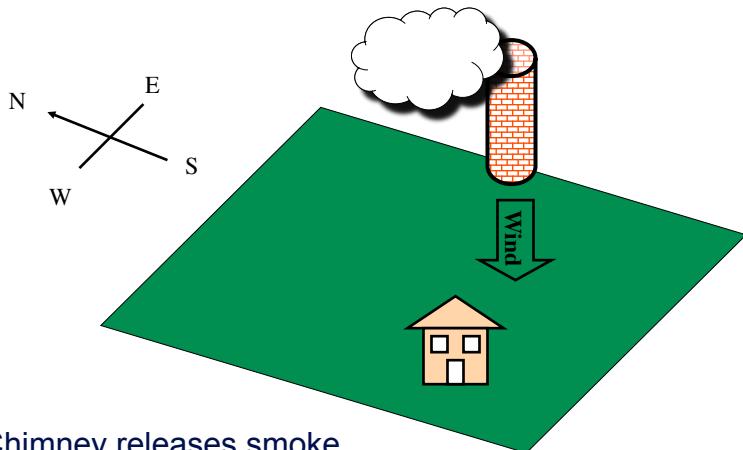
- Problems in **Engineering, Mathematics, Physics, Chemistry, Finance, Biology, Complex systems** represented by **differential** equations [partial, ordinary, integral etc]
- Express relationships between different quantities and how they **change** depending on position in **space and time**
- Example Newton's second Law

$$\vec{F} = m\vec{a} = m \frac{d\vec{v}}{dt} = m \frac{d^2\vec{x}}{dt^2}$$

- Vector ordinary differential equation
 - Ordinary → change of only one variable, time)

Modelling spread of pollution

|epcc|



- Chimney releases smoke
 - how much arrives at house with prevailing north-easterly wind?

Laplace's Equation

|epcc|

- Convection-Diffusion equation
 - Let $u(x,y)$ be the amount of pollution at point (x,y)

$$\left(\frac{\partial^2}{\partial x^2} + \frac{\partial^2}{\partial y^2} \right) u(x, y) = 0$$

- Two dimensional problem, u depends on x and y
- Scalar elliptical partial differential equation
- Often written as

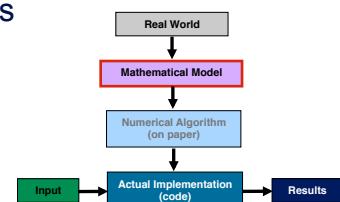
$$\nabla^2 u(x, y) = 0$$

- We can introduce Boundary Conditions

- chimney (source of pollution)

$$\nabla^2 u(x, y) = b(x, y)$$

- Poisson's equation (inhomogeneous eqn)



Linear equations

|epcc|

- Add a wind term

$$\left(a_x \frac{\partial}{\partial x} + a_y \frac{\partial}{\partial y} \right) u(x, y)$$
$$\left(\nabla^2 + \vec{a} \cdot \vec{\nabla} \right) u(x, y) = b(x, y)$$
$$\mathbf{A}u = b$$

Differential operator
Encodes how the system behaves

Linear in u
The solution
Solving Linear equations is **hard!**
Solving non-Linear equations is **harder!**

Boundary conditions
Different for each problem

Linear equations

|epcc|

- Add a wind term

$$\left(a_x \frac{\partial}{\partial x} + a_y \frac{\partial}{\partial y} \right) u(x, y)$$
$$\left(\nabla^2 + \vec{a} \cdot \vec{\nabla} \right) u(x, y) = b(x, y)$$
$$\mathbf{A}u = b$$

Differential operator
Encodes how the system behaves

Linear in u
The solution
Solving Linear equations is **hard!**
Solving non-Linear equations is **harder!**

Boundary conditions
Different for each problem

Linear equations

|epcc|

- Add a wind term

$$\left(a_x \frac{\partial}{\partial x} + a_y \frac{\partial}{\partial y} \right) u(x, y)$$

\downarrow

$$(\nabla^2 + \vec{a} \cdot \vec{\nabla}) u(x, y) = b(x, y)$$
$$\mathbf{A}u = b$$

Differential operator
Encodes how the system behaves

Linear in u
The solution
Solving Linear equations is **hard!**
Solving non-Linear equations is **harder!**

Boundary conditions
Different for each problem

Linear equations

|epcc|

- Add a wind term

$$\left(a_x \frac{\partial}{\partial x} + a_y \frac{\partial}{\partial y} \right) u(x, y)$$

\downarrow

$$(\nabla^2 + \vec{a} \cdot \vec{\nabla}) u(x, y) = b(x, y)$$
$$\mathbf{A}u = b$$

Differential operator
Encodes how the system behaves

Linear in u
The solution
Solving Linear equations is **hard!**
Solving non-Linear equations is **harder!**

Boundary conditions
Different for each problem

Linear equations

|epcc|

- Add a wind term

$$\left(a_x \frac{\partial}{\partial x} + a_y \frac{\partial}{\partial y} \right) u(x, y)$$

$$\left(\nabla^2 + \vec{a} \cdot \vec{\nabla} \right) u(x, y) = b(x, y)$$

$$\mathbf{A}u = b$$

General form

Differential operator

Encodes how the system behaves

Linear in u

The solution

Solving Linear equations is **hard!**

Solving non-Linear equations is **harder!**

Boundary conditions

Different for each problem

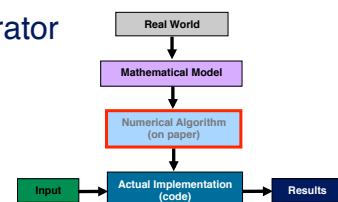
PNA L01 Introduction

31

Numerical solution

|epcc|

- Typically no analytic solution
- Numerical method finds an **approximate** solution
- Replace infinitesimal differential equation with finite difference equation
 - c.f. also finite element method, see later
- Replace continuous domain, D , with a **discrete mesh**, M
 - Mesh spacing, h , $N = M^d$
- Differential operator \rightarrow difference operator
 - many different schemes for this
- u and b \rightarrow vectors of size N
- A is $N \times N$ matrix

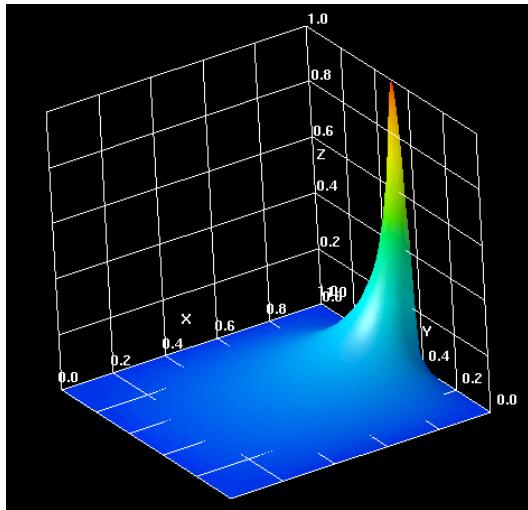


PNA L01 Introduction

32

Solution

|epcc|



Numerical error

Depends on
discretisation scheme

Number of points on
mesh M , or grid
spacing, h

Error ϵ

$$\epsilon = f(h)$$

f scheme dependent

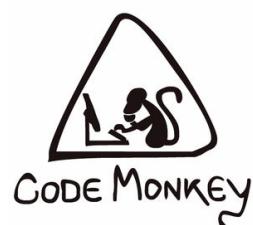
Smaller error \rightarrow more
computation

Implementation

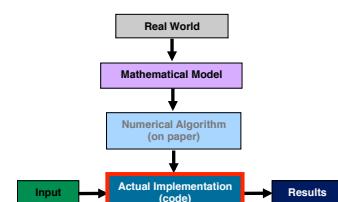
|epcc|



This is your task!
Each exercise looks at
implementing a
different algorithm to
solve the (sometimes
same) problem



Code errors \rightarrow wrong answer
Verification?
Rounding errors see next
lecture



Checking Correctness

|epcc|

- This is a real issue!
 - Verification
- Unit tests - code
 - each module or functional/algorithmic component has a test?
 - what about the tests?
- Integration tests
 - Do the functional pieces interact correctly?
- Replicate known results
 - Analytic solution
 - other codes

Checks & Sums

|epcc|

- Conservation of energy

$$E_K(t) = \sum_i \frac{1}{2} m_i \vec{v}_i(t)$$

$$E_P(t) = \sum_{\langle i,j \rangle} U(|\vec{r}_j(t) - \vec{r}_i(t)|)$$

$$E_T = E_P(t) + E_K(t) = C$$

Constant: Not
time dependant

Similarly Conservation of linear momentum $\vec{P} = \sum_i m_i \vec{v}_i(t)$

Linear momentum should also be conserved in each dimension individually, e.g. x, y, z

Necessary but not sufficient
Implementation could pass, and still be wrong

Verifying the unknown

|epcc|

- Sometimes answer is not known
- Physical symmetries of the problem should be respected
 - simulation should be invariant under changes

$$f(x) \rightarrow f(x') = f(x + a) \quad \text{Translations}$$

$$f(x) \rightarrow f(x') = f(x \cdot \theta) \quad \text{Rotations}$$

$$f(x) \rightarrow f'(x) = f(x) + \phi \quad \text{Rotations of field}$$

- This is problem specific

Necessary but
not sufficient
Implementation
could **pass**, and
still be **wrong**

Summary

|epcc|

- Most problems in Science and Engineering have no analytic solution
- Solve with numerical methods
 - these are approximations
 - Understanding the error
 - Verification
- Introduce the most commonly used algorithmic or computational patterns
- How and when to use them

References

- PA Wiegert, KA Innanen and S Mikkola *An asteroidal companion to the Earth*, *Nature* **387**, 685-686 (12 June 1997)
- V. Szebehely, *Theory of orbits*, (Academic, New York, 1967)
- A View from Berkeley http://view.eecs.berkeley.edu/wiki/Main_Page
- C.H. Papadimitriou, *Computational Complexity*, (Addison-Wesley, Reading, MA 1994)
- Rivest, R.; A. Shamir; L. Adleman (1978). *A Method for Obtaining Digital Signatures and Public-Key Cryptosystems* *Comm. ACM* **21** (2): pp.120–126.
<http://theory.lcs.mit.edu/~rivest/rsapaper.pdf>
- RSA factoring challenge <http://www.rsa.com/rsalabs/node.asp?id=2092>