

Parallel Numerical Algorithms

Lecture 7 Basic Linear Algebra

▶ Lecture will cover

- why matrices and linear algebra are so important
- basic terminology
- Gauss-Jordan elimination
- LU factorisation
- error estimation and iterative improvement

- ▶ In mathematics linear algebra is the study of linear transformations and vector spaces...
- ▶ ...in practice linear algebra is the study of matrices and vectors
- ▶ Many physical problems can be formulated in terms of matrices and vectors

- ▶ Don't let the terminology scare you
 - concepts quite straightforward, algorithms easily understandable
 - implementing the methods is often surprisingly easy
 - but numerous variations (often for special cases or improved numerical stability) lead to an explosion in terminology

Cholesky Decomposition

LU factorisation

Back-substitution

Symmetric Positive Definite

Condition number

Forward-substitution

Gauss-Jordan

Partial pivoting

Crout's algorithm

Singular value decomposition

▶ Matrix

$$A = \begin{pmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{pmatrix}$$

▶ Vector

$$v = \begin{pmatrix} v_1 \\ v_2 \\ v_3 \end{pmatrix}$$

▶ A matrix multiplied by a vector gives another vector

$$Av = w = \begin{pmatrix} a_{11}v_1 + a_{12}v_2 + a_{13}v_3 \\ a_{21}v_1 + a_{22}v_2 + a_{23}v_3 \\ a_{31}v_1 + a_{32}v_2 + a_{33}v_3 \end{pmatrix}$$

▶ Many problems expressible as linear equations

- two apples and three pears cost 40 pence
- three apples and five pears cost 65 pence
- how much does one apple or one pear cost?

▶ Express this as

$$2a + 3p = 40$$

$$3a + 5p = 65$$

▶ Or in matrix form

$$\begin{bmatrix} 2 & 3 \\ 3 & 5 \end{bmatrix} \begin{bmatrix} a \\ p \end{bmatrix} = \begin{bmatrix} 40 \\ 65 \end{bmatrix}$$

- matrix x vector = vector

► For a system of N equations in N unknowns

$$a_{11}x_1 + a_{12}x_2 + \dots + a_{1N}x_N = b_1$$

$$a_{21}x_1 + a_{22}x_2 + \dots + a_{2N}x_N = b_2$$

.

$$a_{N1}x_1 + a_{N2}x_2 + \dots + a_{NN}x_N = b_N$$

- coefficients form a matrix A with elements a_{ij}
- unknowns form a vector x with elements x_i
- solution forms a vector b with elements b_i

► All linear equations have the form $A x = b$

► $A x = b$ implies $A^{-1} A x = x = A^{-1} b$

- simple formulae exist for $N=2$

$$A^{-1} = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{bmatrix}^{-1} = \frac{1}{a_{11}a_{22} - a_{12}a_{21}} \begin{bmatrix} a_{22} & -a_{12} \\ -a_{21} & a_{11} \end{bmatrix}$$

$$\begin{bmatrix} 5 & -3 \\ -3 & 2 \end{bmatrix} \begin{bmatrix} 2 & 3 \\ 3 & 5 \end{bmatrix} \begin{bmatrix} a \\ p \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} a \\ p \end{bmatrix} = \begin{bmatrix} 5 & -3 \\ -3 & 2 \end{bmatrix} \begin{bmatrix} 40 \\ 65 \end{bmatrix} = \begin{bmatrix} 5 \\ 10 \end{bmatrix}$$

► Rarely need (or want) to store the explicit inverse

- usually only require the solution to a particular set of equations

► Algebraic inversion impractical for large N

- use numerical algorithms such as Gaussian Elimination

- Equations are: $\begin{bmatrix} 2 & 3 \\ 3 & 5 \end{bmatrix} \begin{bmatrix} a \\ p \end{bmatrix} = \begin{bmatrix} 40 \\ 65 \end{bmatrix}$
- $2a + 3p = 40$ (i)
- $3a + 5p = 65$ (ii)
- computing $2 \times (ii) - 3 \times (i)$ gives $p = 130 - 120 = 10$
 - substitute in (i) gives $a = 1/2 \times (40 - 3 \times 10) = 5$
- Imagine we actually had
- $2.00000 a + 3.00000 p = 40.00000$ (i)
- $4.00000 a + 6.00001 p = 80.00010$ (ii)
- (ii) - 2 x (i) gives $(6.00001 - 6.00000) p = (80.00010 - 80.00000)$
- cancellations on both sides may give inaccurate numerical results
 - value of p comes from multiplying a huge number by a tiny one
- How can we tell this will happen in advance?

- Characterise a matrix by its *condition number*
- gives a measure of the range of the floating point numbers that will be required to solve the system of equations
- A *well-conditioned* matrix
- has a small condition number
 - and is numerically easy to solve
- An *ill-conditioned* matrix
- has a large condition number
 - and is numerically difficult to solve
- A *singular* matrix
- has an infinite condition number
 - is impossible to solve numerically or analytically

► Easy to compute condition no. for small problems

$$2a + 3p = 40$$

$$3a + 5p = 65$$

- has a condition number of 46 (ratio of largest/smallest eigenvalue)

$$2.00000a + 3.00000p = 40.00000$$

$$4.00000a + 6.00001p = 80.00010$$

- has condition number of 8 million!

► Very hard to compute for real problems

- methods exist for obtaining good estimates

► Gives a measure of the range of the scales of numbers in the problem

- eg if condition number = 46, largest number required in calculation will be roughly 46 times larger than smallest
- if condition number = 10^7 , this may be a problem for single precision where we can only resolve one part in 10^8

► May require higher precision to solve ill-conditioned problems

- in addition to a robust algorithm

► The technique you may have learned at school

- subtract rows of A from other rows to eliminate off-diagonals
- must perform same operations to RHS (i.e. b)

eliminate

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} \xrightarrow{\text{sweep}} a_{ij} \rightarrow a_{ij} - \frac{a_{i1}}{a_{11}} a_{1j}$$

sweep

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ 0 & a'_{22} & a'_{23} & a'_{24} \\ 0 & a'_{32} & a'_{33} & a'_{34} \\ 0 & a'_{42} & a'_{43} & a'_{44} \end{bmatrix}$$

► Pivoting

- using row p as the *pivot row* ($p=1$ above) implies division by a_{pp}
- very important to do row exchange to maximise a_{pp}
- this is *partial pivoting* (full pivoting includes column exchange)

► Gauss-Jordan is a simple *direct* method

- we know the operation count at the outset, complexity $O(N^3)$
- *iterative* methods are optimised for a particular b - see later

► Possible to reduce A to purely diagonal form

- solving a diagonal system is trivial

$$\begin{bmatrix} a'_{11} & 0 & 0 & 0 \\ 0 & a'_{22} & 0 & 0 \\ 0 & 0 & a'_{33} & 0 \\ 0 & 0 & 0 & a'_{44} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} = \begin{bmatrix} b'_1 \\ b'_2 \\ b'_3 \\ b'_4 \end{bmatrix} \rightarrow \begin{aligned} a'_{11}x_1 &= b'_1 \\ a'_{22}x_2 &= b'_2 \\ a'_{33}x_3 &= b'_3 \\ a'_{44}x_4 &= b'_4 \end{aligned}$$

- better to reduce to upper triangular - Gaussian Elimination

- ▶ Operate on active sub-matrix of decreasing size

$$\begin{bmatrix} a'_{11} & a'_{12} & a'_{13} & a'_{14} \\ 0 & a'_{22} & a'_{23} & a'_{24} \\ 0 & a'_{32} & a'_{33} & a'_{34} \\ 0 & a'_{42} & a'_{43} & a'_{44} \end{bmatrix} \rightarrow \begin{bmatrix} a'_{11} & a'_{12} & a'_{13} & a'_{14} \\ 0 & a'_{22} & a'_{23} & a'_{24} \\ 0 & 0 & a'_{33} & a'_{34} \\ 0 & 0 & a'_{43} & a'_{44} \end{bmatrix} \rightarrow \dots$$

- ▶ Solve resulting system with *back-substitution*

– can compute x_4 first, then x_3 , then x_2 , etc...

$$\begin{bmatrix} a'_{11} & a'_{12} & a'_{13} & a'_{14} \\ 0 & a'_{22} & a'_{23} & a'_{24} \\ 0 & 0 & a'_{33} & a'_{34} \\ 0 & 0 & 0 & a'_{44} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} = \begin{bmatrix} b'_1 \\ b'_2 \\ b'_3 \\ b'_4 \end{bmatrix} \quad \begin{aligned} a'_{11}x_1 + a'_{12}x_2 + a'_{13}x_3 + a'_{14}x_4 &= b'_1 \\ a'_{22}x_2 + a'_{23}x_3 + a'_{24}x_4 &= b'_2 \\ a'_{33}x_3 + a'_{34}x_4 &= b'_3 \\ a'_{44}x_4 &= b'_4 \end{aligned}$$

- ▶ Gaussian Elimination is a practical method
 - must do partial pivoting and keep track of row permutations
 - restriction: must start a new computation for every different b
- ▶ Upper-triangular system $Ux = b$ easy to solve
 - likewise for lower-triangular $Lx = b$ using forward-substitution
- ▶ Imagine we could decompose $A = LU$
 - $Ax = (LU)x = L(Ux) = b$
 - first solve $Ly = b$ then $Ux = y$
 - each triangular solve has complexity $O(N^2)$
- ▶ But how do we compute the L and U factors?

► Clearly only have N^2 unknowns

- assume L is *unit* lower triangular and U is upper triangular

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} = \begin{bmatrix} 1 & \cdot & \cdot & \cdot \\ l_{21} & 1 & \cdot & \cdot \\ l_{31} & l_{32} & 1 & \cdot \\ l_{41} & l_{42} & l_{43} & 1 \end{bmatrix} \begin{bmatrix} u_{11} & u_{12} & u_{13} & u_{14} \\ \cdot & u_{22} & u_{23} & u_{24} \\ \cdot & \cdot & u_{33} & u_{34} \\ \cdot & \cdot & \cdot & u_{44} \end{bmatrix}$$

- writing out in full

$$\begin{bmatrix} u_{11} \\ l_{21}u_{11} \\ l_{31}u_{11} \\ l_{41}u_{11} \end{bmatrix} \left| \begin{bmatrix} u_{12} \\ l_{21}u_{12} + u_{22} \\ l_{31}u_{12} + l_{32}u_{22} \\ l_{41}u_{12} + l_{42}u_{22} \end{bmatrix} \right| \begin{bmatrix} u_{13} \\ l_{21}u_{13} + u_{23} \\ l_{31}u_{13} + l_{32}u_{23} + u_{33} \\ l_{41}u_{13} + l_{42}u_{23} + l_{43}u_{33} \end{bmatrix} \left| \begin{bmatrix} \cdot \\ \cdot \\ \cdot \\ \cdot \end{bmatrix} \right.$$

► Can pack LU factors into a single matrix

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} \rightarrow \begin{bmatrix} u_{11} & u_{12} & u_{13} & u_{14} \\ l_{21} & u_{22} & u_{23} & u_{24} \\ l_{31} & l_{32} & u_{33} & u_{34} \\ l_{41} & l_{42} & l_{43} & u_{44} \end{bmatrix}$$

► RHS computed in columns

- once l_{ij} or u_{ij} is calculated, a_{ij} is not needed any more
- can therefore do LU decomposition *in-place*
- elements of A over-written by L and U
- complexity is $O(N^3)$

- ▶ Replaces A by its LU decomposition
 - implements pivoting, ie decomposes row permutation of A
 - computation of l_{ij} requires division by u_{jj}
 - can promote a sub-diagonal l_{ij} as appropriate
 - essential for stability with large N
- ▶ Loop over columns j
 - compute u_{ij} for $i = 1, 2 \dots j$
 - compute l_{ij} for $i = j+1, j+2 \dots N$
 - pivot as appropriate before proceeding to next column
- ▶ For details see Numerical Recipes section 2.3

- ▶ To solve $Ax = b$
 - decompose A into L and U factors via Crout's algorithm
 - replaces A in-place
 - set $x = b$
 - do in-place solution of $Lx = x$ (forward substitution)
 - do in-place solution of $Ux = x$ (backward substitution)
- ▶ Advantages
 - pivoting makes the procedure stable
 - only compute LU factors once for any number of vectors b
 - subsequent solutions are $O(N^2)$ after initial $O(N^3)$ factorisation
 - to compute inverse, solve for a set of N unit vectors b
 - determinant of A can be computed from the product of u_{jj}

► We hope to have solved $Ax = b$

- there will inevitably be errors due to limited precision
- can quantify this by computing the residual vector $r = b - Ax$
- typically quote the root-mean-square *residue*

$$residue = \frac{\|r\|_2}{\|b\|_2}, \quad \|x\|_2 = \sqrt{x^T x} = \sqrt{\sum_{i=1}^N x_i^2}$$

- length defined by L_2 norm (“two-norm”) - other norms exist

► Residue may not be the best measure

- really want difference δx from the perfect solution \hat{x}
 - residual gives the error in the RHS b
- but of course we don’t know the perfect solution!

► Let \hat{x} be perfect solution (which we don’t know!)

- ie $A\hat{x} = b$, so the residual $\hat{r} = b - A\hat{x} = 0$

► Our numerical solution is x

- we have a finite residual $r = b - Ax$
- ie $Ax = b - r$

► If we write $x = \hat{x} + \delta x$

$$Ax = A(\hat{x} + \delta x) = A\hat{x} + A\delta x = b + A\delta x$$

$$\text{but: } Ax = b - r$$

$$\text{so: } A\delta x = -r$$

► Can solve for error δx in numerical solution x

- without knowing the perfect solution!

- ▶ Exact solution $\hat{x} = x - \delta x$
 - but we know that $A \delta x = -r$
 - can easily solve for δx as we already have the LU factors of A
- ▶ Procedure
 - compute LU factors of A and solve $Ax = b$
 - calculate residual vector $r = b - Ax$ and compute residue
 - solve $A \delta x = -r$ and improve the solution $x \rightarrow x - \delta x$
 - re-compute the residue and see how much it has reduced
 - repeat as appropriate
- ▶ Must compute elements r_i in double precision
 - but only need to store using standard precision

- ▶ Dense matrices arise from linear equations
 - standard notation is $Ax = b$
- ▶ Matrices characterised by their condition number
 - equations difficult to solve numerically have large condition number
 - an ill-conditioned matrix
 - may lead to large errors in our solution
 - so **always compute the residue!**
- ▶ Have covered direct solution methods for $Ax = b$
 - all are basically variants of Gaussian Elimination
 - rather than storing A^{-1} , compute the LU factors of A
 - can then solve further equations $Ax = c$, $Ax = d$, ... at little extra cost
 - the larger the condition number, the harder the problem
 - pivoting is essential in practice for numerical stability

- ▶ Solve a dense system of equations
 - you are given a template code in C, Fortran or Java
- ▶ The codes should compile and run as they stand

```
$ pgf90 -o lufact lufact.f90
$ ./lufact

$ pgcc -o lufact lufact.c
$ ./lufact

$ javac lufact.java
$ java lufact
```
- ▶ Supplied code doesn't do anything useful
 - generates a random matrix A and a right-hand-side b
 - but **you** must supply LU decomposition and forward/backward substitution code to solve for x
 - as supplied, codes simply set L and U to zero and $x = b$

- ▶ For simplicity
 - do not implement partial pivoting
 - the system is constructed to have a solution $x_i = 1$
 - unlike real problems, easy to check if you have the right answer!
- ▶ Debugging small problems ($N \leq 6$)
 - code writes out L , U , L multiplied by U , x and Ax
 - if your LU decomposition is correct
 - L will be unit lower triangular, U will be upper triangular
 - L times U will be equal to the original matrix A
 - if your forward/backward substitution is correct
 - x will contain numbers all very close to 1.0 (the solution)
 - Ax will be close to b (since we are trying to solve $Ax = b$)
 - you **MUST** quantify this by computing the residual

- ▶ You should try to
 - get the right answer!
 - quantify the error by computing the residual
 - for this test case we can also compare to the error in the solution
 - see how the residual (a measure of the error) varies with N
 - investigate iterative improvement
 - look at the effect of using double precision

- ▶ Following slides are for information only
 - you are not expected to understand the maths

- ▶ Can compute the *eigenvalues* of A
 - a small eigenvalue indicates a numerical problem
 - an $N \times N$ matrix has, in general, N eigenvalues $\lambda_1, \lambda_2, \dots, \lambda_N$.
 - conventionally order them so that $|\lambda_1| < |\lambda_2| < \dots < |\lambda_{N-1}| < |\lambda_N|$
 - eigenvalues for first problem on slide 7 are 0.15 and 6.85
 - eigenvalues for second problem are 0.000001 and 8.000009
- ▶ A matrix is, for most purposes, “equivalent” to the diagonal matrix formed by the eigenvalues

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} & a_{14} \\ a_{21} & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & a_{34} \\ a_{41} & a_{42} & a_{43} & a_{44} \end{bmatrix} \approx \begin{bmatrix} \lambda_1 & 0 & 0 & 0 \\ 0 & \lambda_2 & 0 & 0 \\ 0 & 0 & \lambda_3 & 0 \\ 0 & 0 & 0 & \lambda_4 \end{bmatrix}$$

- ▶ Individual entries in matrix have little meaning
 - values can be changed by re-scaling equations
 - structure can be changed by re-labelling variables
- ▶ Important global characteristics
 - determinant
 - condition number
- ▶ Can both be expressed in terms of eigenvalues
 - unfortunately, computing the eigenvalues is much harder than solving linear equations!
 - we will see that we can compute the determinant in other ways

► Determinant

- the product of the eigenvalues λ_i
- a zero eigenvalue implies a zero determinant
- comes from a singular matrix, ie a set of equations that are *impossible* to solve (or no unique solution)

► Condition number

- given by ratio of maximum to minimum eigenvalue: λ_N / λ_1
- poorly conditioned matrices have a large condition number
 - and a small determinant
- correspond to equations that are *numerically difficult* to solve

► This is an *Applied* course so

- for most practical purposes a matrix can be thought of as being a diagonal matrix made up of the eigenvalues
- a theorem true for diagonal matrices is probably true in general!

► When does $A^M = AxAxA \dots$ converge to zero?

- when all the eigenvalues have absolute value < 1

$$\begin{bmatrix} \lambda_1 & \cdot & \cdot & \cdot \\ \cdot & \lambda_2 & \cdot & \cdot \\ \cdot & \cdot & \lambda_3 & \cdot \\ \cdot & \cdot & \cdot & \lambda_4 \end{bmatrix}^M = \begin{bmatrix} \lambda_1^M & \cdot & \cdot & \cdot \\ \cdot & \lambda_2^M & \cdot & \cdot \\ \cdot & \cdot & \lambda_3^M & \cdot \\ \cdot & \cdot & \cdot & \lambda_4^M \end{bmatrix}$$

► Often have very nice properties, eg

- Symmetric: $a_{ij} = a_{ji}$ which guarantees real eigenvalues
- Positive Definite: all eigenvalues are greater than zero
- for an SPD matrix $x^T A x$ is positive for all vectors x

$$\begin{bmatrix} x_1 & x_2 & x_3 & x_4 \end{bmatrix} \begin{bmatrix} \lambda_1 & \cdot & \cdot & \cdot \\ \cdot & \lambda_2 & \cdot & \cdot \\ \cdot & \cdot & \lambda_3 & \cdot \\ \cdot & \cdot & \cdot & \lambda_4 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} = \lambda_1 x_1^2 + \lambda_2 x_2^2 + \dots$$

► Physical significance of eigenvalues / vectors

- engineering: frequency of vibration / direction of movement
- quantum mechanics: energy of electron / distribution in space

► If A is Symmetric Positive Definite

- can decompose $A = U^T U$
- Crout's algorithm is simpler as $l_{ij} = u_{ji}$
- called a *Cholesky* decomposition
 - ability to do Cholesky actually proves that A is SPD

► If A is singular or very ill-conditioned

- can do a Singular Value Decomposition
- SVD is, in some sense, the “best” of all possible inverses
- beyond the scope of this course