

## Parallel Numerical Algorithms

### Solution of Boundary Value Problems

1

- ▶ General solution methods

- ▶ Relaxation methods

- Jacobi algorithm
- testing for convergence
- Gauss Seidel
- over-relaxation

- ▶ Notes

- parallelisation
- non-linear equations

- ▶ Pollution problem

- solution using relaxation methods
- 2D equations including wind

2

▶ Direct methods

- give the solution after a fixed number of operations
  - Gaussian elimination
  - LU factorisation

▶ Relaxation methods (this lecture)

- gradually improve solution, starting from an initial guess
- stop when the answer is sufficiently accurate
- simple to implement but may be slow to converge on solution
  - or may fail completely!

▶ Krylov subspace methods (following lectures)

- iterative (like relaxation methods) but more sophisticated
- harder to implement but more efficient and reliable

3

▶ Direct methods explicitly operate on the matrix  $A$

- eg decompose it into  $L$  and  $U$  factors

▶ For PDEs,  $A$  is very sparse indeed

- may contain 99% zeros so clearly we use compressed storage
- we want to take advantage of this when we solve equations

▶ Difficult to exploit sparsity for direct methods

- eg  $L$  and  $U$  may be dense even though  $A$  is sparse
- for large systems of equations, we may run out of memory!

▶ Relaxation and Krylov methods exploit sparsity

- relaxation methods operate on the equations not the matrix
- Krylov methods comprise mostly matrix-vector multiplications
  - can write efficient routines to do  $y = Ax$  when  $A$  is sparse

4

- ▶ Operate directly on the difference equations
  - can forget (almost!) all about the matrix representation  $Au = b$  for this lecture
  - it turns out that relaxation methods can usefully be understood in terms of matrix-vector operations (not immediately obvious)
    - See lecture on “Matrix Splitting Techniques”
- ▶ For illustrative purposes, look at 1D problem
  - for simplicity with no wind
  - exercise will involve extending this to the 2D problem
    - quite straightforward in practice

5

- ▶ 1D diffusion equations are
  - $-u_{i-1} + 2u_i - u_{i+1} = 0, \quad i = 1, 2, \dots, N$
- ▶ Equivalently:  $u_i = \frac{1}{2} (u_{i-1} + u_{i+1})$ 
  - why not make an initial guess at the solution
  - then loop over each lattice point  $i$  and set  $u_i = \frac{1}{2} (u_{i-1} + u_{i+1})$
  - ie we solve the equation exactly at each point in turn
- ▶ Updating  $u_i$  spoils solution we just did for  $u_{i-1}$ 
  - so simply iterate the whole process again and again ...
  - ... and hope we eventually get the right answer!
- ▶ This is called the Jacobi Algorithm
  - the simplest possible relaxation method

6

- ▶ Use superscript  $n$  to indicate iteration number
  - $n$  counts the number of times we update the *whole* solution
  - equivalent to computer time

- ▶ Jacobi algorithm for diffusion equation is:

$$u_i^{(n+1)} = \frac{1}{2} (u_{i-1}^{(n)} + u_{i+1}^{(n)})$$

- ▶ Each iteration, calculate  $u^{(n+1)}$  in terms of  $u^{(n)}$ 
  - don't need to keep copies of all the previous solutions
  - only need to remember two solutions at any time:  $u$  and  $u_{new}$ 
    - corresponding to iterations  $n$  and  $n+1$

7

```

declare arrays:    u(0, 1, ..., M+1)
                  unew(0, 1, ..., M+1)

initialise: set boundaries: u(0)    = fixed value  $u_0$ 
                              u(M+1) = fixed value  $u_{M+1}$ 
                    initial guess: u(1, 2, ..., M) = guess value

loop over n = 1, 2, ...

    update:    loop over internal points: i = 1, 2, ... M
                unew(i) = 0.5*( u(i-1) + u(i+1) )
                end loop over i

    copy back: u(1, 2, ..., M) = unew(1, 2, ..., M)

end loop over n

```

8

### ► Array declarations

- Fortran: `real, dimension(0:M+1) :: u`
- Java: `float[] u = new float[M+2];`
- C: `float u[M+2];`





### ► Arrays explicitly contain boundaries $u_0$ and $u_{M+1}$

- we set them according to boundary conditions
  - but we NEVER update them!
- eg when we copy  $u_{new}$  back to  $u$ , only copy internal values
- in pseudo-code, boundary values for  $u_{new}$  are never set
  - complete solution is therefore only ever present in  $u$
  - might be more elegant to set boundaries in  $u_{new}$  as well

### ► What to choose for initial guess $u_i^{(0)}$ ?

- for a simple implementation just set interior values to zero

9



10

- ▶ The solution *appears* to be getting better
  - must quantify this!
- ▶ For dense systems we used the residual
  - we tried to solve  $Ax=b$ , so  $r = b-Ax$  should be a zero vector
  - in practice, there is a numerical error in solution of each equation
  - error in equation  $i$  is the value of  $r_i$ 
    - residual is computed from the sum of the squares of  $r_i$
- ▶ Can do the same thing for relaxation methods
  - compute the sum of the squares of the error in each equation
  - do this at the end of each iterative loop over  $n$ 
    - stop if this is small enough

11

```
loop over n = 1, 2, ...
  update:    ...
  copy back: ...

  compute residue: rnorm = 0.0
                  loop over i = 1, 2, ..., M
                    rnorm = rnorm + (-u(i-1)+2*u(i)-u(i+1))2
                  end loop over i
                  rnorm = sqrt(rnorm)

  normalise:    res = rnorm / bnorm


  if (res < tolerance) finish

end loop over n
```

12

- ▶ For a perfect solution, residue will be zero
  - in practice we will get a finite value
  - usually stop when it is “small”, eg a tolerance of  $\text{res} < 10^{-6}$
  - there will be a limit to how small the residual can get
    - can easily hit the limits of single precision
    - use double precision everywhere (or at least perform residual calculation using doubles)
- ▶ Normalisation
  - need to divide by the norm of the  $b$  vector
  - we saw before that  $b$  corresponds to the boundary values
  - in 1D:  $\text{bnorm} = \sqrt{u(0) * u(0) + u(M+1) * u(M+1)}$ 
    - in 2D, need to sum values of squares of  $u_{i,j}$  over all edges

13



- ▶ Decreases exponentially
  - with a zero initial guess for  $u$ , should equal 1.0 at iteration zero

14

- ▶ Very simple for Jacobi
- ▶ Decompose the problem domain regularly across processes/threads
  - for MPI we need halo regions due to  $i+1$ ,  $i-1$  references
  - halos are 1 cell wide for 5-point stencil
  - could be wider for larger stencils
  - swap halos between neighbouring processes every iteration
- ▶ Require global sums for, eg, residue calculation

15

- ▶ About to cover some variations on Jacobi
  - which we hope will be faster!
- ▶ How can we tell if a method will work at all?
- ▶ Necessary (but not sufficient) condition
  - if the method arrives at the correct solution it must stay there
- ▶ Is this true for Jacobi?  $u_i^{(n+1)} = \frac{1}{2} (u_{i-1}^{(n)} + u_{i+1}^{(n)})$ 
  - for a solution:  $-u_{i-1}^{(n)} + 2u_i^{(n)} - u_{i+1}^{(n)} = 0$ , ie  $\frac{1}{2} (u_{i-1}^{(n)} + u_{i+1}^{(n)}) = u_i^{(n)}$
  - so,  $u_i^{(n+1)} = u_i^{(n)}$  and we stay at the solution
    - worth checking this for other methods

16



► Why do we need both  $u_{new}$  and  $u$  ?


```
update:    loop over internal points: i = 1, 2, ... M
           unew(i) = 0.5*( u(i-1) + u(i+1) )
           end loop over i
copy back: u(1, 2, ..., M) = unew(1, 2, ..., M)
```

► Why not do the update in place?

```
update:    loop over internal points: i = 1, 2, ... M
           u(i) = 0.5*( u(i-1) + u(i+1) )
           end loop over i
```

– this is called the *Gauss-Seidel* method

17



► Converges twice as fast as Jacobi


– for less work and less storage!

18

- ▶ Order of the update loop is now significant
  - we used normal (*lexicographic*) order: other orderings possible
- ▶ Red-black order divides grid into chequerboard
  - update all the red squares first then all the black ones
  - enables Gauss Seidel method to be parallelised


Processor 1

Processor 2



19

- ▶ Recall how Jacobi solution progressed



- we have increased the value of  $u_i$  by a small amount
  - but we know the real solution is even higher
- why not increase by more than suggested
  - ie multiply the change by some factor  $w > 1$

20

► Gauss-Seidel method:  $u_i = \frac{1}{2} (u_{i-1} + u_{i+1})$

– ie:  $u_i = u_i + [\frac{1}{2} (u_{i-1} - 2u_i + u_{i+1})]$

► Multiply change (in square brackets) by  $w$

– over-relaxed update:  $u_i = u_i + \frac{1}{2} w (u_{i-1} - 2u_i + u_{i+1})$

– or  $u_i = (1-w) u_i + \frac{1}{2} w (u_{i-1} + u_{i+1})$

► Notes

– original method corresponds to  $w = 1$

– if we get to a solution we stay there for any value of  $w$

21

► Relaxation methods deal directly with equations

– doesn't matter that we cannot express them as  $Au = b$

– equally valid for non-linear equations (eg fluid dynamics)

► Non-linear equations can be very unstable

– may need to under-relax to get convergence, ie  $w < 1$

22

► Initialise

- set boundary values (grey)
  - zero on top, bottom and left
  - hump function on right
- zero interior (white)

► Loop over interior


- $i = 1, 2, \dots, M$
- $j = 1, 2, \dots, M$

- update  $u_{i,j}$  as appropriate

► Repeat until converged

► Write results

- include boundaries so that the solution looks nice!




23

► How do we convert from  $(i,j)$  to  $(x,y)$  coordinates?

- for a domain of size 1x1:
  - $x = i * h$  and  $y = j * h$

► What is the hump function?

- $u(1.0,y) = k * (y_2 - y)^2 * (y - y_1)^2$
- a peak, centred at  $(y_2 + y_1)/2$ , dropping to zero for  $y < y_1$  and  $y > y_2$
- for this example, take  $y_1 = 0.4$  and  $y_2 = 0.6$




24

► How do we convert from  $(x,y)$  to  $(i,j)$  coordinates?

- eg what lattice point do we look at to find  $u(0.20,0.33)$ ?
- $(0.20,0.33)$  is unlikely to fall exactly on a lattice point

- the four nearest neighbours are:

- $i = \text{int}(x/h)$
- $j = \text{int}(y/h)$



- do weighted average of these four values (see exercise notes)

25

► 1D Gauss-Seidel update

$$u_i = \left( \frac{1}{2 + ah} \right) (u_{i-1} + (1 + ah) u_{i+1})$$

► 1D Over-Relaxed update

$$u_i = (1-w)u_i + w \left( \frac{1}{2 + ah} \right) (u_{i-1} + (1 + ah) u_{i+1})$$

► 2D Discrete Equations


$$-u_{i,j-1} - u_{i-1,j} + (4 + (a_x + a_y)h)u_{i,j} - (1 + a_x h)u_{i+1,j} - (1 + a_y h)u_{i,j+1} = 0$$

$(a_x, a_y)$  = wind strength from  $x$  (East) and  $y$  (North) respectively

26

- ▶ Have multiplied all the equations by  $h^2$ 
  - equations now explicitly depend on  $h$  for a non-zero wind  $a$
  - straightforward to derive update equations for 2D case
- ▶ A different convention for Krylov methods
  - maintain the  $1/h^2$  factor in matrix  $A$ 
    - therefore need to multiply RHS by same factor
    - happens to be more convenient
- ▶ Finite wind
  - matrix  $A$  is now non-symmetric
  - in 1D, lower-diagonal elements are  $(1+ah)$ , upper elements are 1
  - gives some minor technical issues when normalising the residue
    - see notes
    - if correctly normalised, residue at zero iterations will *always* be 1.0 if the initial guess is a zero solution

27



28

- ▶ Relaxation methods
  - guess at an initial solution
  - update many times and stop when residue is small enough
- ▶ Update rule is very straightforward
  - solve exactly for each individual  $u_i$ 
    - obtain formula by rearranging difference equations so  $u_i$  is on the LHS
- ▶ Interior points updated according to the PDE
  - boundary points set by the boundary conditions
- ▶ Jacobi is the simplest method
  - Gauss Seidel acts “in-place” and requires roughly half the iterations
  - appropriate over-relaxation can accelerate this even more
    - finding the best value of  $w$  requires some experimentation!