


Overview – two short lectures 

- **Lecture 1 Matrix Splitting Techniques**
 - relaxation techniques as matrix methods
 - analysis of convergence
 - mapping PDE solution domains to vectors

- **Lecture 2 Sparse Matrices**
 - A classification of linear systems
 - A history of large matrix computations
 - Sources of linear systems
 - Sparse matrices and PDEs
 - Sparse matrix storage
 - The COO and CSR storage technique
 - Matrix operations

PNA L14 2

Summary of Relaxation Methods

- So far presented as a pragmatic approach
 - write down the discrete PDE that each solution point must satisfy
 - loop over every interior point and solve the PDE locally
 - apply the boundary conditions by hand
 - iterate many times and hope for convergence
- Some extensions to this Jacobi method
 - eg in-place solution (Gauss-Seidel) and over-relaxation
- Not obviously related to matrix methods

Matrix Splitting

- Any linear problem is of the form $Ax = b$
 - A encodes the precise form of the PDE
 - b contains any fixed boundary conditions
 - Split A into three parts
 - Diagonal, Strictly Upper and Strictly Lower triangular: $A = L + D + U$
 - *not* the same as the LU factors!
- $$(L + D + U)x = b, \quad Dx = -(L + U)x + b$$
- view these as iterative expressions, eg Jacobi corresponds to

$$Dx^{(n+1)} = -(L + U)x^{(n)} + b$$
$$x^{(n+1)} = -D^{-1}(L + U)x^{(n)} + D^{-1}b$$

Consider 1D Pollution Problem

- A represents: $-d^2/dx^2$

$$A = \begin{bmatrix} 2 & -1 & \cdot & \cdot \\ -1 & 2 & -1 & \cdot \\ \cdot & -1 & 2 & -1 \\ \cdot & \cdot & -1 & 2 \end{bmatrix}$$

- Splitting into L , D and U

– a Jacobi iteration is: $D x^{(n+1)} = -(L+U) x^{(n)} + b$

$$\begin{bmatrix} 2 & \cdot & \cdot & \cdot \\ \cdot & 2 & \cdot & \cdot \\ \cdot & \cdot & 2 & \cdot \\ \cdot & \cdot & \cdot & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix}^{(n+1)} = \begin{bmatrix} \cdot & 1 & \cdot & \cdot \\ 1 & \cdot & 1 & \cdot \\ \cdot & 1 & \cdot & 1 \\ \cdot & \cdot & 1 & \cdot \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix}^{(n)} + \begin{bmatrix} b_1 \\ 0 \\ 0 \\ b_4 \end{bmatrix}$$

Jacobi Equations

- Equations the same as in previous lectures
 - with u_i replaced by x_i
 - exterior boundary values u_0 and u_{N+1} replaced by b_1 and b_N

$$\begin{aligned} x_1^{(n+1)} &= \frac{1}{2} (b_1 + x_2^{(n)}) \\ x_2^{(n+1)} &= \frac{1}{2} (x_1^{(n)} + x_3^{(n)}) \\ x_3^{(n+1)} &= \frac{1}{2} (x_2^{(n)} + x_4^{(n)}) \\ x_4^{(n+1)} &= \frac{1}{2} (x_3^{(n)} + b_4) \end{aligned}$$

- Procedure

- impose PDE at each interior point
- new value is the average of the old neighbouring points

- Keep both D and L on the LHS: $(D+L) x^{(n+1)} = -U x^{(n)} + b$

$$\begin{bmatrix} 2 & \cdot & \cdot & \cdot \\ -1 & 2 & \cdot & \cdot \\ \cdot & -1 & 2 & \cdot \\ \cdot & \cdot & -1 & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix}^{(n+1)} = \begin{bmatrix} \cdot & 1 & \cdot & \cdot \\ \cdot & \cdot & 1 & \cdot \\ \cdot & \cdot & \cdot & 1 \\ \cdot & \cdot & \cdot & \cdot \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix}^{(n)} + \begin{bmatrix} b_1 \\ 0 \\ 0 \\ b_4 \end{bmatrix}$$

$$\begin{aligned} x_1^{(n+1)} &= \frac{1}{2} (b_1 + x_2^{(n)}) \\ x_2^{(n+1)} &= \frac{1}{2} (x_1^{(n+1)} + x_3^{(n)}) \\ x_3^{(n+1)} &= \frac{1}{2} (x_2^{(n+1)} + x_4^{(n)}) \\ x_4^{(n+1)} &= \frac{1}{2} (x_3^{(n+1)} + b_4) \end{aligned}$$

- equivalent to solving Jacobi equations *in-place* in order 1, 2, ..., N

- Connection to matrix-splitting
 - Jacobi

$$x^{(n+1)} = -D^{-1}(L + U)x^{(n)} + b$$

- Gauss-Seidel

$$x^{(n+1)} = -(D + L)^{-1}Ux^{(n)} + b$$


- Over-relaxed Gauss-Seidel

$$x^{(n+1)} = (D + wL)^{-1}((1 - w)D - wU)x^{(n)} + wb$$

2D Problem as a Matrix Problem

- Easy to write down difference equations in 2D
 - eg: $-u_{2,1} - u_{1,2} + 4u_{2,2} - u_{3,2} - u_{2,3} = 0$
 - but how do we write these as $Au = b$?
- Map $M \times M$ solution u_{ij} to a vector x_i of length N
 - mapping has no effect on the solution
 - but some may be more convenient than others
 - commonly use lexicographic order, ie $u_{ij} \rightarrow x_{i+(j-1)*M}$

- Consider 3x3 problem
 - nine unknowns
 - solution x has nine elements



Focus on Equation for $u_{2,2}$

- u equations:

$$-\nabla^2 u_{2,2} = -u_{2,1} - u_{1,2} + 4u_{2,2} - u_{3,2} - u_{2,3}$$
 - in terms of x :

$$-\nabla^2 x_5 = -x_2 - x_4 + 4x_5 - x_6 - x_8$$

- row 5 of matrix:

$$\begin{bmatrix}
 \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\
 \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\
 \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\
 \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\
 0 & -1 & 0 & -1 & 4 & -1 & 0 & -1 & 0 \\
 \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\
 \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\
 \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\
 \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot & \cdot
 \end{bmatrix}
 \begin{bmatrix}
 x_1 \\
 x_2 \\
 x_3 \\
 x_4 \\
 x_5 \\
 x_6 \\
 x_7 \\
 x_8 \\
 x_9
 \end{bmatrix}
 =
 \begin{bmatrix}
 \cdot \\
 \cdot \\
 \cdot \\
 \cdot \\
 0 \\
 \cdot \\
 \cdot \\
 \cdot \\
 \cdot
 \end{bmatrix}$$

Choice of Ordering

- Has some effect on performance
 - may want to reflect red/black nature of chequerboard update
 - ie try to achieve linear access patterns on x
- Has a major effect in parallel
 - parallel matrix-vector performed by decomposing over vectors
 - eg $y = Ax$ is parallelised by regular domain decomposition of y
 - amount of communications determined by form of mapping
 - optimising the mapping is a *mesh decomposition* problem
 - standard algorithms exist (recursive spectral bisection, ...)

Note: Convergence of Splitting

- General matrix splitting eqns: $E x^{(n+1)} = F x^n + b$
 - actual solution is perfect solution plus correction
 - $x^n = \hat{x} + \delta x^n$ where $E \hat{x} = F \hat{x} + b$
 - substituting into main equation gives $E \delta x^{n+1} = F \delta x^n$
 - error in solution evolves according to $\delta x^n = (E^{-1}F)^n \delta x^0$
- Convergence depends on eigenvalues of $E^{-1}F$
 - must all be less than one in order to get a solution
 - speed of convergence depends on condition number
- Iteration matrix
 - $E^{-1}F$ is $-D^{-1}(L+U)$ for Jacobi and $-(D+L)^{-1}U$ for Gauss Seidel
 - can show that latter is better conditioned
 - heuristically, GS inverts more of the matrix at each step

Summary

- Relaxation methods are very easy to program
 - may achieve good performance using over-relaxation
- Amount to matrix splitting for linear problems
 - allows for a formal analysis of convergence properties
 - equivalence of GS algorithm not immediately apparent
- Can be applied to non-linear problems
 - no longer equivalent to matrix method
 - under-relaxation may be required for stability

Sparse Matrices

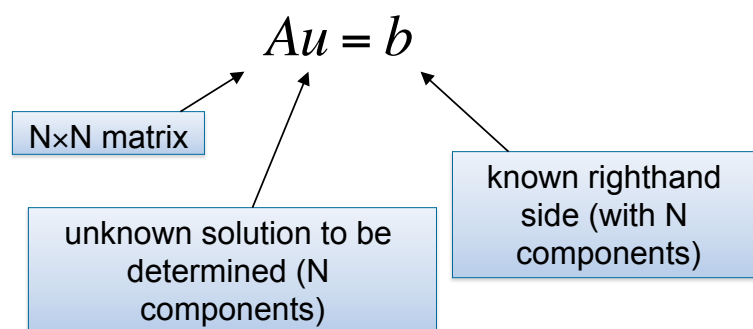
Introduction to Sparse Linear Algebra

Overview

- Motivation
 - A classification of linear systems
 - A history of large matrix computations
 - Sources of linear systems
 - Sparse matrices and PDEs
 - Sparse matrix storage
 - The COO and CSR storage technique
 - Matrix operations

Recall what a linear system is

- Recall that a linear system (of size **N**) can be represented by a matrix equation, of the form:



- This is simply a representation of **N** equations linking **N** unknown quantities: u_1, u_2, \dots, u_N .

Linear equations

- with a wind term $\left(a_x \frac{\partial}{\partial x} + a_y \frac{\partial}{\partial y} \right) u(x, y)$

$$(\nabla^2 + a\nabla)u(x, y) = b(x, y)$$

$$\mathbf{A}u = b$$

Differential operator

Encodes how the system behaves

Linear in u

The solution

Solving Linear equations is *hard*!

Solving non-Linear equations is *harder*!

Boundary conditions

Different for each problem

Linear equations

- with a wind term $\left(a_x \frac{\partial}{\partial x} + a_y \frac{\partial}{\partial y} \right) u(x, y)$

$$(\nabla^2 + a\nabla)u(x, y) = b(x, y)$$

$$\mathbf{A}u = b$$

Differential operator

Encodes how the system behaves

Linear in u

The solution

Solving Linear equations is *hard*!

Solving non-Linear equations is *harder*!

Boundary conditions

Different for each problem

Linear equations

- with a wind term $\left(a_x \frac{\partial}{\partial x} + a_y \frac{\partial}{\partial y} \right) u(x, y)$

$$(\nabla^2 + a\nabla)u(x, y) = b(x, y)$$

$$\mathbf{A}u = b$$

Differential operator

Encodes how the system behaves

Linear in u

The solution

Solving Linear equations is **hard**!

Solving non-Linear equations is **harder**!

Boundary conditions

Different for each problem

Linear equations

- with a wind term $\left(a_x \frac{\partial}{\partial x} + a_y \frac{\partial}{\partial y} \right) u(x, y)$

$$(\nabla^2 + a\nabla)u(x, y) = b(x, y)$$

$$\mathbf{A}u = b$$

Differential operator

Encodes how the system behaves

Linear in u

The solution

Solving Linear equations is **hard**!

Solving non-Linear equations is **harder**!

Boundary conditions

Different for each problem

Linear equations

- with a wind term $\left(a_x \frac{\partial}{\partial x} + a_y \frac{\partial}{\partial y} \right) u(x, y)$

$$(\nabla^2 + a\nabla)u(x, y) = b(x, y)$$

$$\mathbf{A}u = b$$

General form

Differential operator

Encodes how the system behaves

Linear in u

The solution

Solving Linear equations is **hard**!

Solving non-Linear equations is **harder**!

Boundary conditions

Different for each problem

Classification of systems

- Small dense systems
 - are easily stored as dense arrays
 - can be efficiently factorised, e.g. using LU or QR method
 - are well-understood and easily solved, e.g. LAPACK.
- Large sparse systems
 - Contain large numbers of unknowns
 - Possess significant structure e.g. bandedness
 - cannot be efficiently stored in dense arrays
 - present significant practical difficulties.

Computational cost

1950	N=20	(Wilkinson)
1965	N=200	(Forsythe and Moler)
1980	N=2,000	(LINPACK)
1995	N=20,000	(LINPACK)
2002	N=300,000	(LINPACK)

Table 1: Typical matrix computations over last 50 years - an increase in size of $O(10^4)$.


Viable matrix computations

- Matrix factorisation takes $O(N^3)$ flops
 - to solve a problem of N unknowns, containing N^2 coefficients
- Computing power has increased by $O(10^{12})$ in last fifty years
- Notice $(10^4)^3 = 10^{12}$.
- If we could solve system in $O(N^2)$ flops, maybe tackle problems with tens of millions of unknowns!

Sources of linear systems

- Small dense systems ($<O(10^4)$ unknowns)
 - forces between bodies in a mechanical structure
 - interaction between chemicals in a reaction
- Large sparse systems ($>O(10^4)$ unknowns)
 - models from CFD, financial markets, biology, particle physics, environmental:
discretisation of PDEs
 - finite difference, finite element, or finite volume methods

Gear box model: 153,746 unknowns, 9,080,404 nonzeros, T. Davis, University of Florida Sparse Matrix Collection




QCD

- Continuous space-time is replaced with a 4D lattice
- For a lattice with $24^3 \times 48$ space-time points with 4 spin and 3 colour components each matrix size is
 - $7,962,624 \times 7,962,624 = 63,403,380,965,376$ elements
 - This corresponds to 1,014,454,095,446,016 Bytes ≈ 1 PB c.f. HECToR (ARCHER predecessor) which had total of 90 TB RAM
- The matrix represents a Dirac operator
 - essentially a complicated differential operator so mostly we're dealing with nearest-neighbour interactions.


Wilson-Dirac matrix

- Wilson-Dirac operator
 - Hermitian
 - Well-conditioned
 - Well-structured
 - Sparse!



- SM Pickles PhD thesis UoE 1998


Example PDE: spread of pollution



Spread of pollution -- solution of large, sparse linear system.

One-dimensional pollution model

Solution might look like:




PNA L14

29

Discretisation of model

Replace continuous equation by discrete set of values, defined on a mesh ($M=4$):



PNA L14

30

Finite difference matrix

$$\begin{bmatrix} d & u & 0 & & 0 \\ l & d & \ddots & & \\ 0 & \ddots & \ddots & \ddots & \\ & & \ddots & \ddots & 0 \\ 0 & & & 0 & l & d \end{bmatrix} \begin{bmatrix} x_1 \\ \vdots \\ \vdots \\ x_N \end{bmatrix} = \begin{bmatrix} -lb_1 \\ 0 \\ \vdots \\ 0 \\ -ub_N \end{bmatrix}$$

$$l = -\frac{1}{h^2}, \quad d = 2\frac{1}{h^2} + \frac{a}{h}, \quad u = -\frac{1}{h^2} - \frac{a}{h}$$

Sparse matrix storage

- As majority of entries in matrix A are zero, it is inefficient to store a sparse matrix in a 2-D array.
 - makes elimination of redundant calculations more difficult
 - large amounts of memory for storage
- Dozens of storage techniques -- we consider 2 popular choices:
 - COO format (most obvious technique)
 - CSR/ CSC format (as used by Sun S3L, MATLAB, PETSc etc)

COO format

- Co-ordinate (COO) storage involve 3 scalars and 3 arrays:

int nRow	- number of rows
int nCol	- number of columns
int nNz	- number of nonzero elements
double value[nNz]	- contents of matrix location
int rowIdx[nNz]	- row index of each element
int colIdx[nNz]	- likewise for columns

- Data usually in row-major order
- Not particularly efficient
 - redundant storage of column entries (when ordered)
 - easy for user, not optimised for programming
 - expensive to add/modify data or change capacity
 - very easy to perform matrix transpose operation

Example of COO format

- Matrix

$$A = \begin{pmatrix} 1.0 & 0.1 & 0.5 & \bullet \\ 0.4 & \bullet & \bullet & 0.7 \\ \bullet & 0.6 & \bullet & \bullet \\ 0.3 & \bullet & 0.2 & 0.8 \end{pmatrix}$$

```
nRow = 4; nCol = 4; nNz = 9;
value[9] = {1.0, 0.1, 0.5, 0.4, 0.7, 0.6,
0.3, 0.2, 0.8};
colIdx[9] = {0, 1, 2, 0, 3, 1, 0, 2, 3};
rowIdx[9] = {0, 0, 0, 1, 1, 2, 3, 3, 3};
```

Matrix-vector multiplication

```
public void timesv(double[] v, double[] w){  
    for(int j=0; j<this.nRow; j++){  
        w[j] = 0.0;  
  
        for(int j=0; j<this.nNz; j++){  
            w[this.rowIdx[j]] +=  
                this.value[j]*v[this.colIdx[j]];  
        }  
    }  
}
```

double indirection for assignment and data access

CSR format

- Compressed sparse row (CSR) format - 3 scalars and 3 arrays (data stored in row-major order).
 - int nRow - number of rows
 - int nCol - number of columns
 - int nzMax - maximum number of elements
 - double value[nzMax] - contains nonzero entries, stored in row-major order
 - int colIdx[nzMax] - columns index of corresponding element in value
 - int rowStart[nRow+1] - index into value[] and colIdx[] of the first entry in each row
 - (nNz = rowStart[nRow])
- Good for matrix-vector multiplication, and reasonable for transpose operation.
- Flexible design in terms of expansion, modification.

Example of CSR format

- Matrix

$$A = \begin{pmatrix} 1.0 & 0.1 & 0.5 & \bullet \\ 0.4 & \bullet & \bullet & 0.7 \\ \bullet & 0.6 & \bullet & \bullet \\ 0.3 & \bullet & 0.2 & 0.8 \end{pmatrix}$$

```
nRow = 4; nCol = 4; nzMax = 9;  
value[9] = {1.0, 0.1, 0.5, 0.4, 0.7, 0.6,  
0.3, 0.2, 0.8};  
colIdx[9] = {0, 1, 2, 0, 3, 1, 0, 2, 3};  
rowStart[5] = {0, 3, 5, 6, 9};
```

Matrix-Vector multiplication

```
public void timesv(double[] v, double[] w){  
    for(int j=0; j<this.nRow; j++){  
        w[j] = 0.0;  
        for(row=0; this.nrow; row++){  
            for(col=this.rowStart[row];  
                col<this.rowStart[row+1]; col++){  
                w[row] += this.value[j] * v[this.colIdx[j]];  
                j++;  
            }  
        }  
    }  
    return;  
}
```

double indirection for update access only

CSC format

- Compressed sparse column (CSC) similar to CSR
 - int nRow - number of rows
 - int nCol - number of columns
 - int nzMax - maximum capacity
 - double value[nZmax] - nonzero entries in column-major format
 - int rowIndex[nZmax] - row index for each corresponding element in value[]
 - int colStart[nCol+1] - index of first element in each column
- used by packages/libraries derived from FORTRAN which store arrays in column-major format, e.g. MATLAB
- analogous format to CSR matrices

Conclusions

- Large, sparse, linear systems are widely used in scientific computing
- Generally result from discretising PDEs
- Contain significant structure
- Require specialised storage techniques
- Require specialised solution techniques - most commonly iterative methods