

# Parallel Programming with Fortran Coarrays

---

MSc in HPC

David Henty, Alan Simpson (EPCC)  
Harvey Richardson, Bill Long (Cray)

# Coarrays Overview

- Parallel Programming with Fortran Coarrays
  - Fortran Programming Model
  - Basic coarray features
  - Brief mention of Coarray C++
- Further coarray features
- Advanced coarray features
- Experiences with coarrays

# The Fortran Programming Model

---

# Motivation

- Fortran now supports parallelism as a full first-class feature of the language
- Changes are minimal
- Performance is maintained
- Flexibility in expressing communication patterns

# Programming models for HPC

- The challenge is to efficiently map a problem to the architecture we have
  - Take advantage of all computational resources
  - Manage distributed memories etc.
  - Optimal use of any communication networks
- The HPC industry has long experience in parallel programming
  - Vector, threading, data-parallel, message-passing etc.
- We would like to have models or combinations that are
  - efficient
  - safe
  - easy to learn and use


# Why consider new programming models?

- Next-generation architectures bring new challenges:
  - Very large numbers of processors with many cores
  - Complex memory hierarchy
  - We had to deal with 500k cores in 2011
- Parallel programming is hard, need to make this simpler
- Some of the models we currently use are
  - bolt-ons to existing languages as APIs or directives
  - Hard to program for underlying architecture
  - unable to scale due to overheads
- So, is there an alternative to the models prevalent today?
  - Most popular are OpenMP and MPI ...

# Fortran 2008 coarray model


- Example of a Partitioned Global Address Space (PGAS) model
- Set of participating processes like MPI
- Participating processes have access to **local memory** via standard program mechanisms
- Access to **remote memory** is directly supported by the language

# Fortran coarray model





# Fortran coarray model



`a = b[3]`

# Fortran coarrays

- Remote access is a full feature of the language:
  - Type checking
  - Opportunity to optimize communication
- No penalty for local memory access
- Single-sided programming model more natural for some algorithms
  - and a good match for modern networks with RDMA

# Fortran coarrays

---

## Basic Features

(including very basic introduction to coarray C++)

# Coarray Fortran

"Coarrays were designed to answer the question:

*'What is the smallest change required to convert Fortran into a robust and efficient parallel language?'*

The answer: a simple syntactic extension.

It looks and feels like Fortran and requires Fortran programmers to learn only a few new rules."

John Reid,  
ISO Fortran Convener

# Some History

- Introduced in current form by Numrich and Reid in 1998 as a simple extension to Fortran 95 for parallel processing
- Many years of experience, mainly on Cray hardware
- A set of core features are now part of the Fortran standard ISO/IEC 1539-1:2010
- Additional features are expected to be published in a Technical Specification in due course.

# How Does It Work?

- SPMD - Single Program, Multiple Data
  - single program replicated a fixed number of times
- Each replication is called an *image*
- Images are executed asynchronously
  - execution path may differ from image to image
  - some situations cause images to synchronize
- Images access remote data using *coarrays*
- Normal rules of Fortran apply

# What are coarrays?

- Arrays or scalars that can be accessed remotely
  - images can access data objects on any other image
- Additional Fortran syntax for coarrays
  - Specifying a codimension declares a coarray

```
real, dimension(10), codimension[*] :: x  
real :: x(10) [*]
```

- these are equivalent declarations of a array x of size 10 on each image
- x is now remotely accessible
- coarrays have the same size on each image!

# Accessing coarrays

```
integer :: a(4) [*], b(4) [*] !declare coarrays  
b(:) = a(:) [n]             ! copy
```

- integer arrays a and b declared to be size 4 on all images
- copy array a from remote image n into local array b
- `()` for local access `[]` for remote access
- e.g. for two images and `n = 2`:

image 1

a	1	2	3	4
b	8	9	3	8

image 2

a	2	9	3	7
b	12	19	32	13



# Synchronisation

- Be careful when updating coarrays:
  - If we get remote data was it valid?
  - Could another process send us data and overwrite something we have not yet used?
  - How do we know that data sent to us has arrived?
- Fortran provides synchronisation statements
- For example, barrier for synchronisation of all images:

```
sync all
```

- do not make assumptions about execution timing on images
  - unless executed after synchronisation
  - Note there is implicit synchronisation at program start

# Retrieving information about images

- Two intrinsics provide index of this image and number of images
  - `this_image()` (image indexes start at 1)
  - `num_images()`

```
real :: x[*]  
if(this_image() == 1) then  
  read *,x  
  do image = 2,num_images()  
    x[image] = x  
  end do  
end if  
sync all
```

# Making remote references

- We used a loop over images

```
do image = 2,num_images()  
  x[image] = x  
end do
```

- Note that array indexing within the coindex is not allowed so we can not write

```
x[2:num_images()] = x      ! illegal
```


# Data usage

- coarrays have the same size on every image
- Declarations:
  - round brackets () describe rank, shape and extent of local data
  - square brackets [] describe layout of images that hold local data
- Many HPC problems have physical quantities mapped to n-dimensional grids
- You need to implement your view of global data from the local coarrays as Fortran does not provide the global view
  - You can be flexible with the coindexing (see later)
  - You can use any access pattern you wish

# Data usage

- print out a 16 element “global” integer array A from 4 processors
  - 4 elements per processor = 4 coarrays on 4 images


```
integer :: ca(4) [*]  
do image=1,num_images()  
  print *,ca(:)[image]  
end do
```



# 1D cyclic data access

- coarray declarations remain unchanged
  - but we use a cyclic access pattern

```
integer :: ca(4)[*]  
do i=1,4  
  do image=1,num_images()  
    print *,ca(i)[image]  
  end do  
end do
```



# Synchronisation

- code execution on images is independent
  - programmer has to control execution using synchronisation
- synchronise before accessing coarrays
  - ensure content is not updated from remote images before you can use it
- synchronise after accessing coarrays
  - ensure new content is available to all images
- implicit synchronisation after variable declarations at first executable statement
  - guarantees coarrays exist on all images when your first program statement is executed
- We will revisit this topic later

# Example: maximum of array

```
real :: a(10)
real :: maximum[*]
call random_number(a)
maximum = maxval(a)
sync all
if (this_image() == 1) then
  do image = 2, num_images()
    maximum = max(maximum, maximum[image])
  end do
  do image = 2, num_images()
    maximum[image] = maximum
  end do
end if
sync all
```

implicit synchronisation

ensure all images set local maximum

ensure all images have copy of maximum value



# Recap

We now know the basics of coarrays

- declarations
- references with []
- `this_image()` and `num_images()`
- `sync all`

Now consider a full program example...

## Example2: Calculate density of primes

```
program pdensity
  implicit none
  integer, parameter :: n=10000000, nimages=8
  integer start,end,i
  integer, dimension(nimages) :: nprimes[*]
  real density

  start = (this_image()-1) * n/num_images() + 1
  end = start + n/num_images() - 1

  nprimes(this_image())[1] = num_primes(start,end)

  sync all
```

## Example2: Calculate density of primes

```
...  
if (this_image()==1) then  
  
    nprimes(1)=sum(nprimes)  
    density=real(nprimes(1))/n  
    print *, "Calculating prime density on", &  
    &          num_images(), "images"  
    print *, nprimes(1), 'primes in', n, 'numbers'  
    write(*, ' (" density is ", 2Pf0.2, "%") ') density  
    write(*, ' (" asymptotic theory gives ", &  
    &          2Pf0.2, "%") ') 1.0/(log(real(n))-1.0)  
  
end if
```

## Example2: Calculate density of primes

Calculating prime density on 2 images  
664580 primes in 10000000 numbers  
density is 6.65%  
asymptotic theory gives 6.61%

# Coarray C++

---

A basic introduction to syntax for people who know the fundamentals of Fortran coarrays

# Coarray C++

- Brings the coarray programming model to C++
- Announced by Cray in 2013
  - Response to customers challenged by parallelism in C++
  - There was no similar model for C++ (UPC is based on C)
- The implementation is based on template and class definitions from an include file and shared runtime with Cray PGAS languages
- No change to the C++ language/compiler
- As per the Fortran coarray model:
  - 'Images' execute in SPMD fashion
  - Other familiar aspects map to C++ idioms

# Declaring a coarray in C++

- Coarrays are arrays or scalars that can be accessed remotely
  - images can access data objects on any other image
- New syntax to declare coarrays

```
#include <coarray_cpp.h>
using namespace coarray_cpp;

coarray<float[10]> x ;
```

- Declaration of an array x of size 10 on each image
- coarrays have the same size on each image!

# Accessing coarrays

```
coarray<int[4]> a,b;    // declare coarrays  
  
for(i=0;i<4;i++) a[i] = b(n)[i] ; // copy
```

- int arrays a and b declared to be size 4 on all images
- copy array a from remote image n into local array b
- `()` for remote access `[]` for local access
- e.g. for two images and `n = 1`:

image 0

a	1	2	3	4
b	8	6	3	8

image 1

a	2	9	3	7
b	10	19	32	13



# Synchronisation

- Analogously to the Fortran model you can synchronize all the images
- For example, barrier for synchronisation of all images:

```
sync_all();
```

# Retrieving information about images

- Two functions provide index of this image and number of images
  - `this_image()` (image indexes start at 0)
  - `num_images()`

```
coarray<float> x;  
if(this_image() == 0) {  
    std::cin >> x ;  
    for(im=1;im<num_images() ;im++){  
        x(im) = x ;  
    }  
    sync_all();  
}
```

## Example2 prime density: revisited in C++

```
#define NMAX 1000000
#define MAXIMAGES 32

#include <coarray_cpp.h>
using namespace coarray_cpp;
// ...

coarray<int[MAXIMAGES]> nprimes;

start = this_image()*NMAX / num_images() + 1;
end = start + NMAX/num_images() - 1;

nprimes(0)[this_image()] = num_primes(start,end) ;
sync_all();
```

# Launching a coarray program

- The Fortran standard does not specify how a program is launched
- The number of images may be set at compile, link or run-time
- A compiler could optimize for a single image

## Examples on Linux

- Cray XC  
**aprun -n 16000 solver**
- Intel compiler  
**export FOR\_COARRAY\_NUM\_IMAGES=2**  
**./solver**

# Observations so far on coarrays

- Natural extension, easy to learn
- Makes parallel parts of program obvious (syntax)
- Part of Fortran language (type checking, etc)
- No mapping of data to buffers (or copying) or creation of complex types (as we might have with MPI)
- Compiler can optimize for communication
- More observations later...

# Exercise Session 1

- Look at the *Exercise Notes* document for full details
- Write, compile and run a “Hello World” program that prints out the value of the running image’s image index and the number of images
- Extend the simple Fortran code provided in order to perform operations on parts of a picture using coarrays


# Additional Slides

---

## Comparison of Programming Models

# Shared-memory directives and OpenMP


memory





# OpenMP: work distribution


memory



```
!$OMP PARALLEL DO  
do i=1,32  
    a(i)=a(i)*2  
end do
```

# OpenMP implementation

memory




cpus



# Shared Memory Directives

- Multiple threads share global memory
- Most common variant: OpenMP
- Program loop iterations distributed to threads, more recent task features
  - Each thread has a means to refer to private objects within a parallel context
- Terminology
  - Thread, thread team
- Implementation
  - Threads map to user threads running on one SMP node
  - Extensions to distributed memory not so successful
- OpenMP is a good model to use within a node

# Cooperating Processes Models




PROBLEM




processes

# Message Passing, MPI

process




# MPI



# Message Passing

- Participating processes communicate using a message-passing API
- Remote data can only be communicated (sent or received) via the API
- MPI (the Message Passing Interface) is the standard
- Implementation:  
MPI processes map to processes within one SMP node or across multiple networked nodes
- API provides process numbering, point-to-point and collective messaging operations
- Mostly used in two-sided way, each endpoint coordinates in sending and receiving

# SHMEM






# SHMEM

- Participating processes communicate using an API
- Fundamental operations are based on one-sided PUT and GET
- Need to use symmetric memory locations
- Remote side of communication does not participate
- Can test for completion
- Barriers and collectives
- Popular on Cray and SGI hardware, also Blue Gene version
- To make sense needs hardware support for low-latency RDMA-type operations

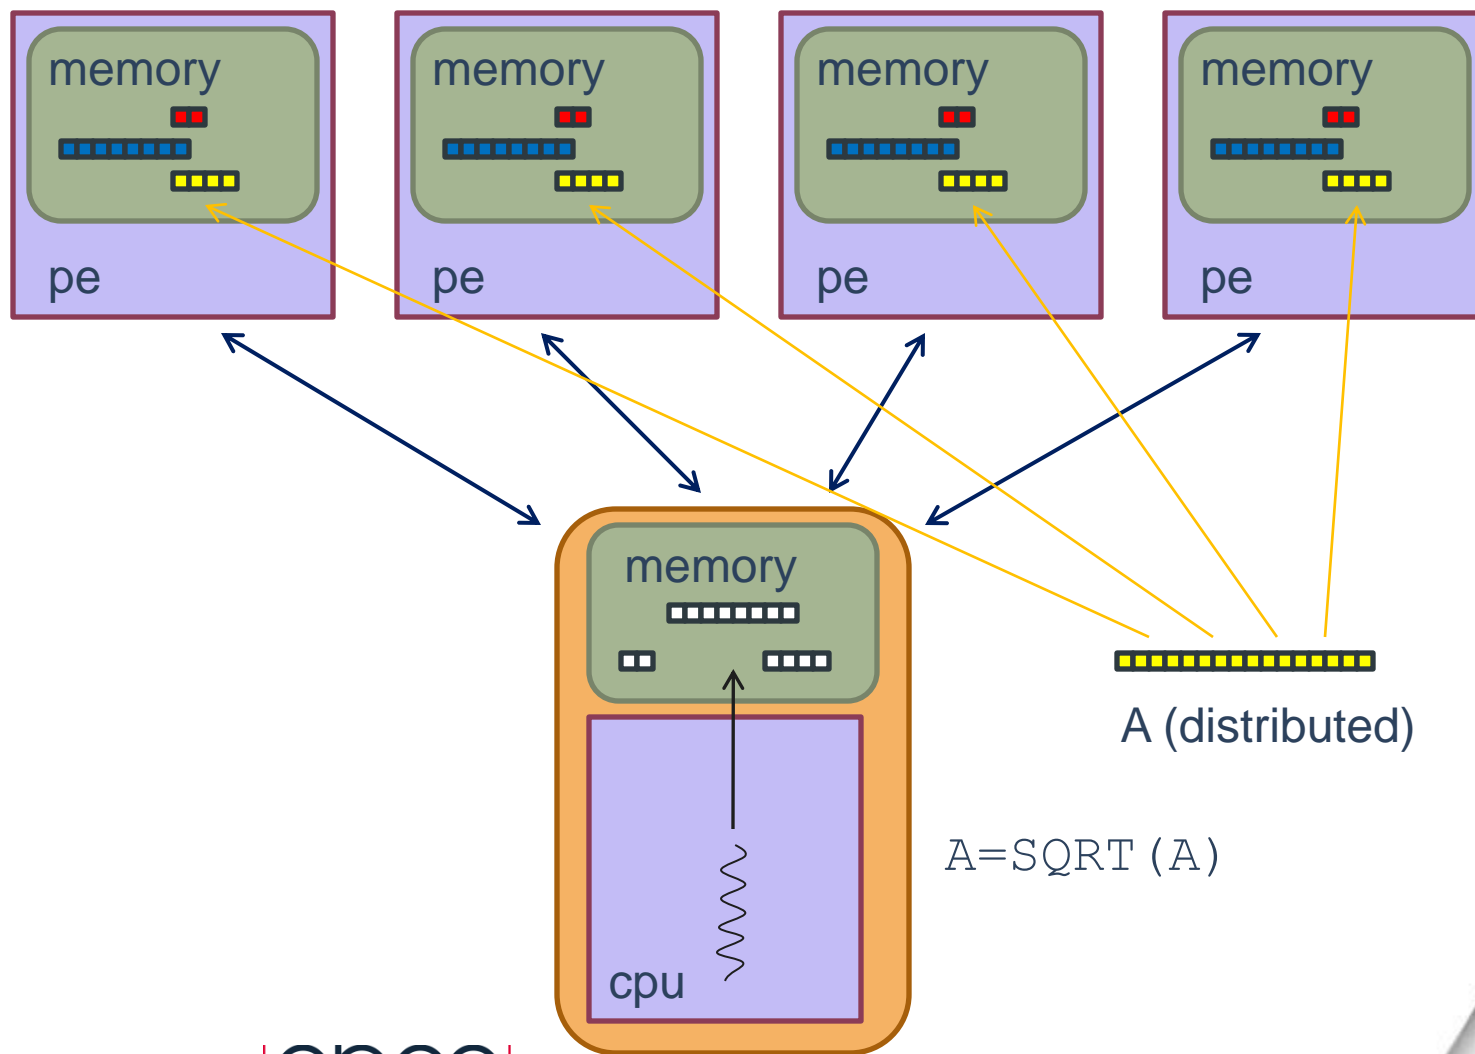
# High Performance Fortran (HPF)

- Data Parallel programming model
- Single thread of control
- Arrays can be distributed and operated on in parallel
- Loosely synchronous
- Parallelism mainly from Fortran 90 array syntax, FORALL and intrinsics.
- This model popular on SIMD hardware (AMT DAP, Connection Machines) but extended to clusters where control thread is replicated


# HPF




# HPF



# UPC



# UPC



```
upc_forall(i=0;i<32;i++;affinity) {  
    a[i]=a[i]*2;  
}
```

# UPC


- Extension to ISO C99
- Participating “*threads*”
- New *shared* data structures
  - shared pointers to distributed data (block or cyclic)
  - pointers to shared data local to a thread
  - Synchronization
- Language constructs to divide up work on shared data
  - `upc_forall()` to distribute iterations of `for()` loop
- Extensions for collectives
- Both commercial and open source compilers available
  - Cray, HP, IBM
  - Berkeley UPC (from LBL), GCC UPC

# Fortran 2008 coarray model


- Example of a Partitioned Global Address Space (PGAS) model
- Set of participating processes like MPI
- Participating processes have access to **local memory** via standard program mechanisms
- Access to **remote memory** is directly supported by the language



# Fortran coarray model



# Fortran coarray model



`a = b[3]`

# Fortran coarrays

- Remote access is a full feature of the language:
  - Type checking
  - Opportunity to optimize communication
- No penalty for local memory access
- Single-sided programming model more natural for some algorithms
  - and a good match for modern networks with RDMA

