



Programming GPUs using Directives

Alan Gray

EPCC

The University of Edinburgh

- Introduction to GPU Directives
- OpenACC
- New OpenMP 4 accelerator support

- Language extensions, e.g. *Cuda* or *OpenCL*, allow programmers to interface with the GPU
 - This gives control to the programmer, but is often tricky and time consuming, and results in complex/non-portable code
- An alternative approach is to allow the compiler to automatically accelerate code sections on the GPU (including decomposition, data transfer, etc).
- There must be a mechanism to provide the compiler with hints regarding which sections to be accelerated, parallelism, data usage, etc
- *Directives* provide this mechanism
 - Special syntax which is understood by accelerator compilers and ignored (treated as code comments) by non-accelerator compilers.
 - Same source code can be compiled for CPU/GPU combo or CPU only
 - c.f. OpenMP for multi-core programming.



- OpenACC standard was announced in Nov 2011
 - By CAPS, CRAY, NVIDIA and PGI
 - Standardisation based on several pre-existing models
 - Examples later

- We will now illustrate accelerator directives using OpenACC
- For definitive guide, full list of available directives, clauses, options etc see

<http://www.openacc-standard.org/>

- With directives inserted, the compiler will attempt to compile the key kernels for execution on the GPU, and will manage the necessary data transfer automatically.
- Directive format:
 - C: `#pragma acc`
 - Fortran: `!$acc`
- These are ignored by non-accelerator compilers

- The programmer specifies which regions of code should be offloaded to the accelerator with the `parallel` construct

C:

```
#pragma acc parallel  
{  
...code region...  
}
```

Fortran:

```
!$acc parallel  
...code region...  
!$acc end parallel
```

- Note: this directive is usually not sufficient on it's own – it needs (at least) to be combined with the `loop` directive (next slide)

- The `loop` construct is applied immediately before a loop (or nest of loops), specifying that it should be parallelised on the accelerator.

C:

```
#pragma acc loop  
for (...) {  
    ...loop body...  
}
```

Fortran:

```
!$acc loop  
do ...  
    ...loop body...  
end do
```


- The `loop` construct must be used inside a `parallel` construct.

C:

```
#pragma acc parallel
{
...
#pragma acc loop
  for(...) {
    ...loop body...
  }
...
}
```

Fortran:

```
!$acc parallel
...
!$acc loop ...
do ...
  ...loop body...
end do
!$acc end loop
...
!$acc end parallel
```

- Multiple `loop` constructs may be used within a single `parallel` construct.

- The `parallel loop` construct is shorthand for the combination of a `parallel` and (single) `loop` construct

C:

```
#pragma acc parallel loop
for (...) {
  ...loop body...
}
```

Fortran:

```
!$acc parallel loop
do ...
  ...loop body...
end do
!$acc end parallel loop
```

```
!$acc parallel loop
do i=1, N
    output(i)=2.*input(i)
end do
!$acc end parallel loop
```

- Compiler automatically offloads loop to GPU (using default values for the parallel decomposition), and performs necessary data transfers.
- Use of `parallel loop` may be sufficient, on its own, to get code running on the GPU, but further directives and clauses exist to give more control to programmer
 - to improve performance and enable more complex cases

Clauses can be added to `loop` (or `parallel loop`) directives

- `gang, vector`
 - Targets specific loop at specific level of hardware
 - `gang` ↔ CUDA block of threads
 - `vector` ↔ CUDA threads in block
 - You can specify both together
 - !\$acc loop gang vector schedules loop over all hardware

To be added to `parallel` (or `parallel loop`) directives

- `num_gangs, vector_length`
 - Tunes the amount of parallelism used
 - equivalent to setting the number of threads per block, number of blocks in CUDA
- Set the number of threads per block by specifying the `vector_length(NTHREADS) clause`
 - NTHREADS must be one of: 1, 64, 128 (default), 256, 512, 1024
- E.g. to specify 128 threads per block

```
#pragma acc parallel vector_length(128)
```

Further clauses to `loop` (or `parallel loop`) directives

`seq`: loop executed sequentially

`independent`: compiler hint

`if(logical)` Executes on GPU if `.TRUE.` at runtime, otherwise on CPU

`reduction`: as in OpenMP

`cache`: specified data held in software-managed data

`cache` e.g. explicit blocking to shared memory on NVIDIA GPUs

- Consider case with 2 loops.

```
!$acc parallel loop
do i=1, N
    output(i)=2.*input(i)
end do
!$acc end parallel loop

write(*,*) "finished 1st region"

!$acc parallel loop
do i=1, N
    output(i)=i*output(i)
end do
!$acc end parallel loop
```

- The `output` array will be unnecessarily copied from/to device between regions
 - Host/device data copies are very expensive

- Allows more efficient memory management

C:

```
#pragma acc data  
{  
...code region...  
}
```

Fortran:

```
!$acc data  
...code region...  
!$acc end data
```



```
!$acc data copyin(input) copyout(output)

!$acc parallel loop
do i=1, N
    output(i)=2.*input(i)
end do
!$acc end parallel loop

write(*,*) "finished first region"

!$acc parallel loop
do i=1, N
    output(i)=i*output(i)
end do
!$acc end parallel loop

!$acc end data
```

- the `output` array is no longer unnecessarily transferred between host and device between kernel calls

Applied to data, parallel [loop] regions

- `copy`, `copyin`, `copyout`
 - copy data "in" to GPU at start of region and/or "out" to CPU at end
 - `copy` means `copyin` and `copyout`
 - Supply list of arrays or array sections (using ":" notation)
 - N.B. Fortran uses start:end; C/C++ uses start:length
 - e.g. first N elements of array: Fortran `1:N`; C/C++ `0:N`
- `create`
 - Do not copy at all – useful for temporary arrays
 - Host copy still exists
- `private`, `firstprivate`
 - as per OpenMP
 - Scalars private by default
- `present`
 - Specify that data is already present on the GPU, so copying should be avoided (example later)

Sharing GPU data between subroutines

```
PROGRAM main
  INTEGER :: a(N)
  ...
  !$acc data copy(a)
  !$acc parallel loop
    DO i = 1,N
      a(i) = i
    ENDDO
  !$acc end parallel loop
  CALL double_array(a)
  !$acc end data
  ...
END PROGRAM main
```

```
SUBROUTINE double_array(b)
  INTEGER :: b(N)
  !$acc parallel loop present(b)
    DO i = 1,N
      b(i) = 2*b(i)
    ENDDO
  !$acc end parallel loop
END SUBROUTINE double_array
```

- The `present` data clause allows the programmer to specify that the data is already on the device, so should not be copied again

- See the documentation for full list of directives and clauses.
- **Runtime Library Routines** are available to, e.g.
 - Retrieve information about the GPU hardware environment
 - Specify which device to use
 - Explicitly initialize the accelerator (helps with benchmarking)
- **Environment Variables**
 - e.g. can be set to specify which device to use
- There are still a number of limitations with the model and current implementations
 - Meaning that the feasibility of use depends on the complexity of code



- All OpenACC partners plus Intel, AMD (and several other organisations including EPCC) formed a subcommittee of the OpenMP committee, looking at extending the OpenMP directive standard to support accelerators.
- More comprehensive and with wider remit than OpenACC
- Accelerator support now included in OpenMP 4
 - But implementations are not yet widely available

- Similar to, but not the same as, OpenACC directives.
- Support for more than just loops
- Less reliance on compiler to parallelise and map code to threads
- Advantage of being portable: not GPU specific
 - suitable for Xeon Phi or DSPs, for example
 - This has the disadvantage that the syntax is more complex than OpenACC when used for GPUs
- Fully integrated into the rest of OpenMP

- A directives based approach to programming GPUs is potentially much more simplistic and productive than direct use of language extensions
 - More portable
 - Higher level: compiler automates much of the work
 - Less flexible and possibly poorer performance
 - Limitations may offer challenges for complex codes
- The OpenACC standard emerged in 2011
 - We went through the main concepts with examples
- OpenMP 4 now incorporates accelerators
 - More comprehensive and with wider participation than OpenACC