



UPC

Introduction & Basics

Nick Johnson

EPCC

nick.johnson@ed.ac.uk

Objectives of the coming three lectures:

- understand the basic principles of UPC
- motivation behind PGAS
- learn about data distribution, synchronisation
- advanced features (dynamic memory allocation, collectives)

→ Practicals will try and emphasise the most important aspects of UPC

Unified Parallel C

Parallel extension to ISO C 99, adding

- explicit parallelism
- global shared address space
- synchronisation

Both commercial and open source compilers available

- Cray, IBM, SGI, HP
- GWU, LBNL, GCCUPC

UPC != PGAS

- PGAS is a programming model
- UPC is *one* implementation of this model

Many other implementations

- Language extension: Coarray Fortran
- New languages: Chapel, X10, Fortress, Titanium
- PGAS-like libraries: OpenSHMEM, Global Arrays

All implementations are different, but follow the same model!

UPC uses threads that operate independently in a SPMD fashion

→ threads execute the same UPC program

Identifiers that return information about the program environment:

THREADS: holds total number of threads

MYTHREAD: stores thread index

→ index runs from **0** to **THREADS-1**

```
#include <upc.h>
#include <stdio.h>

void main() {
    printf("Thread %d of %d says: Hello!", MYTHREAD, THREADS);
}
```

Concept of two memory spaces: **private** and **shared**

objects declared in **private** memory space are only accessible by a single thread

objects declared in **shared** memory space are accessible by all threads

➔ shared memory space is used to communicate information between threads

private variables declared as normal C variable

- multiple instances of variable will exist

```
int x; // private variable
```

shared variables declared with **shared** qualifier


- only allocated once, in shared memory space
- accessible by all threads

```
shared int y; // shared variable
```


If shared variable is scalar, space only allocated on thread 0

```
int x;
```

```
shared int y;
```



all threads can directly access shared data, even if it resides in a remote location

UPC creates logical partitioning of the shared memory space

- objects have *affinity* to one thread
- shared scalars always have affinity to thread 0

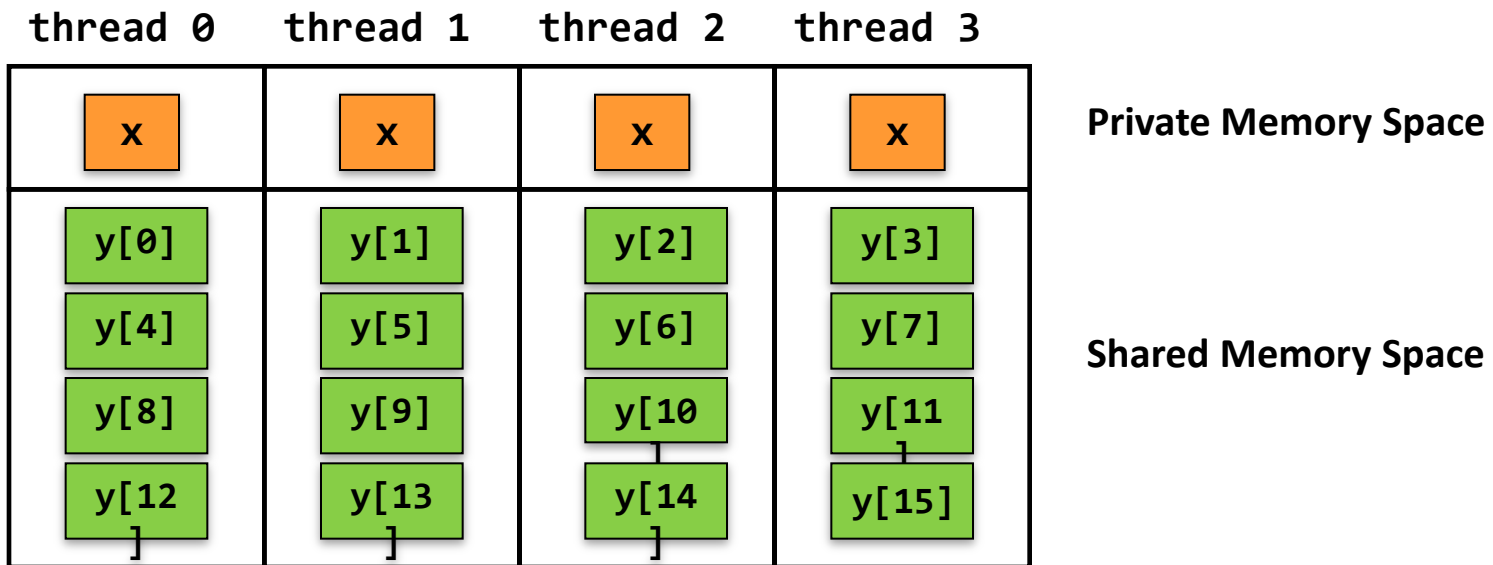
better performance if a thread access data to which it has affinity

- always keep data locality and affinity in mind

If a shared variable is an array, space allocated across shared memory space in a *cyclic* fashion by default

```
int x;
```

```
shared int y[16];
```



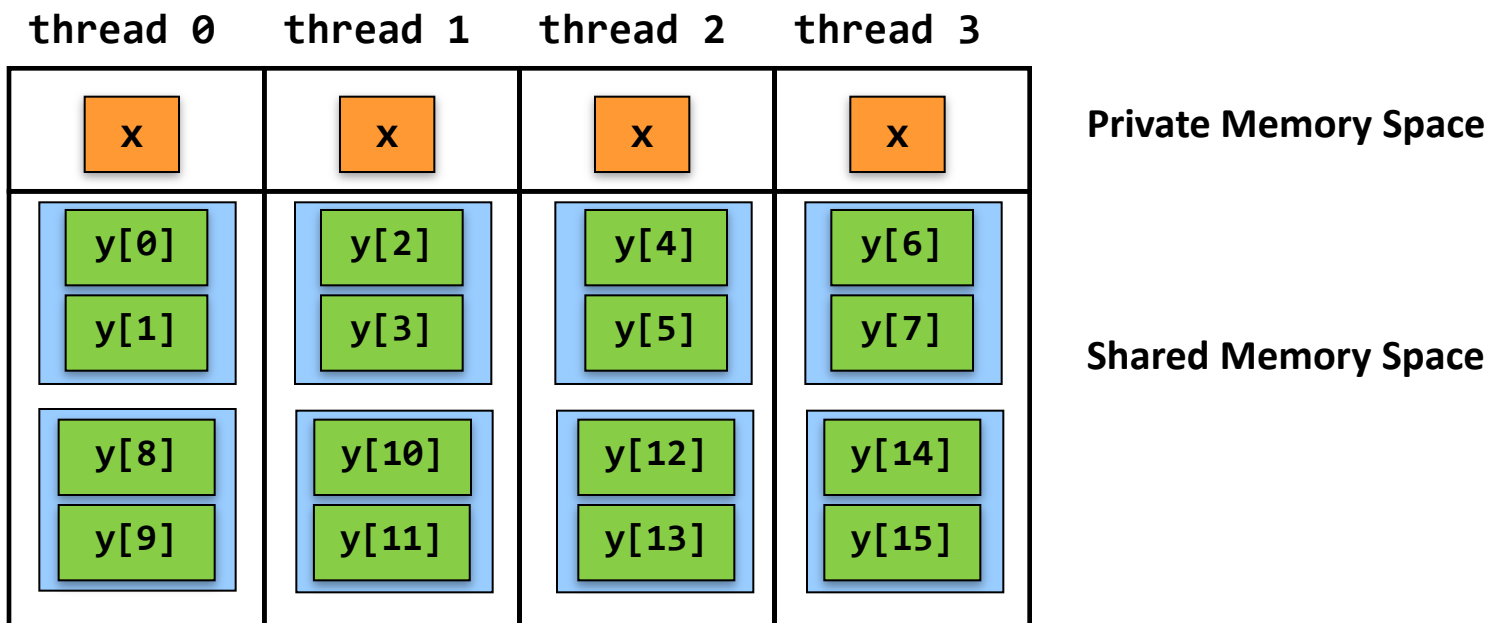
Shared array distribution (2)

Change data layout by adding a “blocking factor” to shared arrays

shared[**blocksize**] **type** array[n]

```
int x;
```

```
shared[2] int y[16];
```



Shared data means shared workload!

If shared data is distributed between threads, threads can distribute work on this data between them

UPC has built-in mechanism for explicitly distributing and sharing work

Statement for work distribution

- allows loop assignment of tasks to threads
- *parallel* for loop

4th parameter defines affinity to thread

- if “affinity % THREADS” matches MYTHREAD, execute iteration for that THREAD

```
upc_forall(expression; expression; expression; affinity)
```

Condition: iterations of `upc_forall` must be independent!

```
#define N 10 * THREADS
shared int vector1[N];
shared int vector2[N];
shared int sum[N];

void main() {
    int i;
    for(i=0; i<N; i++){
        sum[i] = vector1[i] + vector2[i];
    }
}
```

```
#define N 10 * THREADS
shared int vector1[N];
shared int vector2[N];
shared int sum[N];
```

```
void main() {
    int i;
    upc_forall(i=0; i<N; i++; i){
        sum[i] = vector1[i] + vector2[i];
    }
}
```

evaluated as $i \% \text{THREADS}$



Holding data in a shared memory space has implications

1) the lifetime of shared data needs to extend beyond the scope it was defined in (unless this is program scope)

→ storage duration

2) the shared data needs to be keep up-to-date

→ synchronisation

Shared objects *cannot* have **automatic storage duration**

- any variable defined inside a function!

Why?

SPMD model means a shared variable may be accessed outside lifetime of the function!

Conclusion

shared variables must either

- have file scope;
- or be declared as **static** if defined inside a function.

ensures shared objects are accessible throughout program execution

→ objects are not linked to the scope of a thread

→ objects will not simply “disappear” after a thread exists the scope in which the object was defined

Example: maximum of an array

```
#define max(a,b) (((a)>(b)) ? (a) : (b))
```

```
shared int maximum[THREADS];
```

```
shared int globalMax = 0; - - - - - →
```

Here: shared variables have file scope!

```
shared int a[THREADS*10];
```

```
void main(int argc, char **argv) {
```

```
    ... // initialise array a
```

```
    upc_forall(int i=0; i<THREADS*10; i++; i){
```

```
        maximum[MYTHREAD] = max(maximum[MYTHREAD], a[i]);
```

```
    }
```

```
    if (MYTHREAD == 0){
```

```
        for (int thread=0; thread<THREADS; thread++){
```

```
            globalMax = max(globalMax,maximum[thread]);
```

```
        }
```

```
    }
```

```
    ...  
}
```

This code will not work!!

Ensure all threads reach same point in execution

- necessary for memory and data consistency

Barriers used for synchronisation

- blocking
- split-phase (non-blocking)

upc_barrier

→ blocking

upc_notify, upc_wait

→ non-blocking

Example: maximum of an array

```
#define max(a,b) (((a)>(b)) ? (a) : (b))
```

```
shared int maximum[THREADS];
```

```
shared int globalMax = 0;
```

```
shared int a[THREADS*10];
```

Here: shared variables have file scope!

```
void main(int argc, char **argv) {
```

```
    ... // initialise array a
```

```
    upc_barrier;
```

```
    upc_forall(int i=0; i<THREADS*10; i++; i){
```

```
        maximum[MYTHREAD] = max(maximum[MYTHREAD], a[i]);
```

```
    }
```

```
    upc_barrier;
```

Ensure all threads found local maximum

```
    if (MYTHREAD == 0){
```

```
        for (int thread=0; thread<THREADS; thread++){
```

```
            globalMax = max(globalMax, maximum[thread]);
```

```
        }
```

```
    }
```

```
    upc_barrier;
```

Makes sure **globalMax** is found before being used

```
    ...  
}
```

- UPC Language Specification (Version 1.2):
http://upc.gwu.edu/docs/upc_specs_1.2.pdf
- UPC homepage:
<http://upc.gwu.edu/>
- GCCUPC compiler:
- <http://www.gccupc.org>