

UPC

Data distribution,
synchronisation & work sharing

Nick Johnson
EPCC

Nick.Johnson@ed.ac.uk

→ data distribution

- multi-dimensional data

→ synchronisation methods

- blocking versus non-blocking

→ work sharing

- examples: vector addition revisited, matrix-vector multiplication

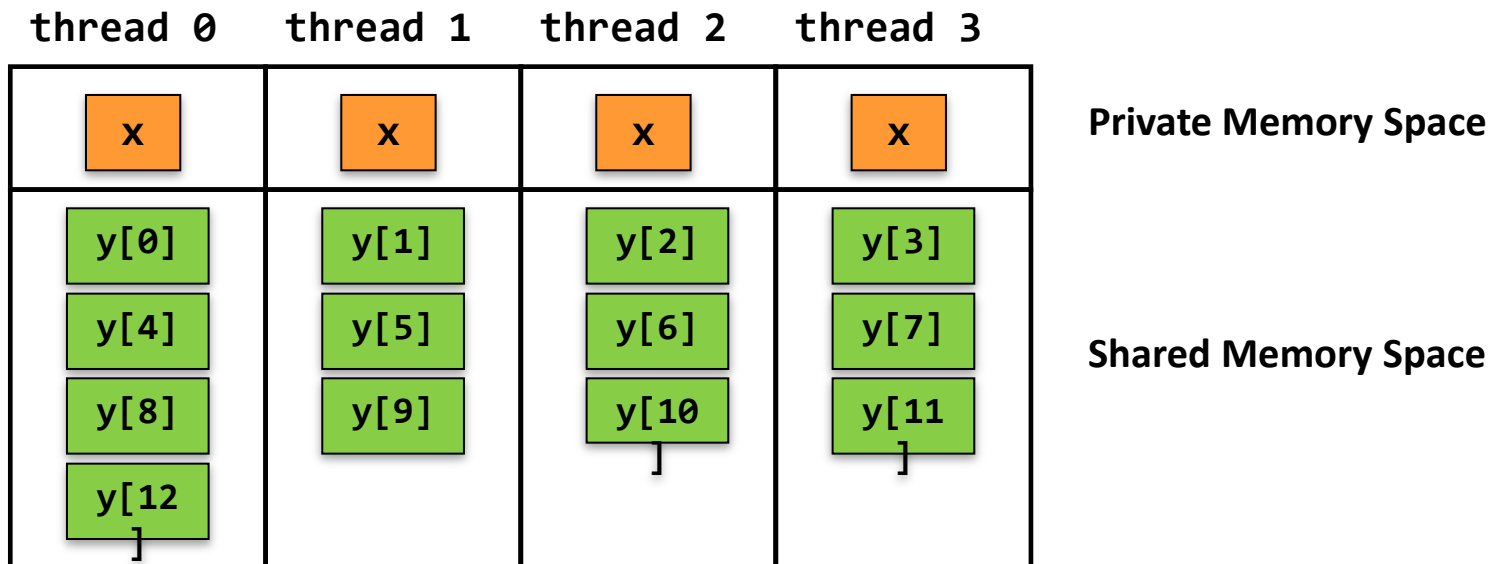
- ✓ private and shared data, logically partitioned memory space
- ✓ data objects have affinity to one thread exactly
- ✓ work sharing through `upc_forall`
- ✓ distribution of shared data
- ✓ storage duration of shared data
- ✓ synchronisation

focus of today's
lecture

Cyclic distribution is the default

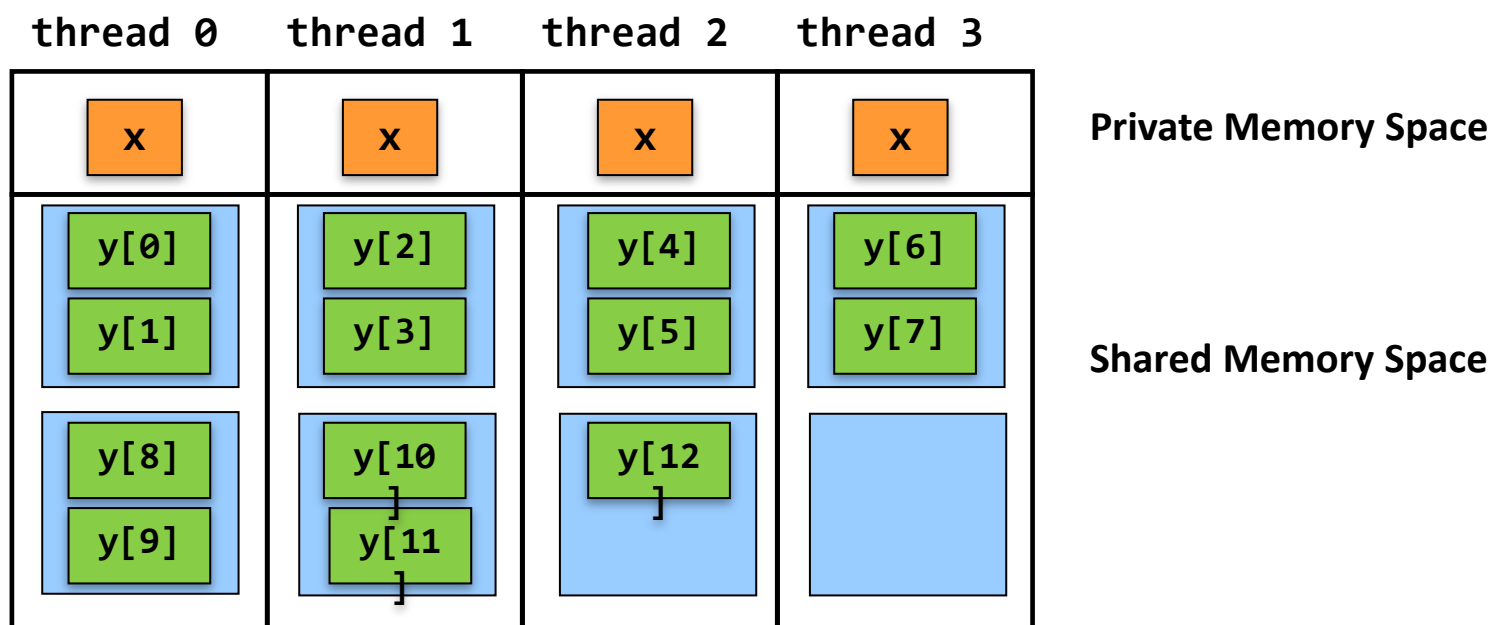
```
int x;
```

```
shared int y[13];
```



If number of elements is not an exact multiple of the thread count, threads can end up having uneven numbers of elements:

```
int x;  
shared[2] int y[13];
```



should be used if default distribution is not suitable

- more on the meaning of “suitable” later on...

→ *four* different cases

shared [4] → defines a block size of 4 elements

shared [0] → all elements are given affinity to thread 0

shared [*] → when possible, data is distributed in contiguous blocks

shared [] → equivalent to shared [0]

UPC can distribute data using ***block cyclic*** distributions

Distributions represent a ***top-down*** approach


- shared objects can be distributed into segments using the blocking factor

→ conceptually opposite of CAF, where shared objects are “created” by merging the pieces from every image in a bottom-up approach

distribution using the * layout qualifier and empty brackets


- Block distribution

```
shared[*] int y[8][8];
```




- Entire array on master

```
shared[] int y[8][8]; or  
shared[0] int y[8][8];
```




Distribution using different blocking factors

```
shared[8] int y[8][8];
```




```
shared[6] int y[8][8];
```




N.B. array layout convention is arr[y][x]!

shared double grid[8][8][8] with THREADS == 3




N.B. array layout convention is arr[x][y][z]!

shared double grid[8][8][8] with THREADS == 4




N.B. array layout convention is arr[x][y][z]!

shared [3] **double** grid[8][8][8] with THREADS == 3




blocking factor depending on dimensions → pencil distribution

shared [8] **double** grid[8][8][8] with **THREADS == 3**




combining thread count *and* blocking factor → slab distribution

shared [8] **double** grid[8][8][8] with THREADS == 4



slabs are contiguous in memory \rightarrow blocking factor product of 2 dimensions

shared [8*8] **double** grid[8][8][8] with **THREADS** == 4



Why is the distribution important?

it is all about performance and minimising the cost of reading and writing data...

accessing shared data which resides in a physically remote location is ***more expensive*** than accessing shared data which has affinity with the thread!

optimise layout of data by using knowledge of problem size and, if possible, the number of threads



number of UPC threads can be specified either at *compile time (static)* or at *runtime (dynamic)*

- GCCUPC Compiler: **-fupc-threads-N** specifies the number of threads at compile time as **N**

Advantages

- dynamic: program can be executed using any number of threads
- static: easier to distribute data based on THREADS

Disadvantages

- dynamic: not always possible to achieve best possible distribution
- static: program needs to be executed with number of threads specified at compile time

“An array declaration is illegal if THREADS is specified at runtime and the number of elements to allocate at each thread depends on THREADS.”

```
shared int x[4*THREADS];  
shared[] int x[8];
```

legal for static and dynamic environments

```
shared int x[8];  
shared[] int x[THREADS];  
shared int x[10+THREADS];
```

illegal for dynamic environment

static compilation can often give *better performance*, as data distribution is much easier to control

dynamic compilation provides *greater flexibility*, however optimal data distribution may not always be achievable

→ tradeoff between performance and convenience

needed to ensure all threads reach same point in execution flow


- memory and data consistency
- two types of synchronisation: blocking and non-blocking

blocking synchronisation makes all threads wait at a barrier until the last thread has reached that barrier before allowing execution to continue

non-blocking synchronisation allows for some local computations to be executed while waiting for other threads


`upc_barrier expopt`

1. all threads execute the code that requires synchronisation
2. once finished they wait at the barrier
3. when the last thread reaches the barrier, all threads are released to continue execution



`upc_notify expopt` and `upc_wait expopt`

1. thread finishes work that requires synchronisation → `upc_notify` to inform others
2. thread performs local computations → once finished, wait
3. when all threads have execute `upc_notify`, thread waiting at barrier can continue



the optional value **exp** can be used to check that all threads have reached the same barrier

if a thread executes a barrier with different **exp** tag than the other threads, the application reports an expression mismatch and aborts

→ very useful for making sure that all threads are on the intended execution path

4th parameter in `upc_forall` loop represents *affinity*

if MYTHREAD executes an iteration

affinity is an integer expression

→ `affinity % THREADS == MYTHREAD`

affinity is a pointer-to-shared

→ object pointed to has affinity to MYTHREAD

→ `upc_threadof(affinity)`

three vectors are distributed in cyclic fashion with the default blocking factor of 1
the modulo function identifies the local elements per thread

- if distribution changes, this code will *fail to identify the local elements*
- it will still produce the correct result!

```
#include <upc.h>
#define N 100*THREADS

shared int v1[N], v2[N], v1plusv2[N];

void main() {
    int i;

    for(i=0; i<N; i++)
        if (MYTHREAD == i%THREADS)
            v1plusv2[i] = v1[i] + v2[i];
}
```

alternative implementation:

iterate in steps of THREADS and eliminate the need for % operation

→ this implementation is again distribution specific

```
#include <upc.h>
#define N 100*THREADS

shared int v1[N], v2[N], v1plusv2[N];

void main() {
    int i;

    for(i=MYTHREAD; i<N; i+=THREADS)
        v1plusv2[i] = v1[i] + v2[i];
}
```

Advantage of affinity parameter: even if the distribution changes, `upc_forall` will behave correctly *and* identify local elements

```
#include <upc.h>
#define N 100*THREADS

shared int v1[N], v2[N], v1plusv2[N];

void main() {
    int i;

    upc_forall(i=0; i<N; i++; i)
        v1plusv2[i] = v1[i] + v2[i];
}
```

affinity parameter is
integer expression

matrix-vector multiplication


- perform as much computation on local data as possible
- work distribution is based on the elements of vector **c**

```
#include <upc.h>
shared int a[THREADS][THREADS];
shared int b[THREADS], c[THREADS];

void main (void)
{
    int i, j;

    upc_forall( i = 0 ; i < THREADS ; i++; &c[i]) {
        c[i] = 0;
        for ( j= 0 ; j < THREADS ; j++)
            c[i] += a[i][j]*b[j];
    }
}
```

affinity parameter is
pointer to shared array



number of remote operations:

$$C_0 = A_{0,0}B_0 + A_{0,1}B_1 + A_{0,2}B_2 \rightarrow 4$$

$$C_1 = A_{1,0}B_0 + A_{1,1}B_1 + A_{1,2}B_2 \rightarrow 4$$

$$C_2 = A_{2,0}B_0 + A_{2,1}B_1 + A_{2,2}B_2 \rightarrow 4$$


distribute matrix a in blocks of size **THREADS**

→ each row will be placed locally onto each thread

```
#include <upc.h>
shared [THREADS] int a[THREADS][THREADS];
shared int b[THREADS], c[THREADS];

void main (void)
{
    int i, j;

    upc_forall( i = 0 ; i < THREADS ; i++;&c[i]) {
        c[i] = 0;
        for ( j= 0 ; j < THREADS ; j++)
            c[i] += a[i][j]*b[j];
    }
}
```



number of remote operations:

$$C_0 = A_{0,0}B_0 + A_{0,1}B_1 + A_{0,2}B_2 \rightarrow 2$$

$$C_1 = A_{1,0}B_0 + A_{1,1}B_1 + A_{1,2}B_2 \rightarrow 2$$

$$C_2 = A_{2,0}B_0 + A_{2,1}B_1 + A_{2,2}B_2 \rightarrow 2$$

correct data distribution is important for performance

keep the number of remote reads/writes as low as possible

UPC gives programmers control over data layout and work sharing

→ it is important to be aware of performance implications

→ aim to keep work sharing loops independent of data distribution