



UPC

UPC Pointers

Dynamic Memory Allocation

UPC collectives

Nick Johnson

EPCC


Nick.Johnson@ed.ac.uk

- C and UPC pointers
 - dynamic memory allocation
 - locks
- UPC collectives

A *pointer* in C is a data type whose value points to another variable's memory address

```
float[2] array;
```

```
float* ptr = &array[0]; // Value of ptr is 213
```




change what object a pointer is referring to through pointer arithmetic:

- is type dependent
- incrementing a *float* pointer will move by **sizeof(float)**

```
float[2] array;
```

```
float* ptr = &array[0]; // Value of ptr is 213
```

```
ptr = ptr++; // Value of ptr is now 217
```



similar concept as in C

pointers are variables that contain addresses of other variables

UPC pointers can

reside in private or shared memory space

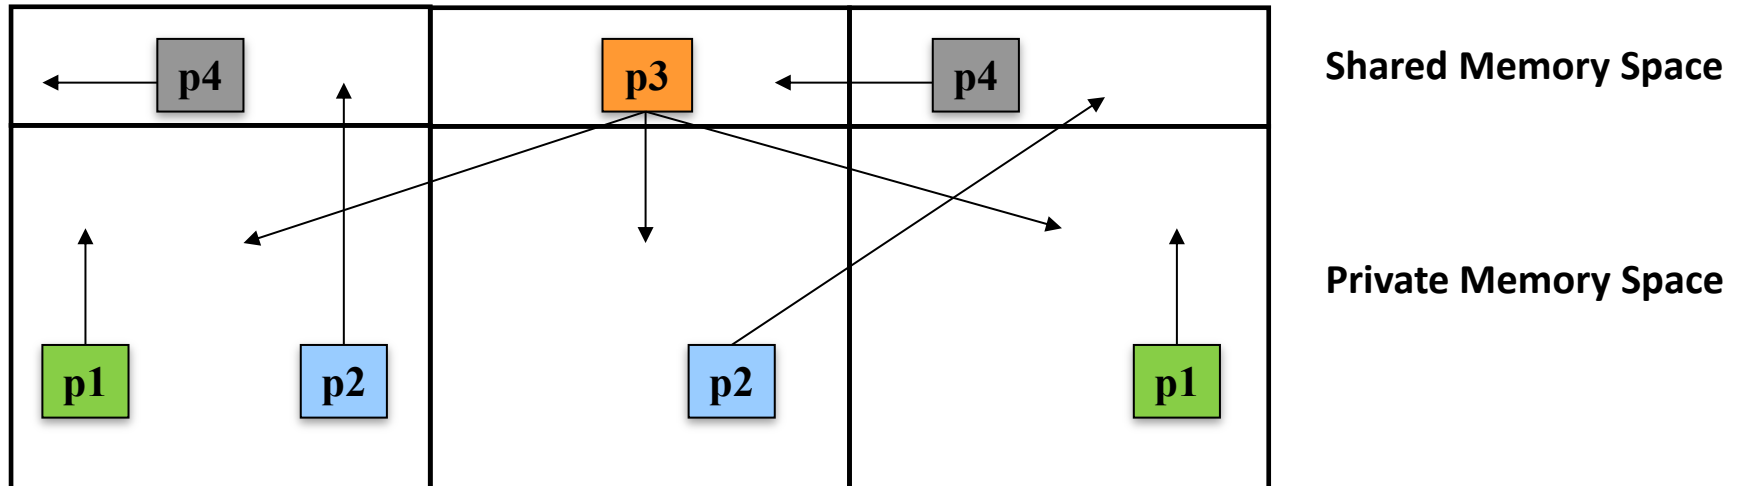
reference private or shared memory space

private to private → `int *p1;`

private to shared → `shared int *p2;`

shared to private → `int *shared p3; (not recommended)`

shared to shared → `shared int *shared p4;`



UPC pointers have three fields

- **thread** : the thread affinity of the pointer
- **address** : the virtual address of the block
- **phase** : indicates the element location within that block





the values of these fields are obtained from the functions

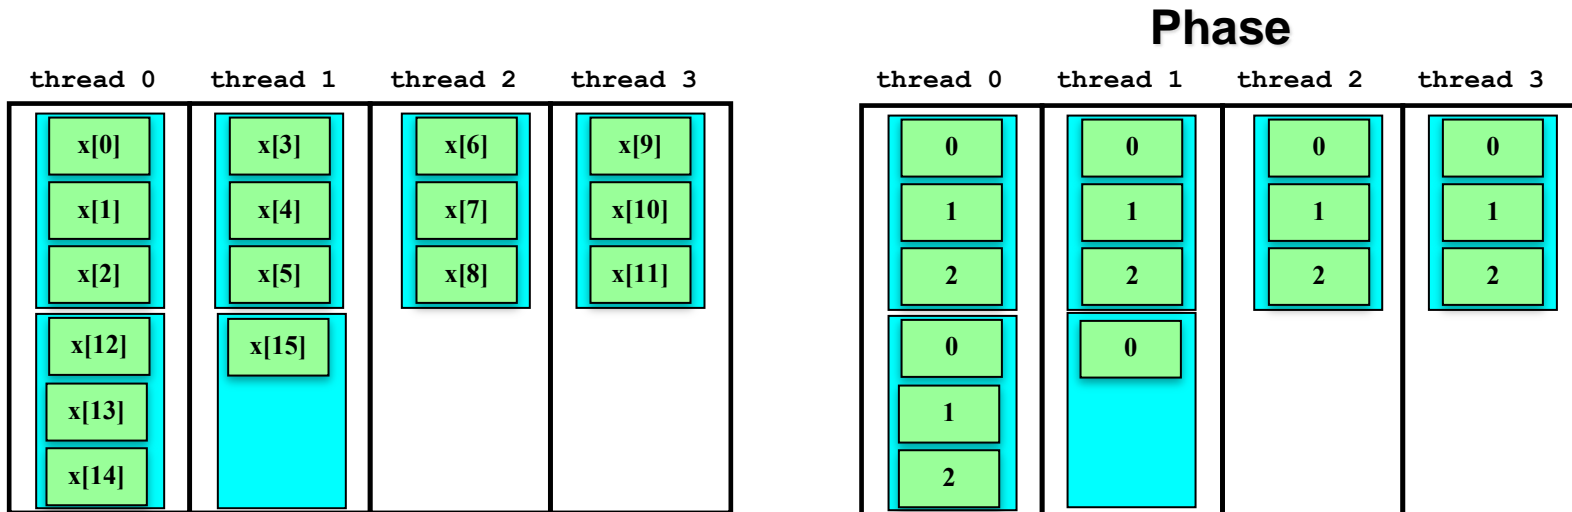
```
size_t upc_threadof (shared void *ptr)
size_t upc_phaseof (shared void *ptr)
size_t upc_addrfield (shared void *ptr)
```

shared [3] float x[16];

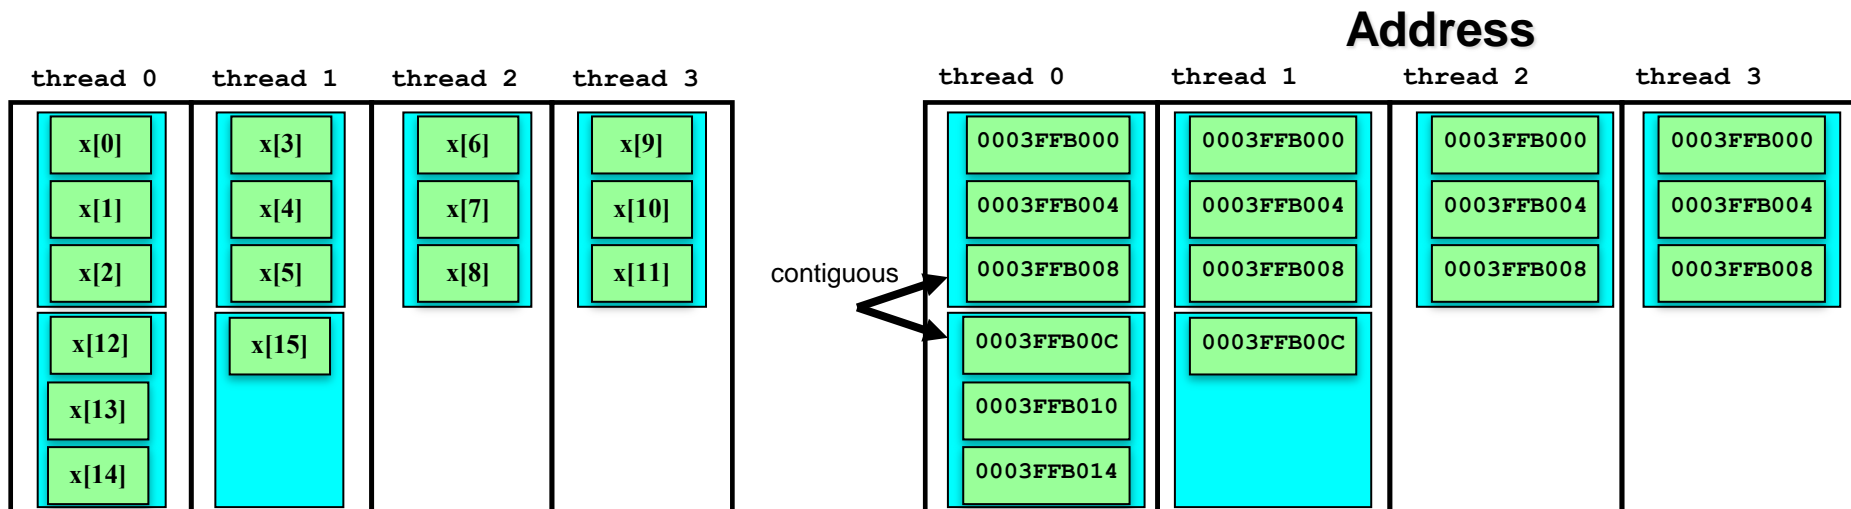
Thread



shared [3] **float** x[16];



```
shared [3] float x[16];
```




Pointer arithmetic takes into account the blocking factor

```
shared int x[16]; //shared int array
```

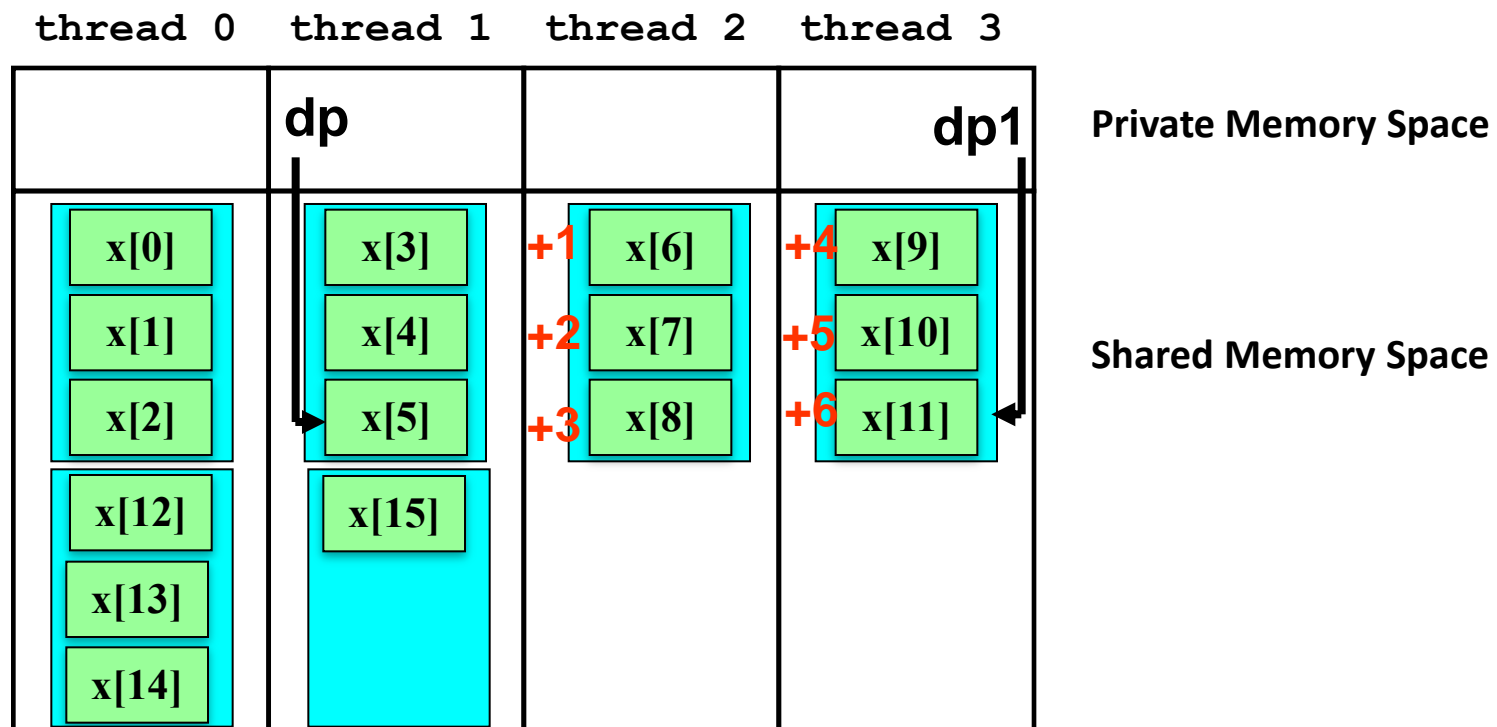
```
shared int *dp = &x[5], *dp1; // (private) pointers to shared int
```

```
dp1 = dp + 9; // default blocking factor 1
```



The pointer will follow its own blocking factor

```
shared [3] int x[16];  
shared [3] int *dp = &x[5], *dp1;  
dp1 = dp + 6; // blocking factor 3
```



- casting a shared pointer to a private pointer is allowed but not the other way around
- casting a shared pointer to private will result in loss of information
 - thread & phase
- casting is only *well defined* if the object pointed to by the shared pointer has local affinity

Casting a shared pointer to a private pointer results in information loss

```
shared [3] int x[16];  
shared int *dp = &x[5];  
int *ptr;  
ptr = (int *) dp;
```

→ `ptr != upc_addrfield(dp)`

	phase	thread	address
dp	2	1	0003FFB008
ptr	00AFF53008		

so far only seen how to allocate memory statically

dynamic memory allocation in UPC is of course possible

provides flexibility

allows object sizes to changes during runtime

→ this is one of the cases where pointers are required

dynamic memory allocation in *private space* is done using standard *C functions*

dynamic memory allocation in *shared space* is achieved using special *UPC functions*

two types of memory allocation functions

- non-collective: **upc_global_alloc**, **upc_alloc**
- collective: **upc_all_alloc**

collective calls are called by ***all*** threads and return the same address value (pointer) to all of them

non-collective calls can be executed by multiple threads. Each call will allocate a different shared block.

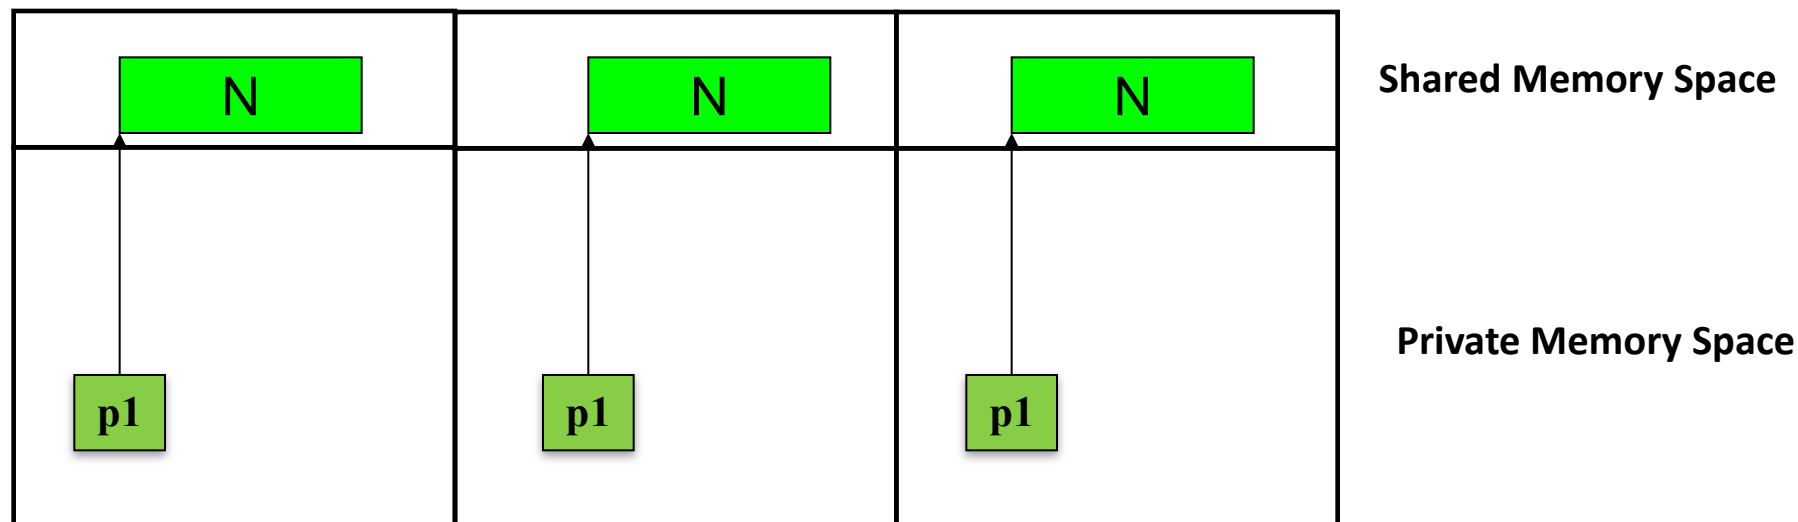
free the allocated memory using **upc_free**

- not a collective call

each thread allocates a memory block in its own shared memory space


```
shared [] int *ptr;
```

```
ptr = (shared [] int *) upc_alloc(N*sizeof(int));
```



```
shared [N] int *ptr;
```


```
ptr = (shared [N] int *) upc_global_alloc(THREADS, N*sizeof(int)) ;
```



allocate contiguous segments of shared memory

```
shared [N] int *ptr;
```

```
ptr = (shared [N] int *)upc_all_alloc(THREADS, N*sizeof(int));
```



Access control mechanism for critical sections

Sections which should be executed by one thread at a time

Serialised execution

UPC data type `upc_lock_t`

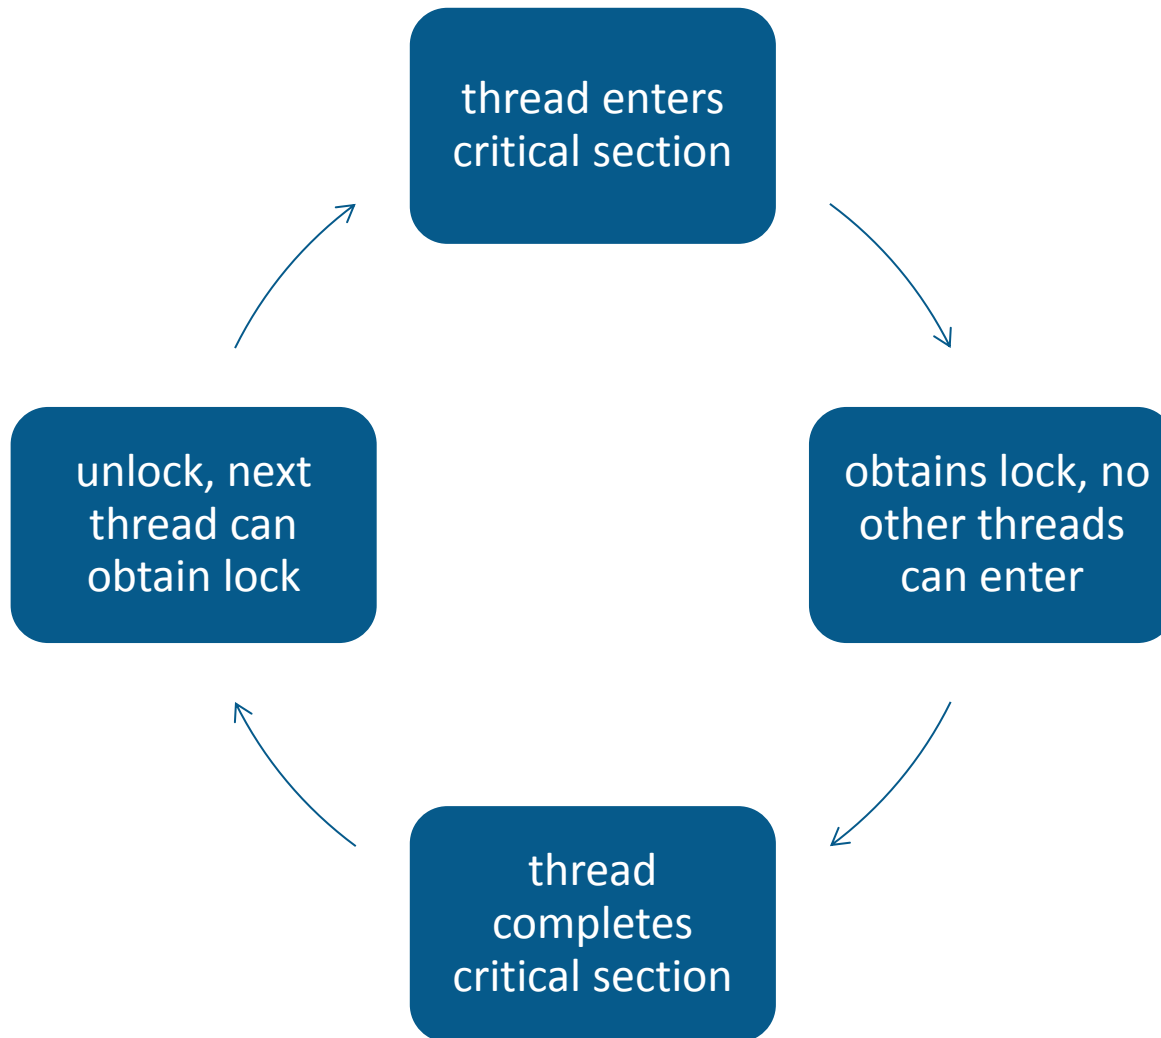
Can have one of two states: **locked** or **unlocked**

Can be seen by all threads → I mentally add **shared** in front of declarations

shared `upc_lock_t`

Technically an opaque type in C

Locks need to be manipulated through pointers



initial state of a new lock object is unlocked

locks can be created collectively

→ return value on every thread points to the same object

→ `upc_lock_t *upc_all_lock_alloc(void);`

or non-collectively

→ all threads that call the function obtain different locks


→ `upc_lock_t *upc_global_lock_alloc(void);`


resources allocated by locks need to be freed

→ `upc_lock_free (upc_lock_t *ptr);`

threads need to lock and unlock

- `upc_lock(upc_lock_t *ptr);` **// blocking**
- `upc_lock_attempt(upc_lock_t *ptr);` **// non-blocking**
- `upc_unlock(upc_lock_t *ptr);`





Model the forks as a shared array of locks, then allocate a lock (fork) dynamically and return a pointer to it:

```
upc_lock_t *shared fork[THREADS]  
fork[MYTHREAD]=upc_global_lock_alloc();
```

Now attempt to get the locks on either side:

```
left_fork=upc_lock_attempt(fork[MYHTREAD]);  
right_fork=upc_lock_attempt(fork[(MYHTREAD+1)%THREADS]);
```

If both forks are unlocked, lock them, eat, and release when finished.

```
upc_unlock(fork[MYTHREAD]);
```

```
upc_unlock(fork[(MYTHREAD+1)%THREADS]);
```

If only one fork is available, unlock (release) it and try again until two are available.

supported by most compilers

- readable code
- but not necessarily optimised for performance

requires separate header file

```
#include <upc_collective.h>
```

Two types of collective operations defined as part of the UPC standard specification:

① relocation collectives

`upc_all_broadcast`, `upc_all_scatter`, `upc_all_gather`,
`upc_all_gather_all`, `upc_all_exchange`, `upc_all_permute`

② computational collectives

`upc_all_reduceT`, `upc_all_prefix_reduceT`, `upc_all_sort`

Supported operations: `UPC_SUM`, `UPC_MULT`, `UPC_AND`, `UPC_OR`,
`UPC_XOR`, `UPC_LOGAND`, `UPC_LOGOR`, `UPC_MIN`, `UPC_MAX`

user specified functions supported via: `UPC_FUNC`, `UPC_NONCOMM_FUNC`

→ Calls to these functions must be performed by all threads

11 variations of the `upc_all_reduceT` and `upc_all_prefix_reduceT`


→ T needs to be replaced with the type used in the reduction operation

T	Type	T	Type
C	signed char	L	signed long
UC	unsigned char	UL	unsigned long
S	signed short	F	Float
US	unsigned short	D	double
I	signed int	LD	long double
UI	unsigned int		

```
shared[]  int A[2];
```

```
shared[2] int B[8];
```


```
upc_all_broadcast(B, A, 2*sizeof(int), UPC_IN_ALLSYNC | UPC_OUT_ALLSYNC);
```



```
shared[] int A[8];
```

```
shared[2] int B[8];
```

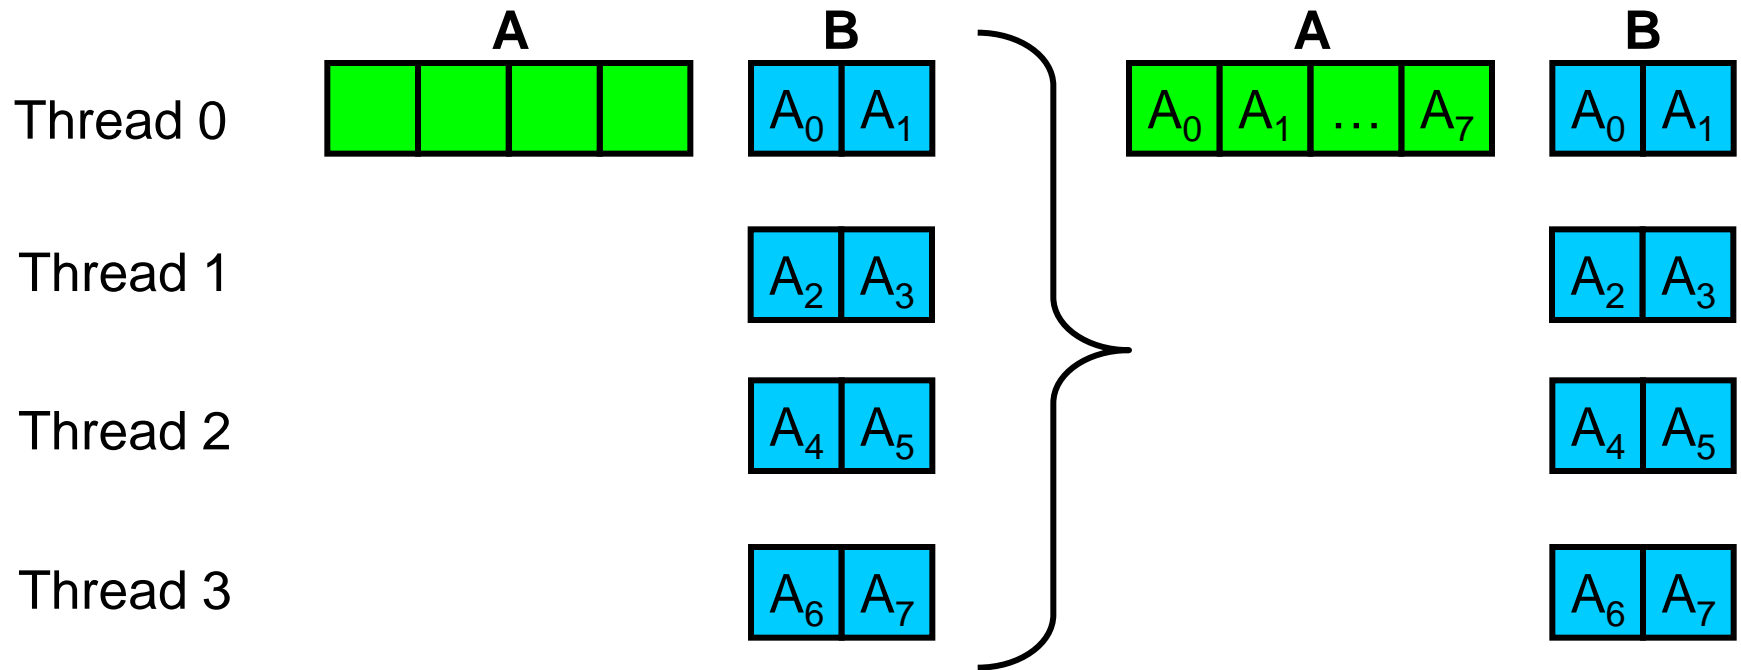
```
upc_all_scatter(B, A, 2*sizeof(int), UPC_IN_ALLSYNC | UPC_OUT_ALLSYNC);
```



```
shared[]  int A[8];
```

```
shared[2] int B[8];
```


```
upc_all_gather(A, B, 2*sizeof(int), UPC_IN_ALLSYNC | UPC_OUT_ALLSYNC);
```

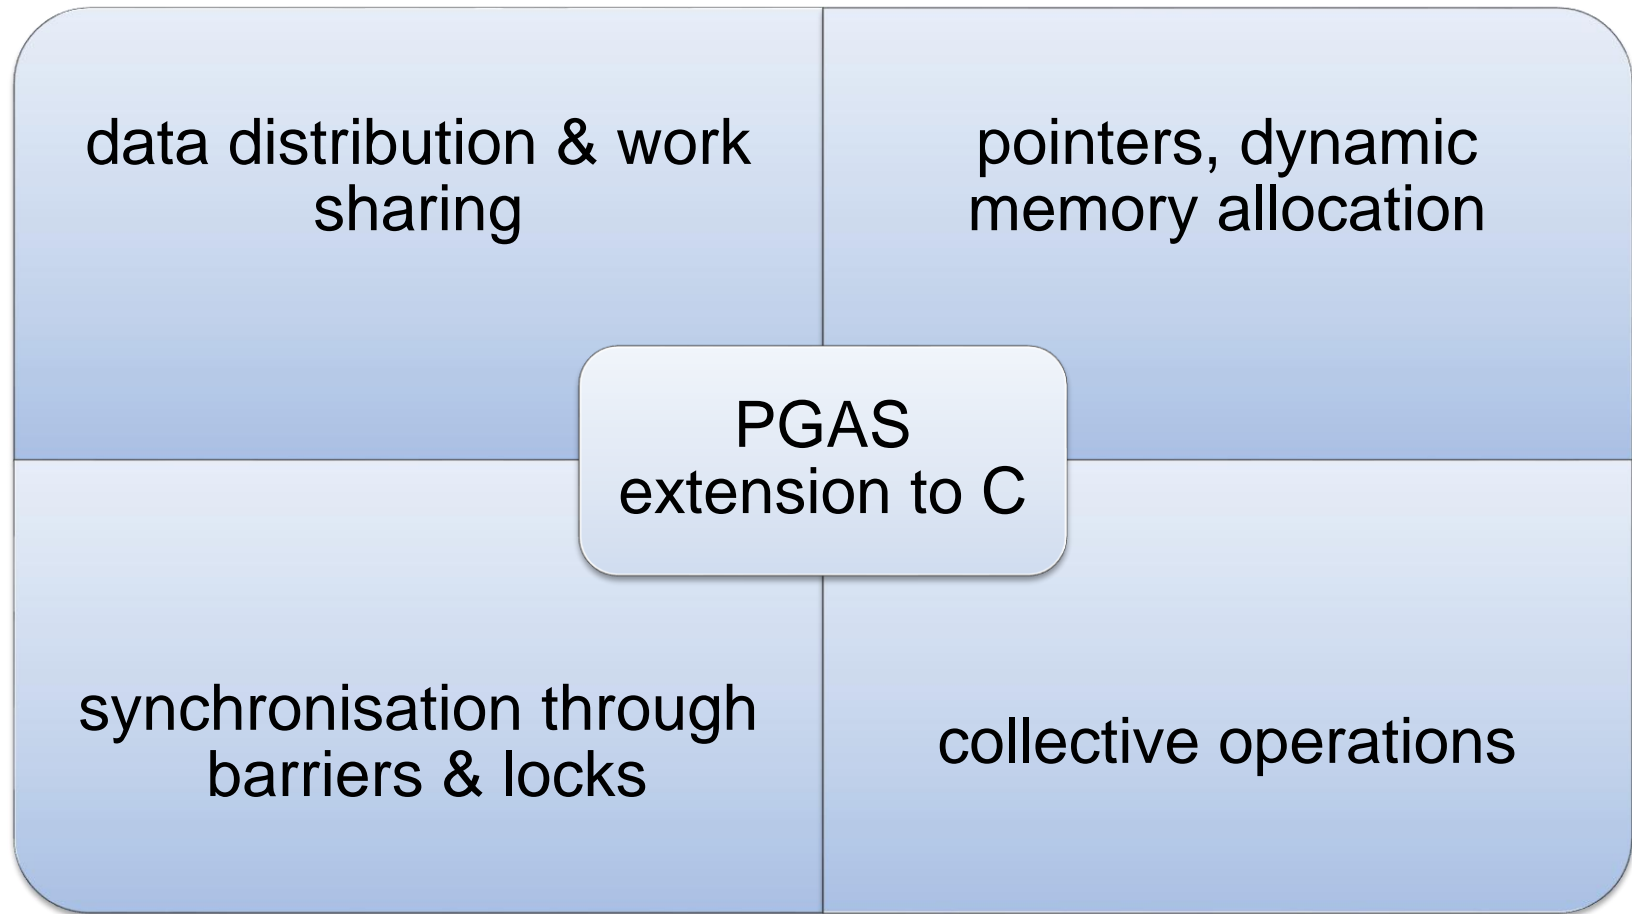


```
shared[4] double A[16];
```

```
shared double B;
```

```
upc_all_reduceD(&B, A, UPC_SUM, 4, 16, NULL, UPC_IN_ALLSYNC | UPC_OUT_ALLSYNC);
```





<http://upc.gwu.edu/documentation.html>

Language Specification (version 1.2)

UPC Manual

UPC Collective Operations Specification (version 1.0)