


Contents	epcc
• A Little Bit of History	
– Non-Parallel Programming Languages	
– Vector Processing	
– Data Parallel	
– Early Parallel Languages	
• Current Status of Parallel Programming	
– Parallelisation Strategies	
– Mainstream HPC	
• Alternative Parallel Programming Languages	
– Single-Sided Communication	
– PGAS	
– Accelerators	
– Hybrid Approaches	
• Final Remarks and Summary	

Parallel Programming Languages 2


Contents

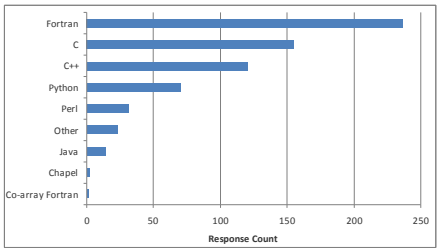


- A Little Bit of History
 - Non-Parallel Programming Languages
 - Vector Processing
 - Data Parallel
 - Early Parallel Languages
- Current Status of Parallel Programming
 - Parallelisation Strategies
 - Mainstream HPC
- Alternative Parallel Programming Languages
 - Single-Sided Communication
 - PGAS
 - Accelerators
 - Hybrid Approaches
- Final Remarks and Summary

Parallel Programming Languages 3

Non-Parallel Programming Languages



- Serial languages are also important for HPC
 - Used for much scientific computing
 - Basis for parallel languages
- PRACE Survey results:
 

Language	Response Count (approx.)
Fortran	230
C	150
C++	120
Python	70
Perl	30
Other	20
Java	10
Chapel	5
Co-array Fortran	2
- PRACE Survey indicates that nearly all applications are written in:
 - Fortran: well suited for scientific computing
 - C/C++: allows good access to hardware
- Supplemented by
 - Scripts using Python, PERL and BASH
 - PGAS languages starting to be used

Parallel Programming Languages 4

Vector Programming

epcc

- *Exploit hardware support for pipelining*
 - and for fast data access
- Early supercomputers were often vector systems
 - Such as the Cray-1
- Allowed operations on vectors
 - A vector is a series of values
 - e.g., a section of a Fortran array
- Typical vector loop

```
DO i = 1, n
    y(i) = a*x(i) + y(i)
END DO
```




Parallel Programming Languages

5

Vector Multiply

epcc

- Multiply process is made up of a number of stages
- Vector hardware allows stages to work independently and to pass results to each other in an “assembly line” manner
- Start-up cost as pipeline fills, but then a result every cycle



Parallel Programming Languages

6

Vectorisation



- Sometimes required restructuring of loops to allow efficient vectorisation
- Directives used to provide information to the compiler about whether a particular operation was vectorisable
- Compilers became increasingly good at spotting opportunities for vectorisation
- Vector supercomputers became less popular as parallel computing grew
- However, many modern CPUs contain vector-like features
 - e.g., INTEL IvyBridge processors in Cray XC30

Data Parallel



- *Processors perform similar operations across data elements in an array*
- Higher level programming paradigm, characterised by:
 - single-threaded control
 - global name space
 - loosely synchronous processes
 - parallelism implied by operations applied to data
 - compiler directives
- Data parallel languages: generally serial language (e.g., Fortran 90) plus
 - compiler directives (e.g., for data distribution)
 - first class language constructs to support parallelism
 - new intrinsics and library functions
- Paradigm well suited to a number of early (SIMD) parallel computers
 - Connection Machine, DAP, MasPar,...

Data Parallel II



- Many data parallel languages implemented:
 - Fortran-Plus, DAP Fortran, MP Fortran, CM Fortran, *LISP, C*, CRAFT, Fortran D, Vienna Fortran
- Languages expressed data parallel operations differently
- Machine-specific languages meant poor portability
- Needed a portable standard: High Performance Fortran
- Easy to port codes to, but performance could rarely match that from message passing codes
 - Struggled to gain broad popularity

Early Parallel Languages




- Connection Machine languages
 - Thinking Machines Corporation provided *data parallel* versions of a variety of sequential languages (*LISP, C*, CM-Fortran)
 - Allowed users to exploit a large number of simple processors in parallel
- OCCAM
 - Early message passing language
 - ...based on Communicating Sequential Processes
 - Developed by INMOS for programming Transputers
 - Explicitly parallel loops via PAR keyword
 - Language constructs for sending and receiving data through named channel
 - Could only communicate with neighbouring processors
 - → Message routing had to be in user software
- Most early languages for parallel computing were vendor-specific

Contents	epcc
<ul style="list-style-type: none"> • A Little Bit of History <ul style="list-style-type: none"> – Non-Parallel Programming Languages – Vector Processing – Data Parallel – Early Parallel Languages • Current Status of Parallel Programming <ul style="list-style-type: none"> – Parallelisation Strategies – Mainstream HPC • Alternative Parallel Programming Languages <ul style="list-style-type: none"> – Single-Sided Communication – PGAS – Accelerators – Hybrid Approaches • Final Remarks and Summary 	
Parallel Programming Languages	11

Parallelisation Strategies	epcc																						
<ul style="list-style-type: none"> • PRACE asked more than 400 European HPC users <ul style="list-style-type: none"> – “Which parallelisation implementations do you use?” 																							
<table border="1"> <thead> <tr> <th>Implementation</th> <th>Response Count (approx.)</th> </tr> </thead> <tbody> <tr> <td>MPI</td> <td>240</td> </tr> <tr> <td>OpenMP</td> <td>120</td> </tr> <tr> <td>Combined MPI+OpenMP</td> <td>90</td> </tr> <tr> <td>MPI, including MPI-2 single-sided</td> <td>50</td> </tr> <tr> <td>Posix threads</td> <td>20</td> </tr> <tr> <td>Other</td> <td>10</td> </tr> <tr> <td>Combined MPI+Posix threads</td> <td>10</td> </tr> <tr> <td>Combined MPI+SHMEM</td> <td>5</td> </tr> <tr> <td>SHMEM</td> <td>5</td> </tr> <tr> <td>HPF</td> <td>5</td> </tr> </tbody> </table>	Implementation	Response Count (approx.)	MPI	240	OpenMP	120	Combined MPI+OpenMP	90	MPI, including MPI-2 single-sided	50	Posix threads	20	Other	10	Combined MPI+Posix threads	10	Combined MPI+SHMEM	5	SHMEM	5	HPF	5	
Implementation	Response Count (approx.)																						
MPI	240																						
OpenMP	120																						
Combined MPI+OpenMP	90																						
MPI, including MPI-2 single-sided	50																						
Posix threads	20																						
Other	10																						
Combined MPI+Posix threads	10																						
Combined MPI+SHMEM	5																						
SHMEM	5																						
HPF	5																						
<ul style="list-style-type: none"> – Unsurprisingly, most popular answers were MPI, OpenMP and Combined MPI+OpenMP <ul style="list-style-type: none"> – Some users of Single-Sided communications 																							
Parallel Programming Languages	12																						

Parallelisation Strategies II

- PRACE also asked users of very largest systems:
 - “Which parallelisation method does your application use?”



- Most popular: “MPI Only” and “Combined MPI+OpenMP”
- 12% used single-sided routines

Mainstream HPC

- For the last 15+years, most HPC cycles on large systems have been used to run MPI programs, written in Fortran or C/C++
 - Plus OpenMP used on shared memory systems/nodes
- MSc in HPC includes compulsory courses in MPI and OpenMP
- However, there are now reasons why this may be changing:
 - Currently, HPC systems have increasingly large numbers of cores, but the individual core performance is relatively static
 - There are new challenges in exploiting future Exascale systems
- So, alongside mainstream HPC, there is also significant activity in:
 - Single-sided communication
 - PGAS languages
 - Accelerators
 - Hybrid approaches
- *Many of these areas are discussed later in this course*

Shared Memory

- *Multiple threads sharing global memory*
- Developed for systems with shared memory (MIMD-SM)
- Program loop iterations can be distributed to threads
 - Each thread can refer to private objects within a parallel context
- Implementation
 - Threads map to user threads running on one shared memory node
 - Extensions to distributed memory not so successful
- Posix Threads/PThreads is a portable standard for threading
- Vendors had various shared-memory directives
- OpenMP developed as common standard for HPC
 - OpenMP is a good model to use within a node
 - More recent task features


Parallel Programming Languages 15

Message Passing

- *Processes cooperate to solve problem by exchanging data*
- Can be used on most architectures
 - Especially suited for distributed memory systems (MIMD-DM)
- The message passing model is based on the notion of *processes*
 - *Process*: an instance of a running program, together with the program's data
- Each process has access only to its own data
 - i.e., all variables are private
- Processes communicate with each other by sending+receiving messages
 - Typically library calls from a conventional sequential language
- During the 1980s, there was an explosion in message passing languages and libraries
 - CS Tools, OCCAM, CHIMP (developed by EPCC), PVM, PARMACS, ...

Parallel Programming Languages 16


MPI: Message Passing Interface



- *De facto* standard developed by working group of around 60 vendors and researchers from 40 organisations in USA and Europe
 - Took two years
 - MPI-1 released in 1993
 - Built on experiences from previous message passing libraries
- MPI's prime goals are:
 - To provide source-code portability
 - To allow efficient implementation
- MPI-2 was released in 1996
 - New features: parallel I/O, dynamic process management and remote memory operations (single-sided communication)
- Now, MPI is used by nearly all message-passing programs

Parallel Programming Languages 17

Contents



- A Little Bit of History
 - Non-Parallel Programming Languages
 - Vector Processing
 - Data Parallel
 - Early Parallel Languages
- Current Status of Parallel Programming
 - Parallelisation Strategies
 - Mainstream HPC
- Alternative Parallel Programming Languages
 - Single-Sided Communication
 - PGAS
 - Accelerators
 - Hybrid Approaches
- Final Remarks and Summary

Parallel Programming Languages 18

Single-Sided Communication

- *Allows direct access to memory of other processors*
 - Each process can access total memory, even on distributed memory systems
- Simpler protocol can bring performance benefits
 - But requires thinking about synchronisation, remote addresses, caching...
- Key routines
 - PUT is a remote write
 - GET is a remote read
- Libraries give PGAS functionality
- Vendor-specific libraries
 - SHMEM (Cray/SGI), LAPI (IBM)
- Portable implementations
 - MPI-2, OpenSHMEM

Parallel Programming Languages 19

Single-Sided Communication

- Single-sided communication is major part of MPI-2 standard
 - Quite general and portable to most platforms
 - However, portability and robustness can have an impact on latency
 - Quite complicated and messy to use
- Better performance from lower-level interfaces, like SHMEM
 - Originally developed by Cray but a variety of similar implementations were developed on other platforms
 - Simple interface but hard to program correctly
- OpenSHMEM
 - New initiative to provide standard interface
 - See <http://www.openshmem.org>

Parallel Programming Languages 20

PGAS: Partitioned Global Address Space



- *Access to local memory via standard program mechanisms plus access to remote memory directly supported by language*
- The combination of access to all data plus also exploiting locality could give good performance and scaling
- Well suited to modern MIMD systems with multicore (shared memory) nodes
- Newly popular approach initially driven by US funding
 - Productive, Easy-to-use, Reliable Computing System (PERCS) project funded by DARPA's High Productivity Computing Systems (HPCS)

PGAS II



- Currently active and enthusiastic community
- Very wide variety of languages under the PGAS banner
 - See <http://www.pgas.org>
 - Including: CAF, UPC, Titanium, Fortress, X10, CAF 2.0, Chapel, Global Arrays, HPF?, ...
- Often, these languages have more differences than similarities...

PGAS Languages



- The broad range of PGAS languages makes it difficult to choose which to use
- Currently, CAF and UPC are probably most relevant as Cray's compilers and hardware now support CAF and UPC in quite an efficient manner
- CAF: Fortran with Coarrays
 - Minimal addition to Fortran to support parallelism
 - Incorporated in Fortran 2008 standard!
- UPC: Unified Parallel C
 - Adding parallel features to C


Parallel Programming Languages

23

Accelerators




- *Use accelerator hardware for faster node performance*
- Recently, most HPC systems are increasing the number of cores, but individual cores are not getting much faster
 - This gives significant scaling challenges
- Accelerators are increasingly interesting
 - ...for some applications
- PRACE Survey
 - “Could your application benefit from accelerators, such as GPGPUs?”
 - 61% thought so



Parallel Programming Languages

24



Programming GPUs |epcc|

- Graphics Processing Units (GPUs) have been increasing in performance much quicker than standard processor cores
- Led to an interest in GPGPU (General Purpose computation on Graphics Processing Units)
 - Where the GPU acts as an *accelerator* to the CPU
- Variety of different ways to program GPGPUs
 - CUDA
 - NVIDIA's proprietary interface to the architecture.
 - Extensions to C (and Fortran) language which allow interfacing to the hardware
 - OpenCL
 - Cross platform API
 - Similar in design to CUDA, but lower level and not so mature
 - OpenACC
 - Directives-based approach
 - OpenMP-style directives - abstract complexities away from programmer

Parallel Programming Languages 26

Hybrid Approaches



- *Use more than one parallelisation strategy within a single program*
- Trying to obtain more parallelisation by exploiting hierarchical parallelisation
- Most commonly, combining MPI + OpenMP
 - Using OpenMP across a shared-memory partition, with MPI to communicate between partitions
 - May make sense to use OpenMP just within a many-core processor
- But can also combine MPI with Pthreads or OpenSHMEM or CAF...
- Using many GPGPUs also often requires the use of MPI alongside the GPGPU programming approach

Contents



- A Little Bit of History
 - Non-Parallel Programming Languages
 - Vector Processing
 - Data Parallel
 - Early Parallel Languages
- Current Status of Parallel Programming
 - Parallelisation Strategies
 - Mainstream HPC
- Alternative Parallel Programming Languages
 - Single-Sided Communication
 - PGAS
 - Accelerators
 - Hybrid Approaches
- Final Remarks and Summary

Why do Languages Survive or Die?



- It is not always entirely clear why some languages and approaches thrive while others fade away...
- However, languages which survive do have a number of common characteristics
 - Appropriate model for current hardware
 - Good portability
 - Ease of use
 - Applicable to a broad range of problems
 - Strong engagement from both vendors and user communities
 - Efficient implementations available

Summary



- Development of portable standards have been essential for uptake of new parallel programming ideas
- Mainstream HPC is currently based on MPI and OpenMP
 - However, there are alternatives
- Exascale challenges have injected new life into development of novel parallel programming languages and approaches
- The remainder of this course focuses on PGAS languages and programming GPGPUs
 - Plus lectures on data parallel programming and single-sided communication

References



- PRACE-PP
 - **D6.1: Identification and Categorisation of Applications and Initial Benchmarks Suite**, *Alan Simpson, Mark Bull and Jon Hill, EPCC*
- PRACE-1IP
 - **D7.4.1: Applications and user requirements for Tier-0 systems**, *Mark Bull (EPCC), Xu Guo (EPCC), Ioannis Liabotis (GRNET)*
 - **D7.4.3: Tier-0 Applications and Systems Usage**, *Xu Guo, Mark Bull (EPCC)*
- ARCHER User Requirements
 - Project Working Group: *Katharine Bowes (EPSRC), Ian Reid (NAG), Simon McIntosh-Smith (Bristol), Bryan Lawrence (NCAS/Reading) and Alan Simpson (EPCC)*