

Data Parallel Programming with HPF

Dr Alan D Simpson
Technical Director, EPCC
a.simpson@epcc.ed.ac.uk
+44 131 650 5120

Contents



- Data Parallel Programming Languages
- Data Parallel Building Blocks
- Data Distribution in HPF
- An HPF Compiler

Data Parallel Paradigm



- Data Parallel programming involves the processing and manipulation of large arrays
- Processors should (simultaneously) perform similar operations on different array elements
- DP programming languages characterised by:
 - single-threaded control
 - global address space
 - loosely synchronous processes
 - parallelism implied by operations applied to data
 - compiler directives
- Transfer low level details from programmer to compiler
 - Encourages wider use of parallelism

Data Parallel Programming with HPF

3

DP: Pros and Cons



- Advantages: Ease of Use
 - simple to write (no explicit message passing),
 - debug and maintain (single thread of control) and
 - port: both old codes and new (standard languages + set of intrinsic functions)
 - global address space allows easy access to all data
- Disadvantages: Flexibility and Control
 - forced to write in DP style, only suitable for certain applications
 - restricted control: distribution of data and work
 - harder to get top performance
 - *reliant on good compilers*

Data Parallel Programming with HPF

4

Applications



- Requirements
 - data must be in large arrays
 - similar operations on each element
 - independent operations
 - problem should be naturally load-balanced
- Examples
 - lattice physics, e.g. QCD,
 - atmospheric modelling,
 - cellular automata,
 - image processing

Data Parallel Programming with HPF

5

Data Parallel Languages



- Many data parallel languages implemented:
 - Fortran-Plus, MP Fortran, CM Fortran, C*, CRAFT, Fortran D, Vienna Fortran
- Languages expressed data parallel operations differently
- Problems:
 - too many languages and dialects
 - languages were machine specific
- Result:
 - lack of portability (users tied to one manufacturer and prospective users unwilling to learn)
 - need portable standard ... HPF

Data Parallel Programming with HPF

6

HPF: High Performance Fortran



- Need for a Data Parallel standard well recognised
- High Performance Fortran Forum formed in 1991
- HPF Specification V1.0 published in 1993
- HPF Forum aimed to define a language that offered:
 - data parallel programming (single threaded, global name space, loosely synchronous parallel computation)
 - top performance on all machines (with code tuning)
- These goals are addressed in the form: Fortran plus
 - compiler directives: data alignment and distribution
 - new intrinsics, library functions and language constructs
- Subsidiary goals:
 - simplicity and portability (existing code and of new code)
 - ease of compiler implementation

Data Parallel Programming with HPF

7

DP Building Blocks



- Data distribution
- Array syntax
- Conditional operations
- Intrinsic functions
 - Shifts, reductions, scans, ...
- Libraries
- Many of these are already part of standard Fortran
 - ...and so are inherited by HPF


Data Parallel Programming with HPF

8

Data Distribution



- Data distribution is the major new feature in Data Parallel languages like HPF
- Data distributions controlled by language constructs and directives, e.g., HPF **DISTRIBUTE** directive



- *More on this topic later in the lecture*

Data Parallel Programming with HPF

9

Elemental Operations



- Whole arrays used in expressions

$$\begin{array}{|c|c|} \hline 2 & 12 \\ \hline 8 & 25 \\ \hline 18 & 42 \\ \hline \end{array} = \begin{array}{|c|c|} \hline 1 & 4 \\ \hline 2 & 5 \\ \hline 3 & 6 \\ \hline \end{array} * \begin{array}{|c|c|} \hline 2 & 3 \\ \hline 4 & 5 \\ \hline 6 & 7 \\ \hline \end{array}$$

c a b

- Mathematical functions on array-valued arguments
- Makes code more compact and readable
- Array operations in parallel

Data Parallel Programming with HPF

10

Fortran Array Syntax



- Arrays syntax allows arrays to become basic units

FORTTRAN 77

```
REAL a(100), b(100)
DO 1 i = 1,100
  a(i) = 2.0
  1 b(i) = b(i)*a(i)
```

HPF/Fortran 2008

```
REAL,DIMENSION(100):: a,b
a = 2.0
b = b * a
```


- This tends to make HPF code quite compact

Array Sections



- Features to allow some operations to act on only a subset of array elements
 - regular patterns
- Array sections and array constants can be defined through triplet notation in Fortran/HPF:

start index : end index : stride



- For example:

```
REAL, DIMENSION(1024) :: a
a(1:1024) = 1.0           ! same as a = 1.0
a(:1024:2) = 2.0         ! a(1)=a(3)=...= 2.0
a = (/ (i,i=1,1024) /)   ! a(1)=1,a(2)=2,...
a = (/ (i,i=1,4096,4) /) ! a(1)=1,a(2)=5,...
```

HPF/Fortran Array Sections

Diagram illustrating HPF/Fortran Array Sections on a 5x5 grid:

- $a(:,1)$: First column of the array.
- $a(2:4,2:4)$: 3x3 subarray starting at row 2, column 2.
- $a(1::2,1::2)$: Strided subarray starting at row 1, column 1, with a stride of 2 in both dimensions.

Axis 1 and Axis 2 are indicated by arrows pointing downwards and to the right respectively.

Data Parallel Programming with HPF 13

Conditional Operations

- Operations on subset of array, subject to some conditional mask
 - arbitrary patterns specified
- In Fortran and HPF, **WHERE** uses a logical mask
 - equivalent to **IF** on each element
 - there is also an **ELSE WHERE** clause

```

REAL, DIMENSION(7,7) :: a, b
...
WHERE (a /= 0) b = 1/a
...
WHERE (a > 2.0)
  b = a * 9.0
ELSE WHERE
  b = a / 3.0
END WHERE


```

Data Parallel Programming with HPF 14

Shift Operations

epcc

- Move whole array in particular direction
- For example, a cyclic shift one place to the right would produce the following result



- Also "end-off" shift; provide boundary conditions
- Many efficient parallel algorithms can be implemented in terms of shifts
- Useful for, e.g., image processing, cellular automata


Data Parallel Programming with HPF

15

CSHIFT in Fortran/HPF

epcc

- Regular movement of data



```
res = CSHIFT(src, SHIFT=-1, DIM=2)
```

which is roughly equivalent to

```
res(i,j) = src(i,j-1)
```

Data Parallel Programming with HPF

16

Reduction Operations



- Produce one result from many array elements, some global property of the array
- Provided to enable the programmer to, e.g.
 - find the largest/smallest value in an array
 - calculate the sum/product of all the elements of an array
 - logical AND, OR, EOR
 - count number of true elements in a logical array
- Very important in data parallel programming and often implemented efficiently

Data Parallel Programming with HPF

17

Reduction Operations



- Typically involve every element → global operation

```

REAL :: a(1024), b(4,1024)
scalar = SUM(a)           !sum of all elements
scalar = COUNT(a==0)      !number of zero elements
a = PRODUCT(b, DIM=1)    !product in first dim
scalar = MAXVAL(a, MASK=a.LT.0)

```

- Logical Reductions: ANY, ALL

```

LOGICAL a(n)
REAL, DIMENSION(n) :: b, c
IF (ALL(a))              ! Global AND
IF (ALL(b == c))         ! true if all elements equal
IF (ANY(a))              ! Global OR
IF (ANY(b < 0.0))        ! true if any element < 0.0

```

Data Parallel Programming with HPF

18

Other Operations



Intrinsic Functions and Libraries

- High level parallel constructs
- Useful for constructing parallel algorithms; saves reinventing the wheel.
- Efficient implementations exist: provided by manufacturer to make best use of host machine
- Examples
 - scan operations:
 - scan through arrays and produce some cumulative result
 - generalised communications:
 - user defined mapping and combination of elements to form a result

Data Parallel Summary



- Data parallel programming
 - single-threaded control
 - global address space
 - loosely synchronous processes
 - parallelism implied by operations applied to data
 - compiler directives
- Data parallel building blocks
 - data distribution
 - array syntax and conditional operations
 - intrinsic functions and libraries
- Fortran is an excellent basis for a data parallel programming language, such as HPF
 - widely used by computational scientists
 - contains many key data parallel building blocks

Data Distribution in HPF



- Data Parallel Programming involves the processing and manipulation of large arrays across many processors
- Generally, each processor has its own local memory
 - large arrays should be distributed across all processors
- Arrays distributed in different ways depending on how they are to be used
 - chose distribution which maximises the ratio of local work to communication
- Current compilers are not sophisticated enough to automatically distribute arrays over processors
- HPF has directives to control data mapping and help compiler
 - can reduce communication overhead and produce more efficient code
- Data distribution is the key new feature of HPF

Data Parallel Programming with HPF

21

DISTRIBUTE Directive



- Compiler directive to specify type of distribution to use
- Specifies the distribution for each dimension of array
- Formats:


```
!HPF$ DISTRIBUTE a(distribution)
```

or

```
!HPF$ DISTRIBUTE (distribution) :: a, b
```

 - **!HPF\$** is used for all HPF compiler directives --- this is a comment to non-HPF compilers
 - ***distribution*** is a comma-separated list of the distributions for each array dimension

Data Parallel Programming with HPF

22

BLOCK Distribution



```
REAL, DIMENSION(16) :: a
!HPF$ DISTRIBUTE (BLOCK) :: a
```

P1	P2	P3	P4
1	5	9	13
2	6	10	14
3	7	11	15
4	8	12	16



- *Example:* regular domain decomposition (Game of Life)

CYCLIC Distribution



```
REAL, DIMENSION(16) :: a
!HPF$ DISTRIBUTE (CYCLIC) :: a
```

P1	P2	P3	P4
1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16




- *Example:* ray tracing; Mandelbrot

Degenerate Distribution



```
REAL, DIMENSION(4, 8) :: a
!HPF$ DISTRIBUTE (*, CYCLIC) :: a
```



- First dimension can be called *degenerate*, *serial*, *collapsed*,...
- *Example*: serial dimension is an array of attributes associated with a point on the grid

Data Parallel Programming with HPF

25

General Distributions



- **CYCLIC (n)** : array split into contiguous chunks of size **n** and then distributed cyclically

```
REAL, DIMENSION(16) :: a
!HPF$ DISTRIBUTE (CYCLIC(2)) :: a
```



- **BLOCK (n)** : similar to **CYCLIC (n)** except only 1 chunk per (abstract) processor
 - for above example of 16 elements on 4 processors, $n \geq 16 / 4$
 - this means **BLOCK (4)**, **BLOCK (5)**, etc., would be OK, while **BLOCK (3)** would be *non-conforming*

Data Parallel Programming with HPF

26

Example 1: Distributed



```

PROGRAM disfunc
IMPLICIT NONE
INTEGER :: i
INTEGER, PARAMETER :: c1=5, c2=10
INTEGER, DIMENSION(10000) :: x, y

! Distribute data and hence work
!HPF$ DISTRIBUTE (CYCLIC) :: x, y
! CYCLIC for load balance

x = (/ (i, i = 1, 10000) /)

y(1:400) = x(1:400)*x(1:400)*x(1:400) + c1
y(401:1000) = x(401:1000)*x(401:1000) + c2
y(1001:10000) = x(1001:10000)
END

```

Data Parallel Programming with HPF

27

Data Distribution: Summary



- Data distribution is one of the major HPF extensions to Fortran
- **DISTRIBUTE** directive indicates type of distribution appropriate for each array dimension
- Data distribution helps compiler with work sharing
- Choosing the appropriate distribution can make a significant difference to performance

Data Parallel Programming with HPF

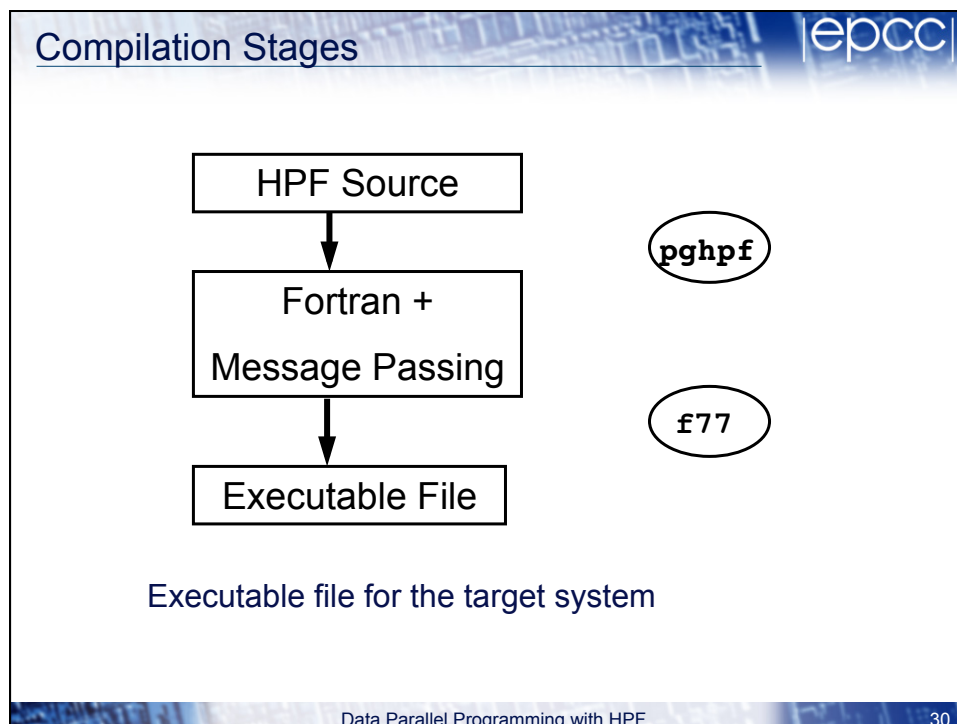
28

The Portland Group Compiler

pghp

- Parallelising translator
- Compiles to SPMD FORTRAN + message passing
- Works on a variety of systems:
 - CRAY, Fujitsu, HP, IBM, SGI, Sun, RISC Workstations, ...
- Supports various communication protocols: MPI, PVM, RPM (optimised PVM), SHMEM, ...
- Full HPF 1.1 functionality with some HPF 2.0 + approved extensions

Data Parallel Programming with HPF 29



Summary



- High Performance Fortran is a standardised data parallel Fortran
 - defined by HPF Forum (HPFF)
- HPF uses key features of Fortran
 - with extensions such as: data distribution, additional intrinsics and library functions,...
- The Portland Group compiler: **pghpf**
 - good functionality and portability

Data Parallel Programming with HPF

31

HPF in CP Lab



- Using the PGI HPF compiler, **pghpf**
- Compile command:


```
pghpf -Mautopar -lrt -o hello hello.f90
```
- Interactive run command:


```
./hello -pghpf -np 4
```
- *Note:* **pghpf** recognises **.f90** as free format Fortran
- Do **NOT** use **.hpf** as **pghpf** then assumes fixed format

Data Parallel Programming with HPF

32

Batch Jobs



- Running batch jobs via Sun Grid Engine:

```
qsub -cwd -pe mpi 4 ./hello.sge
```

 - where the script file **hello.sge** should contain the following line:

```
./hello -pghpf -np $NSLOTS
```
 - the environment variable **NSLOTS** is set equal to the number of processors requested by qsub
 - output will appear in a file called **hello.sge.oxxxx**
- There is a template file **hpfbatch.sge** which you can rename

```
cp hpfbatch.sge hello.sge
```