



## Lecture 3 - overview



- Outline
  - Array declarations and array I/O
  - Array sub-objects and assignments
  - Conditional operations – **where** and **forall** statements
  - Array intrinsic procedures
  - Static, assumed, automatic and dynamic arrays
  - Pure and elemental procedures
  - Advantages/ disadvantages of Fortran array syntax
  - Derived data types
  - Pointers

## Array declaration

- Recall, arrays are declared with **dimension** attribute


```
implicit none
integer, dimension(4) :: n4
integer, dimension(1:4) :: n4 ! same as previous line
```
- Provides 4 elements
  - Elements: `n4(1)`, `n4(2)`, `n4(3)`, `n4(4)`
  - First element is, by default, 1 \*\*\* not 0 \*\*\*
- Arrays can have more than one dimension

```
complex, dimension(1:10, 1:20) :: z
```
- Terminology
  - Number of dimensions is the *rank* (here 2)
  - Number of elements in given dimension is the *extent*
  - Sequence of the extents is the *shape* (here 10, 20)

## Array I/O

- Recall that arrays are stored in memory by columns
- Can write out the 3 (rows) by 2 (columns) array `a` using

```
real, dimension(3,2) :: a
write(*,*) a
```



- The output will be: 1, 3, 5, 2, 4, 6
- If we wish to write out the array in the order we typically use in matrix notation (maths) we need to use

```
do i = 1,3
  write(*,*) (a(i,j), j = 1,2) ! implied do loop
end do
```

- Output:  
1, 2  
3, 4  
5, 6

## Array sub-objects

- Consider the array

```
real, dimension(1:10) :: a
```


- Array in whole, or in part, can be addressed using a triplet:

- `a([lower] : [upper] [:stride])`


- Examples

- `a(7:)` ! `a(7), a(8), a(9), a(10)`
  - `a(2:10:2)` ! `a(2), a(4), a(6), a(8), a(10)`
  - `a(:)` ! whole array
  - `a(10:2:-2)` ! `a(10), a(8), a(6), a(4), a(2)`


## 2-d array sub-objects




`a(:, 1)`



`a(2:4, 2:4)`



`a(1::2, 1::2)`



## Elemental assignments

- Array assignments and arithmetic
  - `a(1:10) = 0.0` ! assign 0 to each element
  - `a(:) = a(:) + 1.0` ! add 1 to each element
  - `a = 0.0` ! array or scalar? depends  
! on declaration of a
- Can use a constructor (`/ ... /`)
  - `real, dimension(2) :: s = (/ -1.0, 1.0 /)`
  - `a(1:5) = (/ 1.0, -1.0, 1.0, -1.0, 1.0 /)`
- With an implied do loop...
  - `real, dimension(10000) :: &`  
`a = (/ (i, i = 1, 10000) /)`
  - sets `a(1) = 1.0, a(2) = 2.0, ...`

## Conformable arrays

- Arrays used in Fortran 2008 expressions must be *conformable*
  - They must have the same shape, that is,
  - The same rank (number of dimensions)
  - The same extent in every dimension
- For example


```
integer, dimension(10)    :: a, b
integer, dimension(10,2) :: c
a(:) = b(:)               ! correct
a(1:5) = b(6:10)          ! correct
a(1:3) = b(1:4)           ! clearly wrong
c(:, 1) = a(:)            ! correct
c(:, :) = b(:)            ! clearly wrong
                           ! different ranks
                           ! will not compile
```

## Conditional operations

- A conditional ( 'if' ) type construct for arrays

```
where (logical array-expression)
  array assignments
[else where] (logical array-expression)
  array assignments
end where
```
- Example: can create a mask for all elements of  $a > 0.0$ 
  - **where** uses a logical mask equivalent to applying an "if" on each element of the array
  - Also an **else where** ( ) clause

```
real, dimension(7,7) :: a,b
where(a > 0.0)
  a = 1.0/a    ! LHS and RHS must be
  b = a*b      ! the same shape
elsewhere
  b = 0.0
end where
```



## Forall

- Consider

```
do i = 1, n
  a(i, i) = x(i)
end do
```

  - performed serially (one iteration after the other)
- Instead can use Fortran 95 statement **forall**

```
forall (i = 1:n) a(i, i) = x(i)
```

  - performed conceptually 'at the same time' – in parallel
- Several distinct stages involved in the **forall**
  - Computation of valid index set (1:n above)
  - Define active index set (1:n above)
  - Evaluate right hand sides (x(1), x(2), ..., x(n) )
  - Assign to left hand sides (a(1,1), a(2,2), ..., a(n,n) )

## forall

- More complicated structures are allowed

```
forall (i = 1:n)
  where (a(i, :) == 0) a(i, :) = i
  b(i, :) = i / a(i, :)
end forall
```

- Useful for expressing parallelism in code – see e.g. HPF
- **Please note:** `forall` is nearly always less efficient than using array syntax – only use when array syntax is not appropriate

## Fortran intrinsic procedures

- Many intrinsic procedures are pure elemental
  - Elemental means they act on each element of array independently
  - Pure means they have no side-effects (more on this later)
  - Examples include: `sin`, `cos`, `min`, `max`, `exp`, `sqrt`, `log`...
- Consider

```
real, dimension(5) :: a
a = (/ 1.0, 4.0, 9.0, 16.0, 25.0 /)
a = sqrt(a)
```

- `a = sqrt(a)` finds the square root of all the elements of array `a`
- after application `a(1)=1.0`, `a(2)=2.0`, `a(3)=3.0`, `a(4)=4.0` and `a(5)=5.0`

- Array specific inquiry functions

**size, lbound, ubound, shape**

- Used to obtain information about the bounds of an array
- Functions **size**, **lbound** and **ubound** have an optional, **dim** argument, e.g. **size(array [,dim])**

- Consider the two dimensional array, **a**, declared with

**real, dimension(20,10) :: a**

- **size(a, dim = 1)** gives the size of the array along the first dimension, i.e. = 20 for this example; **size** returns integer
- **lbound(a, dim = 1)** gives the lower bound of the array along the first dimension, i.e. = 1; **lbound** returns integer or array
- **ubound(a, dim = 2)** gives the upper bound of the array along the second dimension, i.e. = 10
- **shape(a)** gives the shape of the array, i.e. 20 10

- Array transformation functions

- Array construction functions: **spread**, **pack**, **reshape**, ...
- Vector and matrix multiplication: **dot\_product**, **matmul**
- Reduction functions: **sum**, **product**, **any**, **maxval**, **minval**...
- Geometric location functions: **minloc** **maxloc**
- Array manipulation functions: **transpose**, **cshift**, ...

- Complete list – see Metcalf and Reid or the Standard documents



## Example: reshape


- The **reshape** array intrinsic function allows the shape of an array to be altered

– Syntax **reshape** (source, shape) e.g.

```
a(1:8) = (/ 1, 2, 3, 4, 5, 6, 7, 8 /)
```

```
b = reshape(a, (/ 4, 2 /) )
```

– Visually



## Reduction operations

- Typically involve every element, i.e, a global operation

```
integer, dimension(1024)    :: a
integer, dimension(4, 1024) :: b
integer :: ival
ival = sum(a)                ! Sum of all elements
ival = count(a == 0)         ! Number of zero elements
ival = product(b, dim = 1)   ! Product in first dim
ival = minval(a, mask = a < 0) ! Smallest negative
```


- Logical reductions: **any**, **all**

```
logical, dimension :: c(n)
if (all(c)) ...                ! Global logical and
if (all(b(1,:) == a(:))) ...   ! True if all elements equal
if (any(c)) ...                ! Global logical or
if (any(b < 0.0)) ...          ! True if any element < 0.0
```



## Shift Operations


- Move whole array in particular direction
- For example, a cyclic shift one place to the right would produce the following result



- Also "end-off" shift; provide boundary conditions
- Many efficient parallel algorithms can be implemented in terms of shifts
- Useful for, e.g., image processing, cellular automata

## Cshift

Intrinsic for regular movement of data



```
res1 = cshift(src, shift=-1, dim=2)
```

is roughly equivalent to

```
res1(i,j) = src(i,j-1)
```

## Pure/elemental procedures (F95 onwards)

- Procedures with no side-effects can be declared *pure*
  - No function dummy arguments are altered, i.e. all variables, intent(in)
  - No variables accessed by host association are altered
  - No I/O, stop

```
pure real function eqn_of_state(t, p)
```

- Elemental functions

- declared with scalar dummy arguments but can be called with array actual arguments
- must be pure

```
elemental real function eqn_of_state(t, p)  
  real, intent(in) :: t, p
```

```
a          = eqn_of_state(t0, p0)  
b(:)       = eqn_of_state(t1(:), p1(:))  
c(:, 2:4) = eqn_of_state(t2(:, 2:4), p2(:, 2:4))
```

## Static arrays

- So far have looked only at static arrays
  - Fixed size at declaration
  - Number of elements cannot be altered during program execution
  - Can be wasteful in terms of storage, often defined to be larger than required
  - In Fortran 2008 these are known as **explicit-shape** arrays
  - Declared using

```
real, dimension(1:10) :: a      ! 1D array  
real, dimension(1:10,1:5) :: b ! 2D array
```

## Assumed & automatic arrays

- Fortran allows array sizes to flexible
  - Size is determined on entering a sub-program
  - Can only be used within a sub-program
  - Allows flexibility and re-use of sub-program
  - The dimensions of these arrays are *known* at compile time
  - In Fortran 2008 known as **assumed-shape** or **automatic** arrays
  - Declared using

```
real, dimension(:) :: a      ! 1D array
real, dimension(:, :) :: b   ! 2D array
real, dimension(:, :, :) :: c ! 3D array
```

## Assumed & automatic arrays

```
subroutine reverse(aa)
  implicit none
  real, dimension(:), intent(inout) :: aa ! assumed
  real :: work(size(aa)) ! automatic
  integer :: i,imax

  imax = int(size(work))
  do i = 1,imax
    work(imax-i+1) = aa(i)
  end do
  aa = work      ! Copy back for output
end subroutine reverse
```

- Size of **aa** is *assumed* - determined by the calling program
- Size of **work** is *automatic* - as it depends on **aa**

## Array dummy arguments

- Ranks of actual and dummy arguments must agree
- Shape may be *assumed*

```
subroutine swap(a, b)
  real, dimension(:), intent(inout) :: a, b
  – Array inquiry functions such as size(a) and lbound(a),  
  ubound(a) make it easier to write generic routines
```

- Local arrays which may need to vary are *automatic*

```
subroutine swap(a, b)
  real, dimension(:), intent(inout) :: a, b
  real, dimension(size(a))           :: work
  work = a
  a = b
  b = work
end subroutine swap
  – Care with lower and upper bounds
```

## Dynamic arrays

- Dynamic arrays
  - Size can be allocated during execution
  - Very flexible but may slow run-time performance
  - Lack of array bounds checking during compilation
  - In Fortran 2008 these are known as **allocatable** arrays
  - Declared using

```
real, dimension(:), allocatable :: names ! 1D array
real, dimension(:,:), allocatable :: grid ! 2D array
  – Storage space for dynamic array is allocated using allocate

allocate(work(n, 2*n, 3*n))
  – Storage space is deallocated using deallocate

deallocate(work)
  – Beware, deallocated data is lost permanently!
```


## Example: dynamic arrays

```
module work_array
  implicit none
  integer :: n
  real, dimension(:,:,:), allocatable :: work
end module work_array

program main
  use work_array
  implicit none
  read(*,*) n
  allocate(work(n,2*n,3*n))
  . . .
  deallocate(work)
end program main
```

## Pros/Cons of array syntax

- Fortran 2008 has important features for scientific computing
  - Array sub-objects and assignment
  - Operations **where** and **forall**
  - Pure and elemental procedures
  - New intrinsic transformation functions
- These allow concise code to be written
- However, overblown array syntax can
  - Harm readability (e.g., obscure underlying algorithm)
  - Have undesirable impact on performance (possible to create very complex expressions the compiler can't unravel)
- Use with discretion
- Be aware of possible performance issues



## Derived data types



- Fortran 2008\* allows the use of derived data types
- In many algorithms data types can be grouped together to form an aggregate structure
- Often useful to be able to manipulate objects that are more sophisticated than the intrinsic types
- Allows linked data structures, lists, trees etc
- Derived data types are often needed in Coarray Fortran
- Imagine we wish to specify objects representing persons
  - Each person is uniquely distinguished by a name, age and ID number
- We can define a corresponding “person” data type as follows:

\* allowed since Fortran 90



## Derived data types

```
type person
  character (len=10) :: name
  real          :: age
  integer       :: id
end type person
```

- To declare a variable of type `person` use the syntax  
`type(person) :: you, me`
- Can assign values to the variable `you` using  
`you = person("Joe Bloggs", 21, 1234)`
- `you` is a variable containing 3 elements: `name`, `age` & `id`

## Derived data types

- The elements of derived type `person` may be accessed by using the variable name and the element name separated by the `%` character, e.g.  
`you%name` ! contains the name of you  
`you%age` ! contains the age of you  
`you%id` ! contains the id of you
- `%` is known as the *component selector*
- Can perform computations using derived type variables as follows:
  - Difference in age between variables `you` and `me`  
`real :: age_diff`  
`age_diff = you%age - me%age`



## Pointers



- Pointers were included in the Fortran 90 standard
  - Included in Fortran 2008 and useful for Coarray Fortran
- Similar but not identical to pointers used in C/C++/Java
- In other languages a pointer stores an *address* rather than the actual value
- In Fortran, a pointer is an *alias* for the variable it points to
  - Pointer variables do not store addresses, instead they evaluate to the variable that they point to
  - No need to dereference pointer variables in Fortran
- Fortran pointers can also point to a section of data structure, e.g. to a row or column of an array or part of a derived data type

## Pointers



- To declare a pointer variable, add the keyword, **pointer**, to the variable declaration, e.g.  
`real, pointer :: f`
- We also need to identify the variables to which **f** can point. Use the keyword, **target**, to do this, e.g.  
`real, target :: a`
- Assignment of pointer variables is achieved using the pointer assignment operator, **=>**, e.g.  
`f => a`
- The memory location associated with **a** is now also the memory location associated with **f**
  - If the value of **a** changes then the value of **f** also changes and vice versa as they now share the same memory location
- Can use **associated** to test whether a pointer points to something

## Pointers – simple example

```
program pointer_example
  implicit none
  real, target :: a=3.141, b=1.414, c=0.866 ! initial values
  real, pointer :: r, s ! pointers to real values

  ! Pointer associations result in a = r = 3.141 and b = s = 1.414
  r => a
  s => b

  ! Conventional assignments, variables a, b, r, and s will be 3.141
  s = r

  ! Can use pointer association to set r = c = 0.866 without changing
  ! a, b, or s.
  r => c

end program pointer_example
```

## Pointers and arrays

- Can also have pointers to arrays

```
real, dimension(8), target :: mydata
real, dimension(:), pointer :: index
```

- Pointer arrays just have a rank
  - They are declared without any bounds
  - Bounds are picked up from the target array
  - Target and pointer ranks must match

```
mydata = (/ (i, i=1,8) /)
index => mydata(1:8:2)

write(*,*) "Size of target array = ", size(mydata) ! 8
write(*,*) "Size of pointer array = ", size(index) ! 4
```

## Summary

- This session we have looked at
  - Array operations in Fortran 2008: assignments, I/O, sub-objects, conditionals, array intrinsic functions
  - Advantages/disadvantages of using Fortran array syntax
  - Advanced features of the language
    - Derived data types
    - Pointers