



Profiling, Analysis, and Performance

Adrian Jackson
EPCC
adrianj@epcc.ed.ac.uk
[@adrianjhc](https://twitter.com/adrianjhc)

- Analysis and browsing
 - Code browsing
 - Code analysis
- Performance profiling
 - Standard profilers
 - Timing by hand

- By “browsing and analysis” we mean
 - Reading and understanding code
 - Perhaps reformatting to make it more comprehensible
 - Analysing program structure
 - Essentially, understanding what it does and how to change it
- Understanding a 20 line program is quite easy
 - Reading it should probably be enough
- Understanding a 20,000 line program is almost impossible without tool help
 - Code browsers or IDEs are essential for this

- Understand source code syntax
 - Perhaps display code using “syntactic highlighting”
- Provide call-tree or class hierarchy views
 - Simple overview of program structure and control flow
- Declaration and reference lookups
 - Refactoring tools
- Trace variables through function/subroutine calls
 - Key data flow: when and how are data read? changed?
- Often offer customisable code formatting / reformatting
 - Very useful for later comprehensibility

- OO-Browser
 - <http://sourceforge.net/projects/oo-browser/>
 - Browser/analyser for C, C++, CLOS/Lisp, Eiffel, Java, Objective-C, Python, Smalltalk
 - Can work as part of Emacs under X Window or MS Windows
 - Free download, commercial support
- f90browse
 - Sun WorkShop (old Cray Xbrowse)
 - Browser/analyser for Fortran 90/77
 - Uses compiler information files (CIF) produced by f90 -db
- F90tohtml
 - Creates website based on source code
 - <http://code.google.com/p/f90tohtml/>
- Cleanscape Xlint
 - <http://www.cleanscape.net/>
 - Based on Xlint output from FortranLint product
 - Displays call tree, source code, Xlint output
- OpenGrok
 - <http://opengrok.github.io/OpenGrok/>
 - A fast search engine for programs
 - Read-only web interface for version control systems history log of a file
 - A stand alone GUI that can be used

- IDEs very useful
 - Eclipse, NetBeans, etc..
 - Usefulness depends on code and language use
 - Good refactoring and browsing tools
 - On-fly checking through compilation
- Your favourite debugger
- Debuggers don't have to be used for buggy code
- Many debugger features very useful
 - Breakpoint setting
 - Single stepping
 - Variable monitoring
 - Run-to-mark
- GUI-driven debuggers best, but command line can be useful too

- Less than a full browser
 - Typically command-line not GUI driven
 - Kind of commands to call from your Makefile
 - Possibly build into testing structure
- Offer “second opinions” on code structure
- Can offer more in-depth analysis than browsing tools

```
#include<stdio.h>

int main(){
    int i, a;

    for(i=0; i<100; i++){
        a += i;
    }

    printf("a: ", a);

    return 0;
}
```

- lint
 - Checks C program “correctness”
 - Identifies semantic errors, poor programming practices, portability and robustness issues
 - Splint is new open source version: <http://www.splint.org/download.html>

```
(8) warning: variable may be used before set: a
```

```
function returns value which is always ignored  
printf
```

```
too many arguments for format  
printf                               assemloop.c(11)
```

- flint – cleanscape has a version
 - Same for Fortran
- ftnchek - <http://www.dsm.fordham.edu/~ftnchek/> - F77
 - Similar “correctness” checker
 - Identifies semantic errors
- cscope – open source
 - Interactive browsing/analysis tool for C
- cppcheck, splint, cscout
- lint4j, jlint, findbugs, any IDE you are using

- cflow – gnu <http://www.gnu.org/software/cflow/>
 - Generates function call graph with code line numbers
 - Works on C source or (non-stripped) object files

```
1  main: int(),  
    <assemloop.c 3>  
2    printf: <>
```

- cxref – open source <http://www.gedanken.org.uk/software/cxref/>
 - “Big brother” of cflow
 - Includes external function calls
 - Also includes variables etc.

NAME	FILE	FUNCTION	LINE				
__func__	modassemloop.	main	1*				
a	modassemloop.	main	3*	6=			
i	modassemloop.	main	2*	5=	5	6	7=
main	modassemloop.	---	1*				

Code analysis tools cont.

assembloup.c:

NAME	FILE	FUNCTION	LINE			
BUFSIZ	stdio_iso.h	---	105*			
EOF	stdio_iso.h	---	133*			
FILE	stdio_iso.h	---	75-			
FILENAME_MAX	stdio_iso.h	---	137*			
FOPEN_MAX	stdio_iso.h	---	136*			
L_ctermid	stdio.h	---	126*			
L_cuserid	stdio.h	---	127*			
L_tmpnam	stdio_iso.h	---	144*			
NULL	stdio_iso.h	---	101*			
P_tmpdir	stdio.h	---	133*			
SEEK_CUR	stdio_iso.h	---	140*			
SEEK_END	stdio_iso.h	---	141*			
SEEK_SET	stdio_iso.h	---	139*			
TMP_MAX	stdio_iso.h	---	142*			
_ALIGNMENT_REQU	isa_defs.h	---	308*			
_BIG_ENDIAN	isa_defs.h	---	297*			
_BIT_FIELDS_HTO	isa_defs.h	---	300*			
_CHAR_ALIGNMENT	isa_defs.h	---	303*			
_CHAR_IS_SIGNED	isa_defs.h	---	302*			
_CONSOLE_OUTPUT	isa_defs.h	---	316*			
_DMA_USES_VIRTA	isa_defs.h	---	314*			
_DOUBLE_ALIGNME	isa_defs.h	---	307*			
_FILEDEFED	stdio_iso.h	---	74*			
_FILE_OFFSET_BI	feature_tests	---	96*	98	98	
<hr/>						
	stdio_iso.h	---	52	87		
	stdio.h	---	90	141	297	

- Compiler checking
 - -Wall for gnu compilers
 - -fcheck=all – gfortran
 - -pedantic

```
adrianj@gateway:~$ gcc -Wall -pedantic -o main main.c
```

```
main.c: In function âmainâ:
```

```
main.c:10:3: warning: too many arguments for format [-Wformat-extra-args]
```

- Use different compilers
 - Cray, Nag, etc...
 - Try something other than gnu if possible
 - Different compilers have different checking functionality etc...

```
adrianj@hector-xe6-10:~> craycc -o test test.c
```

```
    a += i;
```

```
CC-7212 craycc: WARNING File = test.c, Line = 7
```

```
Variable "a" is used before it is defined.
```

```
Total warnings detected in test.c: 1
```

- Once you've a working code you may want to optimise it
 - *"We should forget about small efficiencies, say about 97% of the time: premature optimization is the root of all evil. Yet we should not pass up our opportunities in that critical 3%. A good programmer will not be lulled into complacency by such reasoning, he will be wise to look carefully at the critical code; but only after that code has been identified"* — Donald Knuth
 - *"The First Rule of Program Optimization: Don't do it. The Second Rule of Program Optimization (for experts only!): Don't do it yet."* — Michael A. Jackson
- First step is Profiling
 - Discover if it's efficient or quick
- Highlight areas that are computational expensive
 - Often not where you'd expect
- Focus work on areas where impact will be biggest
 - Need hard data to make these choices
- Profiling tells you *where* the code is spending all of its time
 - Optimisation without profiling is leaping before looking!
 - Avoid working on speeding up part of an algorithm if the application only spends small part of its time there

- Code profiling is the first step for anyone interested in performance optimisation
- Profiling works by instrumenting code at compile time
 - Thus it's (usually) controlled by compiler flags
 - Can reduce performance
- Standard profiles return data on:
 - Number of function calls
 - Amount of time spent in sections of code
- Also tools that will return hardware specific data
 - Cache misses, TLB misses, cache re-use, flop rate, etc...
 - Useful for in-depth performance optimisation

- Standard Unix profilers are `prof` and `gprof`
- Many other profiling tools use same formats
- Usual compiler flags are `-p` and `-pg`:
 - `f90 -p mycode.F90 -o myprog` for `prof`
 - `cc -pg mycode.c -o myprog` for `gprof`
- When code is run it produces instrumentation log
 - `mon.out` for `prof`
 - `gmon.out` for `gprof`
- Then run `prof/gprof` *on your executable program*
 - eg. `gprof myprog` (*not* `gprof gmon.out`)

- `prof myprog` reads `mon.out` and produces this:

%Time	Seconds	Cumsecs	#Calls	msec/call	Name
32.4	0.71	0.71	14	50.7	<code>relax_</code>
28.3	0.62	1.33	14	44.3	<code>resid_</code>
11.4	0.25	1.58	3	83.	<code>__f90_close</code>
5.9	0.13	1.71	1629419	0.0001	<code>_mcount</code>
5.0	0.11	1.82	339044	0.0003	<code>__f90_slr_i4</code>
5.0	0.11	1.93	167045	0.0007	<code>__inrange_single</code>
2.7	0.06	1.99	507	0.12	<code>_read</code>
2.7	0.06	2.05	1	60.	<code>MAIN_</code>

- `gprof myprog` reads `gmon.out` and produces something very similar
- `gprof` also produces a program calltree sorted by inclusive times
- Both profilers list all routines, including obscure system ones
 - Of note: `mcount()`, `_mcount()`, `moncontrol()`, `_moncontrol()`, `monitor()` and `_monitor()` are all overheads of the profiling implementation itself
 - `_mcount()` is called every time your code calls a function; if it's high in the profile, it can indicate high function-call overhead

-
- Java also has profiling tools
 - jvisualvm
 - yourkit
 - Jprofiler
 - Jensor
 - Primarily connect to code in JVM
 - Profile memory usage and time spent in code section
-

- TAU – Tuning and Analysis Utilities
 - <http://www.cs.uoregon.edu/research/tau/home.php>
 - Portable profiling and tracing toolkit for performance analysis of programs written in Fortran, C, C++, Java, Python
 - Primarily for parallel (or multi-threaded) programs
- Heavier weight than standard profilers
 - Produces its own compiler wrappers
 - Based on code instrumentation
 - Instrumentation can be inserted in the source code automatically (like prof and gprof), dynamically, at runtime in the Java virtual machine, or manually using the instrumentation API
 - Can also access hardware counters through PAPI or PCL
 - Can be run through Eclipse, either Java TAU version or through Eclipse Parallel Toolkit
- Graphical visualisation tool
- Memory profiling built in

- Vampirtrace
 - Profiling library for MPI communications
 - Also allows instrumenting of code
 - Produces trace files in Open Trace Format (OTF)
 - Can be viewed in KOJAK, Vampir, TAU, etc...
- Lots of other examples
 - Craypat, ...

- PAPI – Performance API
 - <http://icl.cs.utk.edu/papi/index.html>
 - Set of tools built on it which can give you the info you want (i.e. perfsuite, TAU, KOJAK, etc...)
- AMD Tool - CodeXL
 - <http://developer.amd.com/tools-and-sdks/heterogeneous-computing/codexl/>
 - Hardware counters and general profiling
 - Open source
- Intel Tool – Vtune
 - [http://software.intel.com/en-us/intel-vtune-amplifier-xe /](http://software.intel.com/en-us/intel-vtune-amplifier-xe/)
 - Costs money
 - Includes old performance tuning utility (PTU)

- Instead of profiling a whole code, sometimes you want to time just one or two routines
- Add explicit timing calls to your code
 - More control over overheads
 - Less impact on overall code performance
 - Switch in and out with conditional compilation flags
- However, beware portability!
 - There is generally no such thing as a “standard”, “portable” timing function or routine
 - Resolution very important for timers
 - HPC platforms in particular have their own ways of timing
 - Platform X will have a good, high-resolution timer but it will be completely non-portable to platform Y

- In Fortran, two fairly standard library functions:
 - `dttime()`
 - Returns elapsed time since last call to `dttime`
 - `etime()`
 - Returns elapsed time since start of execution
 - NB: for parallel or multithreaded systems, read the manual pages carefully to check what the implications are
 - Although this is no guarantee... A quote from a certain HPC system:
"For `etime`, the elapsed time is:
 - For multiple processors: the wall-clock time while running your program"
 - **This lies**; it actually sums time over all threads; *not* what you want...

- **system_clock** – returns elapsed wallclock time
 - Takes 3 integer arguments:
 - **count, intent(out)**
 - **count_rate: intent(out)**
 - **count_max: intent(out)**
 - counts time at a rate of count_rate counts per second up to count_max
 - call `system_clock(count, count_rate, count_max)`
 - Elapsed time = $(\text{count2} - \text{count1}) * 1.0 / \text{count_rate}$
- Best function for C is probably
 - `gettimeofday(struct timeval *tp, void *)`;
 - Again, a library function
 - Need to `#include <sys/time.h>`
- Anyone writing MPI programs can use
 - `double MPI_Wtime(void)`
`double precision MPI_WTIME()`
 - Defined parts of the MPI 1 standard library
 - Not always implemented...
 - Return times local to the calling process

- Simple Unix/Linux solution for whole program runs is ‘**time**’:


```
$ time ./mg-test  
[program output]  
real    0m3.832s  
user    0m2.330s  
sys     0m0.210s
```

- Always check accuracy of your timers
 - e.g. above, real (elapsed) time is +/- 1s while user and system (CPU) times are +/- 0.01s
- Portable timing routines are good examples of things to put in your own utility libraries

-
- Understanding big codes is hard
 - But there are plenty of helpful tools
 - Learn to use them!
 - Profiling: look before you leap
 - Don't attempt any optimisation work until you've profiled your code
 - Learn to use and understand prof and gprof
 - More hardware specific data is available
 - Timing by hand is very useful, but beware portability of timers


- Displaying and manipulating performance data from HPC code
 - Easy to draw a simple graph
 - “Real” story not always obvious from basic graphs
 - Need to understand data and what you want to show
- Basic example, principles apply to all cases
 - Always important to understand that data can be shown in many different ways
 - Errors should always be considered
 - Performance variation


- Once HPC code run data must be analysed
 - Often most important step
 - Visualisation can be useful
- Tables, graphs, pictures, animations, etc...
 - Many ways to show data, each useful for particular scenarios
 - Always focus on what you want to display
 - Often, hardest part of visualisation is choosing which technique/display method to use
- We often focus on performance data
 - How data collected is important
 - Errors must be considered here too




Performance Graph Example


Number of Processors	Total Runtime (seconds)
1	23.6
2	20.5
4	17.8
8	14.7
16	13.9








Time for Completion with Ideal Plot




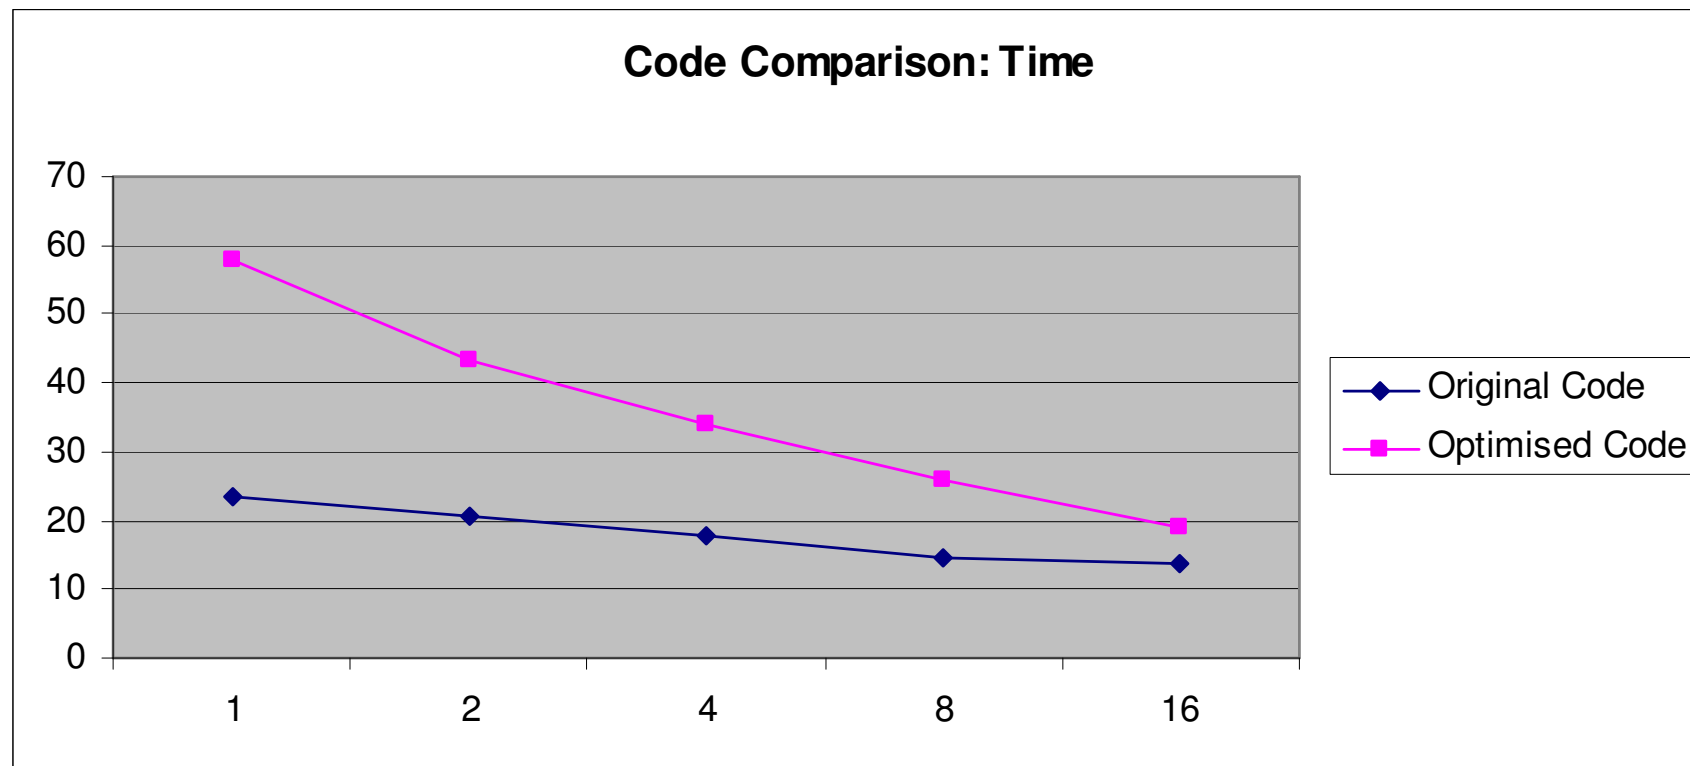









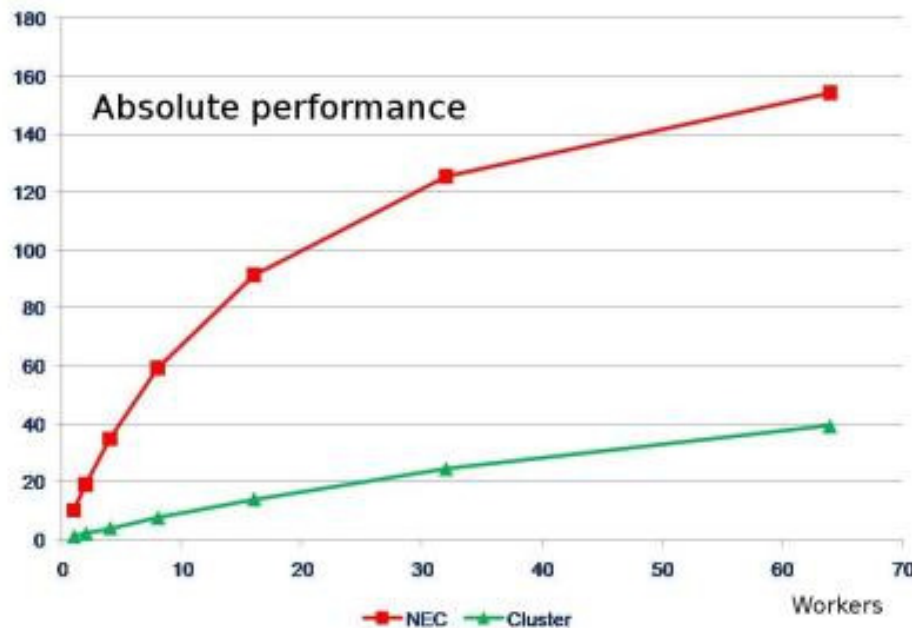
Code Comparison: Log Speedup



Number of Processors	Original Code Runtime (seconds)	New Code Runtime (seconds)
1	23.6	57.9
2	20.5	43.2
4	17.8	33.9
8	14.7	25.8
16	13.9	19.1

- <http://blogs.fau.de/hager/category/fooling-the-masses/>
- David H. Bailey: “Twelve Ways to Fool the Masses When Giving Performance Results on Parallel Computers” (1991)
 - Quote only 32-bit performance results, not 64-bit results.
 - Present performance figures for an inner kernel, and then represent these figures as the performance of the entire application.
 - Quietly employ assembly code and other low-level language constructs.
 - Scale up the problem size with the number of processors, but omit any mention of this fact.
 - Quote performance results projected to a full system.
 - Compare your results against scalar, unoptimized code on Crays.
 - When direct run time comparisons are required, compare with an old code on an obsolete system.
 - If MFLOPS rates must be quoted, base the operation count on the parallel implementation, not on the best sequential implementation.
 - Quote performance in terms of processor utilization, parallel speedups or MFLOPS per dollar.
 - Mutilate the algorithm used in the parallel implementation to match the architecture.
 - Measure parallel run times on a dedicated system, but measure conventional run times in a busy environment.
 - If all else fails, show pretty pictures and animated videos, and don’t talk about performance.



- **Stunt 1: Report speedup, not absolute performance!**



- Stunt 2: Slow down code execution!

$$S(N) = \frac{s + (1 - s)N^\alpha}{s + (1 - s)N^{\alpha-1} + c_\alpha(N)}$$

- Stunt 3: The log scale is your friend!



- Context:
 - Always important to report details of performance data collection:
 - Machine hardware
 - OS
 - Third party tools
 - Software version
 - Technique for collecting data (average, fastest, slowest)